

Games

Introduction to Artificial Intelligence

G. Lakemeyer

Winter Term 2016/17

Games

Games like chess have been studied in AI since the 50-ies.

They are a special **variant of search problems**.

The states are usually **accessible**.

The **actions** are the possible moves of a player.

Complication: More than 1 actor.

From one player's point of view, his or her actions have **uncertain** outcomes, since the reaction of the opponent is usually not foreseeable.

In this sense games belong to the class of **contingency problems** (uncertain knowledge of action effects).

Note: The opponent usually chooses actions to do maximal damage.

What Makes Games Interesting

Problem: Almost all interesting games are unsolvable in practice.

Chess:

Each player has about 50 moves with about 35 actions per move, i.e. there are about 35^{100} nodes in the search tree (with “only” 10^{40} legal chess positions).

Good game programs have the following features:

- a) Early **pruning** of useless sub-trees.
- b) good **evaluation function** of states without doing a complete search.

2-Person Games – Some Terminology

Players: MAX and MIN with MAX doing the first move.

Initial state: e.g. initial board position, assignment of MAX and MIN.

Operators: legal moves.


Terminal test: when a game ends.

Terminal state game over.

Utility function (payoff): gives a numeric value to the outcome of a game; often simply +1 (win), -1 (loss), 0 (draw). Backgammon has a range between +192 and -192.

Strategy: in contrast to regular search, MAX needs to find a path which leads to a winning terminal state for *every* possible reaction of MIN.

An Example: Tic-Tac-Toe




Each level of the search tree (also called **game tree**) is labelled with the player whose turn it is (MAX and MIN levels). When the complete game tree can be generated, the optimal strategy for MAX can be computed.

Minimax

MAX

MIN



- 1 Generate a complete game tree.
- 2 Apply utility function to each terminal state.
- 3 Starting from the terminal states, calculate the values for the parent node as follows:
 - If the parent node is at a MIN level, then assign it the **minimum** of the values of the children.
 - If the parent node is at a MAX level, then assign it the **maximum** of the values of the children.
 - Then at the root MAX chooses the action which leads to the child with maximum utility (**Minimax decision**).

Note: Minimax assumes that MIN plays perfectly rational, i.e. always chooses the optimal move.

Minimax-Algorithm

use, for ~~α~~-
DFS.

function MINIMAX-DECISION(*game*) *returns an operator*

for each *op* **in** OPERATORS[*game*] **do**

 VALUE[*op*] \leftarrow MINIMAX-VALUE(APPLY(*op*, *game*), *game*)

end

return the *op* with the highest VALUE[*op*]

function MINIMAX-VALUE(*state*, *game*) *returns a utility value*

~~cut-off test~~
~~**if** TERMINAL-TEST[*game*](*state*) **then**~~

Eval

~~**return** UTILITY[*game*](*state*)~~

else if MAX is to move in *state* **then**

return the highest MINIMAX-VALUE of SUCCESSORS(*state*)

else

return the lowest MINIMAX-VALUE of SUCCESSORS(*state*)

Evaluation Functions

When the search space is large, the game tree can only be generated up to a certain depth.

Minimax works then as well. Simply replace `TERMINAL-TEST` by `CUT-OFF-TEST` and the utility function `UTILITY` by the evaluation function `EVAL`.

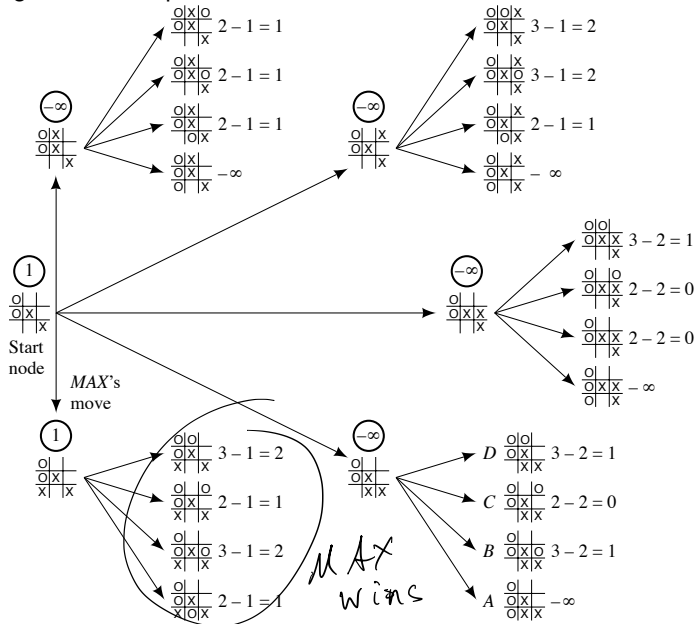
The trick is then to correctly evaluate the goodness of the leaves.

Simple criteria for **chess**:

- Material value: pawn = 1, knight = 3, bishop = 3, rook = 5, queen = 9.
- Other features: safety of the king, good pawn structure.
- Rule of thumb: 3-point advantage = certain victory.

Example: Tic-Tac-Toe (2)

Final stage, search depth 2



Form of Evaluation Functions

The choice of an evaluation function is critical.

It should be **easy to compute** and **accurately reflect the chance of winning**.

Chance of winning for a given material value means the probability to win averaged over all positions with the same material value.

Usually evaluation functions are **weighted linear functions**:

$$w_1 f_1 + w_2 f_2 + \dots + w_n f_n$$


(E.g. MAX = White: $w_1 = 3$, f_1 = number of white knights on the board.)

Assumption: The criteria are independent of each other. (Simplification)


When to Cut Off Search?

Cut-off only in **quiescent** states, i.e. those which do not lead to dramatic subsequent changes, since the evaluation function is otherwise misleading.


Example:




(a) White to move
Fairly even



(b) Black to move
White slightly better



(c) White to move
Black winning




(d) Black to move
White about to lose

white has
mat. adv.
of +3

mat. adv.
is +5

Horizon Problem




Black to move

- Black has a slight advantage in material value;
- black loses the game eventually (pawn turns into a queen eventually);
- fixed depth search may not detect this since the “disaster” can be pushed over the horizon.

Alpha Beta

MAX

MIN



Player

Opponent


..

..

..

Player

Opponent



Motto: if you find out that things are bad, do not waste effort to see how bad it really is.

Alpha-Beta Search Algorithm

function MAX-VALUE(*state*, *game*, α , β) **returns** the minimax value of *state*

inputs: *state*, current state in game

game, game description

α , the best score for MAX along the path to *state*

β , the best score for MIN along the path to *state*

if CUTOFF-TEST(*state*) **then return** EVAL(*state*)

for each *s* **in** SUCCESSORS(*state*) **do**

$\alpha \leftarrow \text{MAX}(\alpha, \text{MIN-VALUE}(s, \text{game}, \alpha, \beta))$

if $\alpha \geq \beta$ **then return** β

end

return α

function MIN-VALUE(*state*, *game*, α , β) **returns** the minimax value of *state*

if CUTOFF-TEST(*state*) **then return** EVAL(*state*)

for each *s* **in** SUCCESSORS(*state*) **do**

$\beta \leftarrow \text{MIN}(\beta, \text{MAX-VALUE}(s, \text{game}, \alpha, \beta))$

if $\beta \leq \alpha$ **then return** α

end

return β

Initialization: $\alpha = -\infty$, $\beta = +\infty$.

Max-Value is applied to nodes at MAX levels, *Min-Value* at MIN levels.

Alpha-Beta Search Algorithm Version 2

function ALPHA-BETA-SEARCH(*state*) **returns** an action

inputs: *state*, current state in game

$v \leftarrow \text{MAX-VALUE}(\text{state}, -\infty, +\infty)$

return the *action* in SUCCESSORS(*state*) with value *v*

function MAX-VALUE(*state*, α , β) **returns** a utility value

inputs: *state*, current state in game

α , the value of the best alternative for MAX along the path to *state*

β , the value of the best alternative for MIN along the path to *state*

if TERMINAL-TEST(*state*) **then return** UTILITY(*state*)

$v \leftarrow -\infty$

for *a*, *s* in SUCCESSORS(*state*) **do**

$v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(s, \alpha, \beta))$

if $v \geq \beta$ **then return** *v*

$\alpha \leftarrow \text{MAX}(\alpha, v)$

return *v*

function MIN-VALUE(*state*, α , β) **returns** a utility value

inputs: *state*, current state in game

α , the value of the best alternative for MAX along the path to *state*

β , the value of the best alternative for MIN along the path to *state*

if TERMINAL-TEST(*state*) **then return** UTILITY(*state*)

$v \leftarrow +\infty$

for *a*, *s* in SUCCESSORS(*state*) **do**







$v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(s, \alpha, \beta))$

if $v \leq \alpha$ **then return** *v*

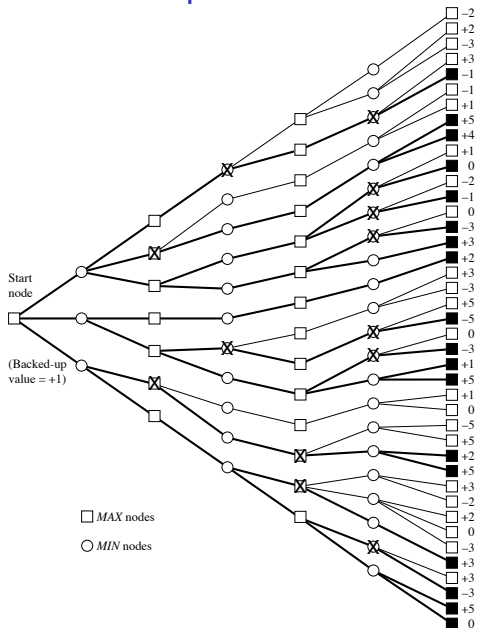
$\beta \leftarrow \text{MIN}(\beta, v)$

return *v*


Alpha-Beta Example



Alpha-Beta Example




Alpha-Beta: the Ideal Case



© 1992 Patrick H. Winston

© G. Lakemeyer


$$19 \rightarrow 24, 5 \rightarrow 11$$
$$5 \rightarrow 10, 10 \rightarrow 16$$
$$5 \rightarrow 11, 11 \rightarrow 16$$
$$\sum_i p(d_i) \neq \max_{s \in S(C, d_i)} \{ \text{value}(s) \}$$

$d_i = i$ -th distinct word

$S(C, d_i)$ = set of legal states reachable for all d_i



Chess Programs

Flo
rating

