

# Neural Networks

## Introduction to Artificial Intelligence

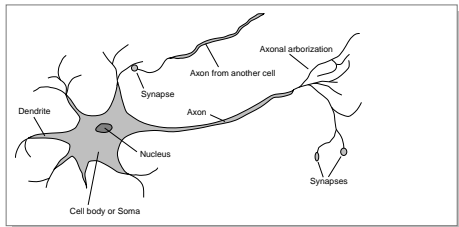
G. Lakemeyer

Winter Term 2016/17

# Brain vs. Computer

## Analogy to the human brain:

many processors (**neurons**)  
and connections (**synapses**),  
which process information  
locally and in parallel.



## von Neumann computer vs. the brain:

	Computer	Human Brain
Computational units	1 CPU, $10^5$ gates	$10^{11}$ neurons
Storage units	$10^9$ bits RAM, $10^{10}$ bits disk	$10^{11}$ neurons, $10^{14}$ synapses
Cycle time	$10^{-8}$ sec	$10^{-3}$ sec
Bandwidth	$10^9$ bits/sec	$10^{14}$ bits/sec
Neuron updates/sec	$10^5$	$10^{14}$

**Note:** The brain cycle is slow ( $10^{-3}$ s), but updates work in parallel. A serial simulation on a computer needs several hundred cycles.

# Advantages of Neural Networks

- High processing speed through massive parallelism;
- still works if parts of the network are damaged (robustness);
- graceful degradation;
- designed for inductive learning.

It seems reasonable (obvious?) to try and recreate these advantages by designing artificial neural networks.

Research in this area started already in the 40-ies.  
(McCulloch und Pitts 1943)

# Some Terminology of Neural Networks

**Units:** represent the nodes (neurons) in the network.

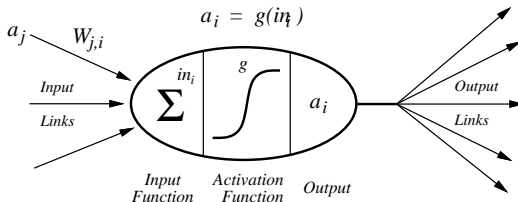
**Input/output units:** Special nodes connected to the external environment.

**Links:** Edges between nodes in the network. Each node has input and output edges.

**Weights:** Each edge has a weight, usually a real number.

**Activation level:** The value computed by a unit given its weighted inputs. It is passed to the neighbor nodes via the output edges.

# How a Unit Works

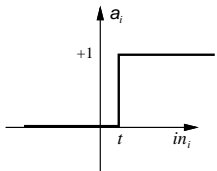


$$a_i := g(in_i) = g \left( \sum_j w_{j,i} \times a_j \right)$$

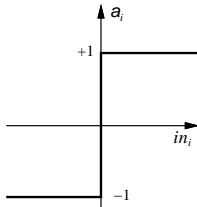
$g$  is usually a **nonlinear** function.

# Activation Functions

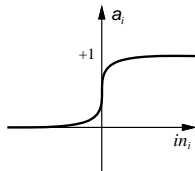
Some important activation functions are:



(a) Step function



(b) Sign function



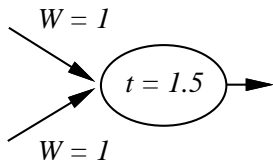
(c) Sigmoid function

**Note:**  $t$  in *step* represents a **threshold**, that is, the minimum total weighted input necessary for the neuron to fire (similar to actual neurons in the brain).

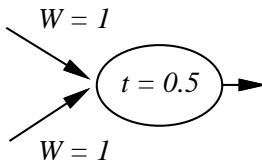
Mathematically, one can always use a threshold of 0 by having an additional input with activation level  $a_0 = -1$  and  $W_{0,i} = t$ . Then

$$a_i = \text{step}_t \left( \sum_{j=1}^n W_{j,i} \times a_j \right) = \text{step}_0 \left( \sum_{j=0}^n W_{j,i} \times a_j \right)$$

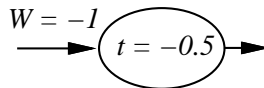
# Representing Boolean Functions



**AND**



**OR**



**NOT**

Thus neural nets can at least represent any arbitrary Boolean function.

[We will see below that they can represent much more.]


# Network Topologies

**Feed-forward:** Units and links represent a directed acyclic graph.

**Recurrent:** Units and links represent arbitrary directed graphs.

Nets are often arranged in **layers**, that is, nodes of one layer are linked only to nodes of the next layer. There are no links between nodes of a layer. (When counting layers, the input layer is ignored.)

## Example of a feed-forward 2-layer network:



Of all networks feed-forward nets are the best understood. (The output is solely a function of the inputs and the weights.) Recurrent networks are often unstable, oscillate, or have chaotic behavior.

**Note:** The brain is massively recurrent.



# Recurrent Network Types – Hopfield Nets

- Recurrent nets with symmetric bi-directional links ( $W_{i,j} = W_{j,i}$ ).
- All units are both input and output units.
- Nets function as **associative memory**.
- After training a new input is matched against the “closest” example seen during training.

## Example:

Training with  $n$  photographs.

New input = part of a photograph used during training.

Then the network reconstructs the whole image.

**Note:** Each weight is a partial representation of every image!


## Theorem

A Hopfield Net can reliably store  $0,138 \times N$  examples, where  $N$  is the number of units.

# Recurrent Network Types – Boltzmann Machines

- Also use symmetric weights.
- There are units which are neither input nor output units.
- Stochastic Activation function.
- State transitions are like simulated annealing search for the configuration that best approximates the training set.  
(Formally identical to a certain kind of belief nets.)

# Feed-Forward (FF) Nets



## Hidden Units:

Units without direct connection to the external environment. Hidden units are organized in one or more hidden layers.

## Perceptrons:

FF Networks without hidden units.

FF Nets represent complex nonlinear functions.

- With 1 hidden layer every **continuous function** is representable.
- With 2 hidden layers **every function** is representable.

When the topology and activation function  $g$  are fixed, the representable functions have a specific parametrized form (parameters = weights).

## Example:

$$a_5 = g(W_{3,5}a_3 + W_{4,5}a_4) = g(W_{3,5}g(W_{1,3}a_1 + W_{2,3}a_2) + W_{4,5}g(W_{1,4}a_1 + W_{2,4}a_2))$$

$\Rightarrow$  **Learning** = Search for the correct parameters = **nonlinear regression**.

# Optimal Network structure?

Choosing the right network is a difficult problem!

Also the optimal net may be exponentially large (relative to the input).

- Net is too small: the desired function is not representable.
- Net is too big: Net memorizes examples without generalization (analogous to memorizing in decision trees) **Overfitting**

There is no good theory of how to choose the right network, only some heuristics.

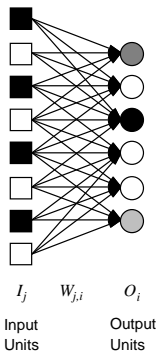
**Heuristics of optimal brain damage:**

Start with a network with a maximal number of connection. After the 1st training reduce the number of connections using information theory. Iterate.

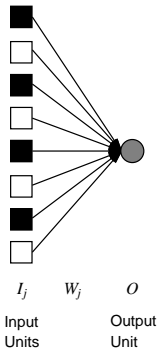
## Example:

Network to recognize zip codes. 3/4 of the initial connections could be removed. (There are also methods to move from a network with few nodes to a network with more nodes.)

# Perceptrons



**Perceptron Network**



**Single Perceptron**

$$O = \text{Step}_0 \left( \sum_j W_j I_j \right) = \text{Step}_0(\mathbf{W} \cdot \mathbf{I})$$

# How Do Perceptrons Learn?

**function** NEURAL-NETWORK-LEARNING(*examples*) **returns** *network*

*network*  $\leftarrow$  a network with randomly assigned weights

**repeat**

**for each** *e* **in** *examples* **do**

$\mathbf{O} \leftarrow$  NEURAL-NETWORK-OUTPUT(*network*, *e*)

$\mathbf{T} \leftarrow$  the observed output values from *e*

    update the weights in *network* based on *e*,  $\mathbf{O}$ , and  $\mathbf{T}$

**end**

**until** all examples correctly predicted or stopping criterion is reached

**return** *network*

Error =  $T - O$  with

$O$  = predicted (actual) output

$T$  = correct output

Update-Rule:  $W_j := W_j + \alpha \times I_j \times \text{Error}$

**Theorem:** [Rosenblatt 1960]

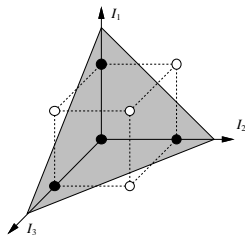
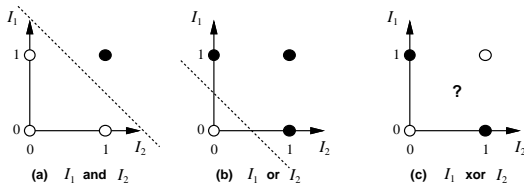
Learning using the update rule always converges to the correct output for representable functions.

# What can Perceptrons represent?

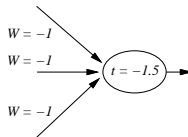
Answer: Not much!

Perceptrons can represent exactly the class of **linearly separable functions**.

Examples:

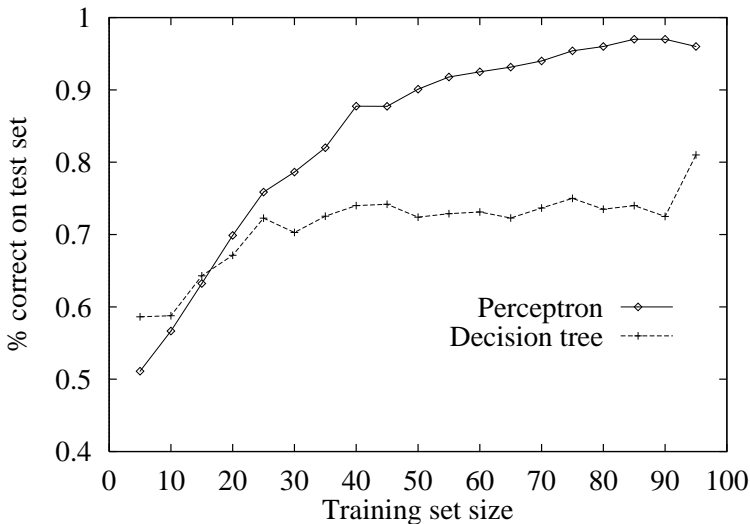


(a) Separating plane



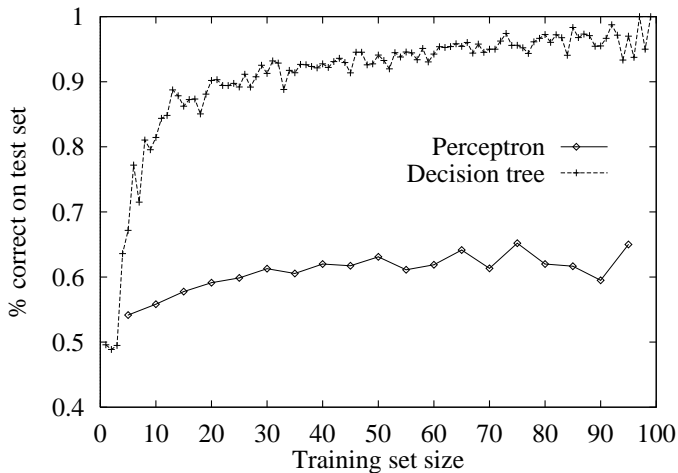
(b) Weights and threshold

# Learning Behavior I: Majority function

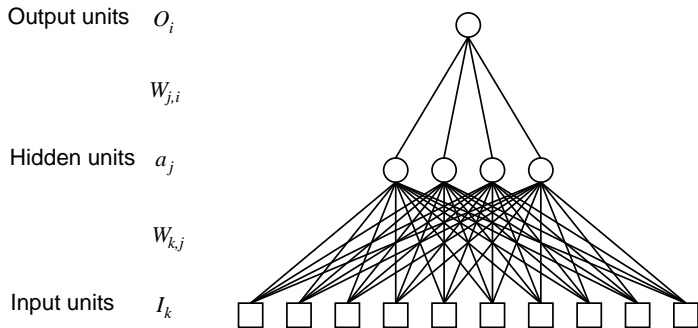




# Learning Behavior II: Restaurant Example



# Multi-layer Feed-Forward Networks



(This topology is suitable for the restaurant example.)

# Back-Propagation

**function** BACK-PROP-UPDATE(*network*, *examples*,  $\alpha$ ) **returns** a network with modified weights

**inputs:** *network*, a multilayer network  
*examples*, a set of input/output pairs  
 $\alpha$ , the learning rate

**repeat**

**for each** *e* **in** *examples* **do**

/\* Compute the output for this example \*/

$\mathbf{O} \leftarrow \text{RUN-NETWORK}(\text{network}, \mathbf{I}^e)$

/\* Compute the error and  $\Delta$  for units in the output layer \*/

$\text{Err}^e \leftarrow \mathbf{T}^e - \mathbf{O}$

/\* Update the weights leading to the output layer \*/

$W_{j,i} \leftarrow W_{j,i} + \alpha \times a_j \times \text{Err}_i^e \times g'(in_i)$

**for each** subsequent layer **in** *network* **do**

/\* Compute the error at each node \*/

$\Delta_j \leftarrow g'(in_j) \sum_i W_{j,i} \Delta_i$

/\* Update the weights leading into the layer \*/

$W_{k,j} \leftarrow W_{k,j} + \alpha \times I_k \times \Delta_j$

**end**

**end**

**until** *network* has converged

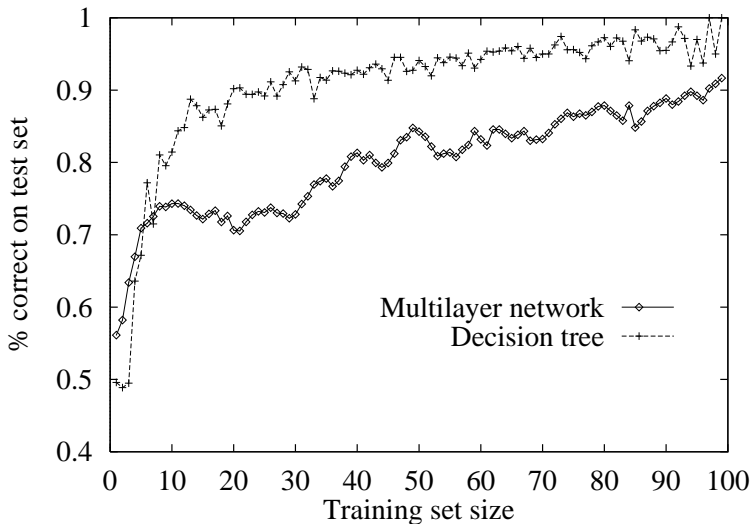
**return** *network*

$\text{Err}_i^e = \text{Error } T_i^e - O_i^e$  of the *i*-th output unit.


$\text{Err}^e$  = error vector.

$\Delta_i = \text{Error}_i \times g'(in_i)$

# Restaurant Example



# Restaurant Example



# Summary

- What are they good for?

Neural nets are useful for **attribute-based** representations (like decision trees), in particular also for attributes with **continuous** values.

- Size of a net:

$2^n/n$  units to represent any Boolean function.

( $\Rightarrow 2^n$  weights). In practice one often needs much less.

- Efficiency of Learning:

no useful theoretical results.

Also a problem with local minima. **Simulated annealing** sometimes helpful (see Boltzmann Machines)

- Generalization: Good if the right network is used :-)

- Noise: no problem.

- Transparency:


bad! Often one has no idea why a network yields a certain output.  
Decision trees are much better in this regard!

- Using additional knowledge: bad. No theoretical results.

# Application: Pronunciation

**NETtalk** synthesizes speech from written text (in English) [Sejnowski and Rosenberg 1987]

**Input:** seven-character window sliding over the written text.:



uses **29 input units** per character (for the 26 letters of the alphabet, blank, comma, and 1 unit for other characters)

**Hidden layer:** 80 units

**Output:** Units to encode phonemes.

**Training set:** text with 1024 words. After 50 epochs 95% correctness.

**Problems:** words with identical spelling but different pronunciation (e.g. lead).

**Result on test set:** 78% correctness.

(Today there are better systems not based on neural nets, but NETtalk paved the way for the commercial success of neural nets.)

# Application: Recognizing Zip Codes

Neural net to recognize handwritten numbers (zip codes) [Le Cun 89]

**Input:**  $16 \times 16$  pixels per digit.

**3 hidden layers:** 768; 192; 30 units.

**Output:** 10 units for digits 0–9.

Limitation of connections was crucial:

- each unit of the 1st layer was connected with a  $5 \times 5$  array of the input (25 connections).
- First layer was organized in 12 groups with 64 units each. All units of the same group use identical weights.
- The whole network used only 9760 different weights instead of 200,000.

**Training set:** 7300 examples.

**Test set:** 2000 examples; 99% correctness!

Is used by US Postal Service, realized on VLSI.