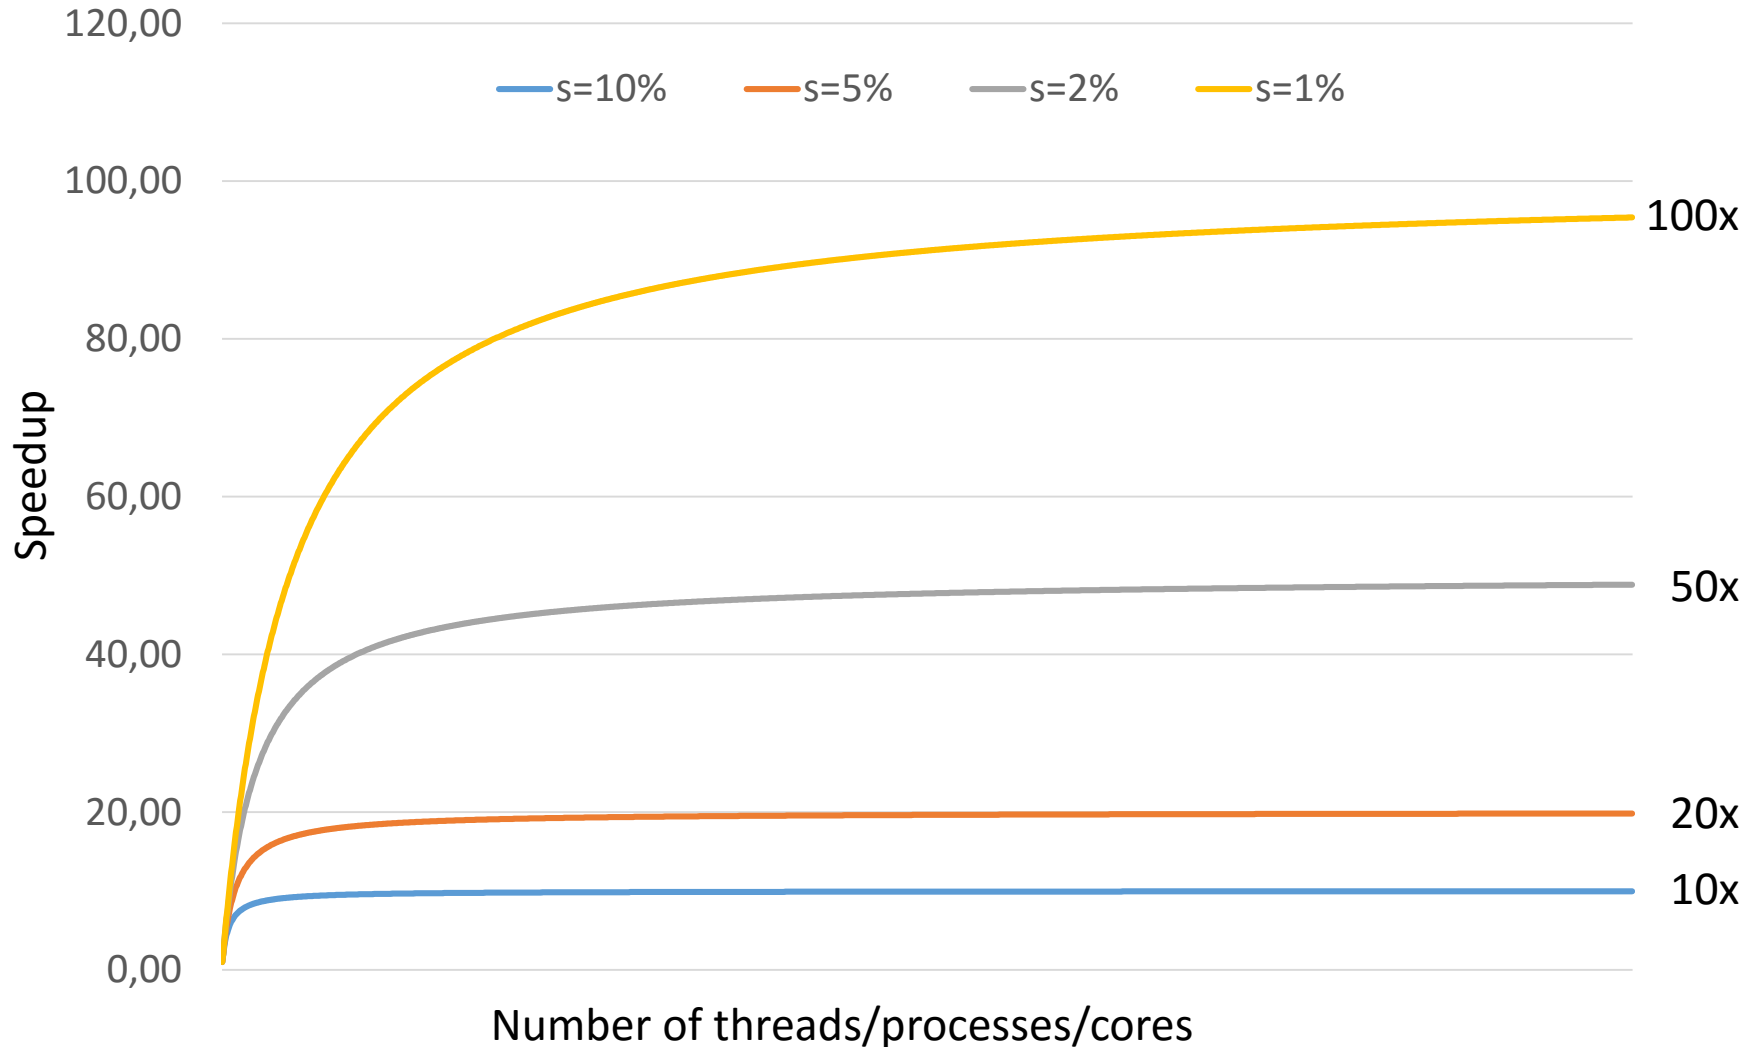


Parallel Algorithms




- **Myth: Parallel programming is too hard for me. There are deadlocks, race conditions, non-deterministic behavior. That's too much for me!**
- **Well, you are right. Partially.**
- **Structured programming and patterns help tame the problem.**
 - It does not magically go away, but is much less problematic
- **Apply the well-known Software Engineering practices:**
 - Identify a (parallel) pattern
 - Use the pattern library and apply a matching pattern from it

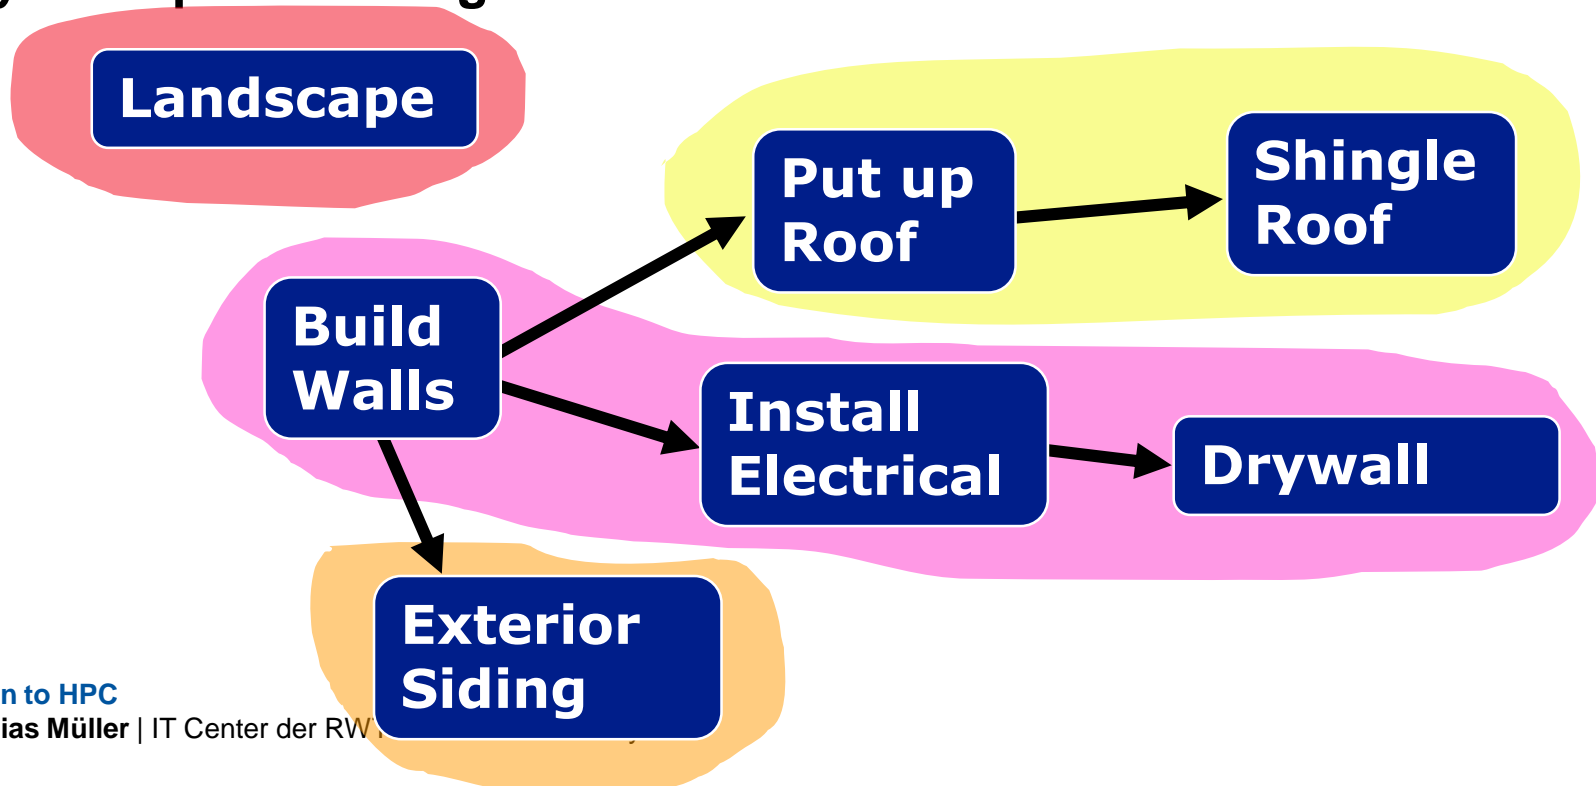
- **Patterns are “best practices” for solving specific problems**
 - Known from Software Engineering
- **Patterns can be used to organize your code, leading to algorithms that are more scalable and maintainable.**
- **A pattern (often) supports a particular “algorithmic structure” with an efficient implementation.**
- **Good parallel programming models support a set of useful parallel patterns with low-overhead implementations.**

Finding Concurrency


- Look for tasks that can be executed simultaneously (task parallelism)



- Catchy example: building a house




- Decompose data into distinct chunks to be processed independently (data parallelism)



- Catchy example: work on a farm

- each worker takes an (unpicked) row and picks the crop
- possibly more rows than worker
- process continues until all rows have been harvested




Example:
ventricular
assist device
(VAD)

■ **The following additional parallel patterns can be used for “structured parallel programming”:**


- Superscalar sequence
- Speculative selection
- Map
- Recurrence
- Scan
- Reduce
- Pack/expand
- Fork/join
- Pipeline
- Partition
- Segmentation
- Stencil
- Search/match
- Gather
- Merge scatter
- Priority scatter
- Permutation scatter
- Atomic scatter

Parallel Patterns: Overview


Superscalar sequence




Speculative selection




Map



Stencil




Fork-Join




Pipeline




Geometric decomposition




Partition




Gather




Scatter




Category Reduction




Recurrence




Reduction




Scan




Expand



Pack



Split



Patterns:

Parallel Control


- **Fundamental pattern to express and exploit parallelism**

- control flow *forks* into multiple flows that *join* later
- well-versed to go from serial to parallel
- can be nested

- **As an algorithmic pattern it is a natural fit for parallel divide-and-conquer approaches**

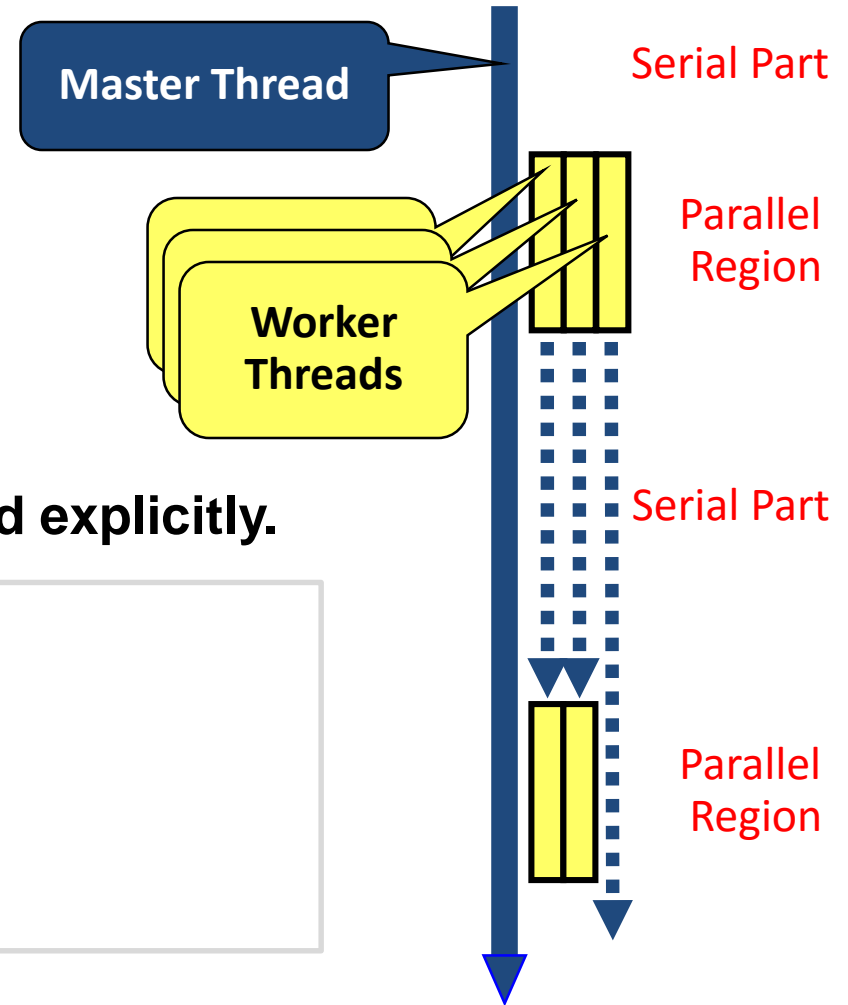
- **(Subtle) Differences among Intel® Cilk Plus and OpenMP*:**

- Intel Cilk Plus only uses tasks and an implicit threadpool
- OpenMP allows for tasks and threads
- OpenMP requires a parallel region to apply any other parallel construct



- OpenMP programs start with just one thread: the *master*.
- *Worker* threads are spawned at *parallel regions*, together with the master they form the *team* of threads.
- The parallelism has to be expressed explicitly.

```
#pragma omp parallel
{
    ...
    structured block
    ...
}
```



■ Spawn = asynchronous function call

→ Arguments to spawned functions are evaluated before fork

■ Example (with optional assignment):

```
x = cilk_spawn f(*p++);  
y = g(*p--);  
cilk_sync;  
z = x+y;
```

■ Expression spawning possible via lambda expression:

```
cilk_spawn [&]{  
    for(int i=0; i<n; ++i )  
        a[i] = 0;  
} ();  
...  
cilk_sync;
```

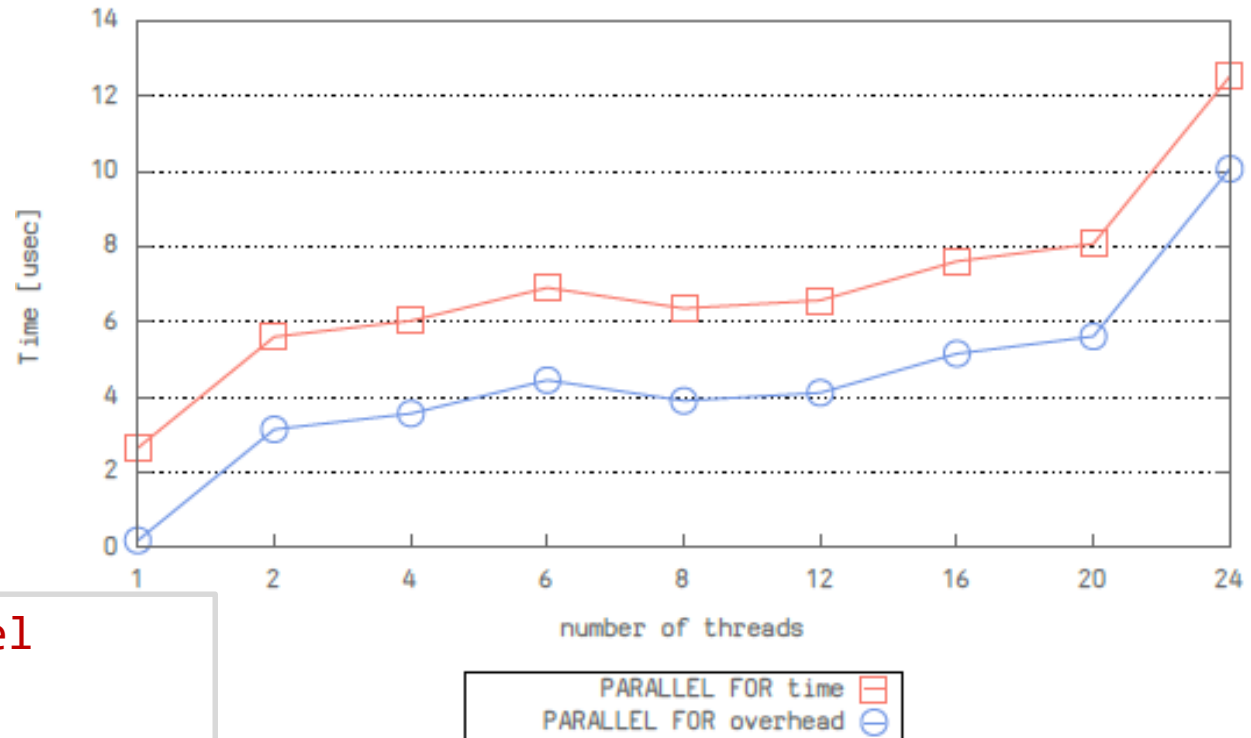
■ How expensive is Fork-Join? OpenMP Overhead Analysis via EPCC micro-benchmarks for a parallel for construct:

Intel® Cilk™ Plus


```
cilk_spawn a();  
cilk_spawn b();  
c();  
cilk_sync();
```

OpenMP*

```
#pragma omp parallel  
{  
    ...  
}
```



- Replicate a function over every element of an “index set”
- The index set typically corresponds to the elements of an array



- Some programming languages have support for the map pattern
 - e.g., Fortran, Intel® Cilk™ Plus

```
module example
  implicit none
contains
  subroutine example
    integer,parameter :: len=100
    real,dimension(len) :: in=/(i,i=1,len)/
    real,dimension(len) :: out
    integer :: i

    ! explicit do loop
    do i=1,size(in)
      output(i)=sqrt(in(i))
    end do

    ! explicit array subsets
    output(:)=sqrt(in(:))

    ! whole array operation
    output=sqrt(in)
    write (*,*), output
  end subroutine
end module
```



```
module example
  implicit none
contains
  subroutine example
    integer,parameter :: len=100
    real,dimension(len) :: in=/(i,i
    real,dimension(len) :: out
    integer :: i

    ! explicit do loop
    do i=1,size(in)
      output(i)=sqrt(in(i))
    end do

    ! explicit array subsets
    output(:)=sqrt(in(:))

    ! whole array operation
    output=sqrt(in)
    write (*,*), output
  end subroutine
end module
```

```
void ex_vector() {
  const size_t len = 100;
  size_t i = 0;
  std::vector<float> in(len);
  std::vector<float> out(len);

  std::generate_n(in.begin(), len, gen);

  std::transform(in.begin(), in.end(),
                out.begin(), sqrt);

  for (auto it = out.begin();
        it != out.end(); ++it, ++i)
    printf("%f%s",
          *it,
          (i%4==3) ? "\n" : " ");
}
```

```
module example
  implicit none
contains
  subroutine example
    integer,parameter :: len=100
    real,dimension(len) :: in
    real,dimension(len) :: output
    integer :: i

    ! explicit do loop
    do i=1,size(in)
      output(i)=sqrt(in(i))
    end do

    ! explicit array operation
    output(:)=sqrt(in(:))

    ! whole array operation
    output=sqrt(in)
    write (*,*), output
  end subroutine
end module
```

```
void ex_vector() {
  const size_t len = 100;
  size_t i = 0;
  std::vector<float> in(len);
  std::vector<float> out(len);
```

Did you see this
pattern in any other
programming
language you have
used?

```
    in.begin(), len, gen);
    t.begin(), in.end(),
    t.begin(), sqrt);
    t.begin();
    ut.end(); ++it, ++i)

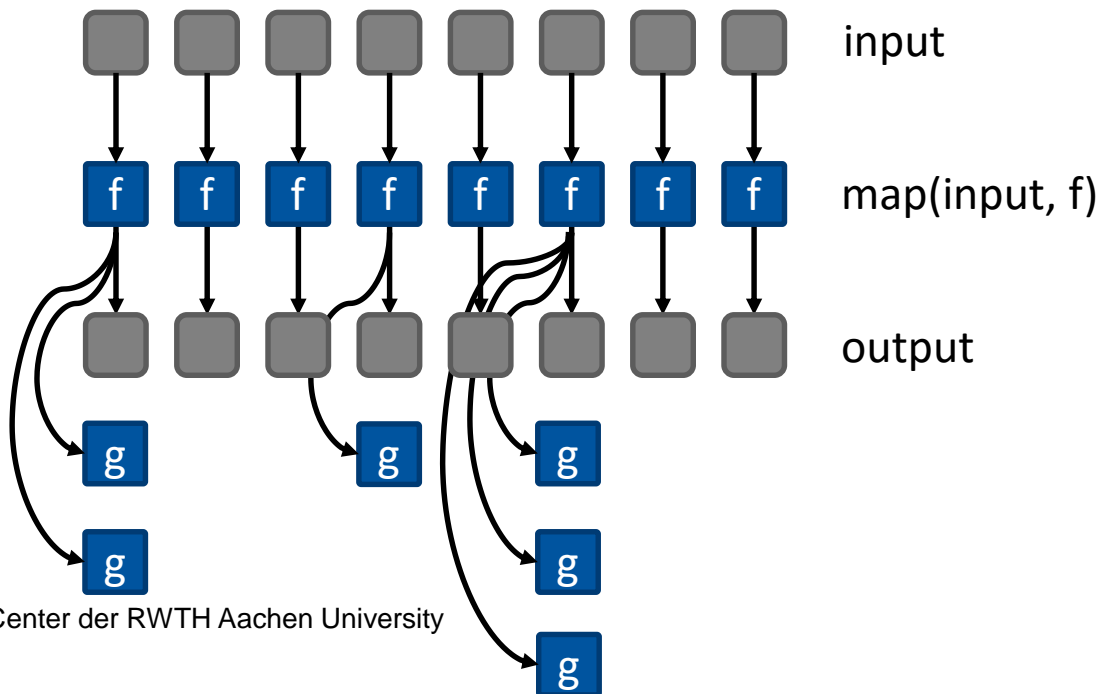
    ? "\n" : " ");
```

■ The Workpile pattern generalizes the Map pattern

- The mapped function can create new parallel work while it is applied
- Number of items to work on is not known in advance

■ Example: tree traversal

- Apply Map pattern to root node; create new tasks for a node's children



■ Lets solve Sudoku puzzles with brute multi-core force

	6						8	11			15	14			16
15	11				16	14				12			6		
13		9	12					3	16	14		15	11	10	
2		16		11		15	10	1							
	15	11	10			16	2	13	8	9	12				
12	13			4	1	5	6	2	3					11	10
5		6	1	12		9		15	11	10	7	16			3
	2				10		11	6		5			13		9
10	7	15	11	16				12	13						6
9						1			2		16	10			11
1		4	6	9	13			7		11		3	16		
16	14			7		10	15	4	6	1				13	8
11	10		15				16	9	12	13			1	5	4
		12		1	4	6		16				11	10		
		5		8	12	13		10			11	2			14
3	16			10			7			6				12	

(1) Find an empty field

(2) Insert a number

(3) Check Sudoku

(4 a) If invalid:

**Delete number,
Insert next number**

(4 b) If valid:

Go to next field

- This parallel algorithm finds all valid solutions

first call contained in a
`#pragma omp parallel`
`#pragma omp single`
 such that one tasks starts the
 execution of the algorithm

`#pragma omp task`
 needs to work on a new copy
 of the Sudoku board

`#pragma omp taskwait`

`#pragma omp taskwait`

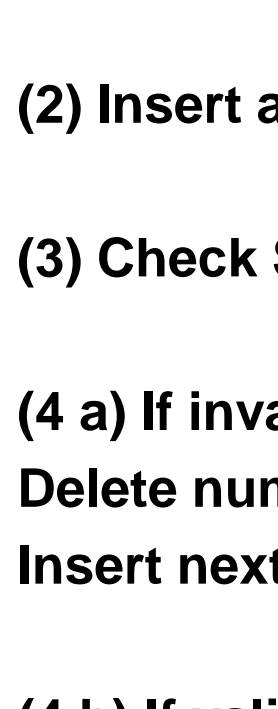
first call contained in a

```
#pragma omp parallel
#pragma omp single
```

such that one tasks starts the
execution of the algorithm


#pragma omp task
needs to work on a new copy
of the Sudoku board

```
#pragma omp taskwait
wait for all child tasks
```


- 
- { (1) Search an empty field**
 - { (2) Insert a number**
 - { (3) Check Sudoku**
 - { (4 a) If invalid:
Delete number,
Insert next number**
 - { (4 b) If valid:
Go to next field**

Sudoku on 2x Intel Xeon E5-2650 @2.0 GHz

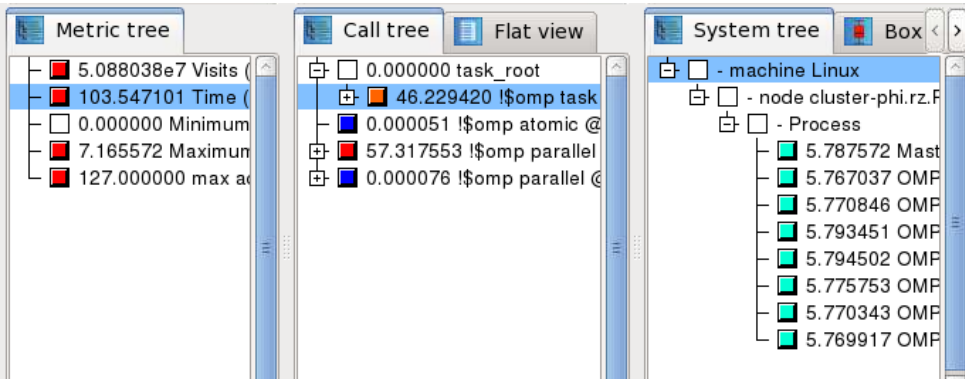
■ Intel C++ 13.1, scatter binding ■ speedup: Intel C++ 13.1, scatter binding



Event-based profiling gives a good overview :






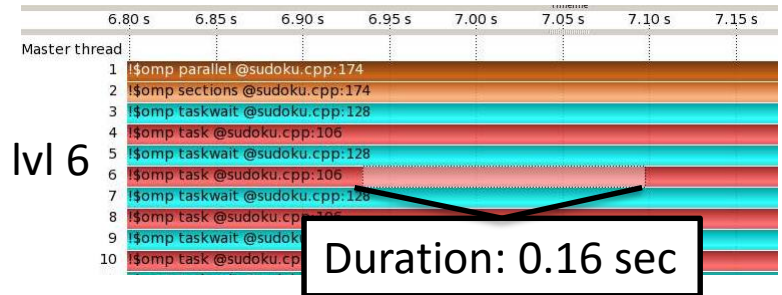
Every thread is executing ~1.3m tasks...



... in ~5.7 seconds.

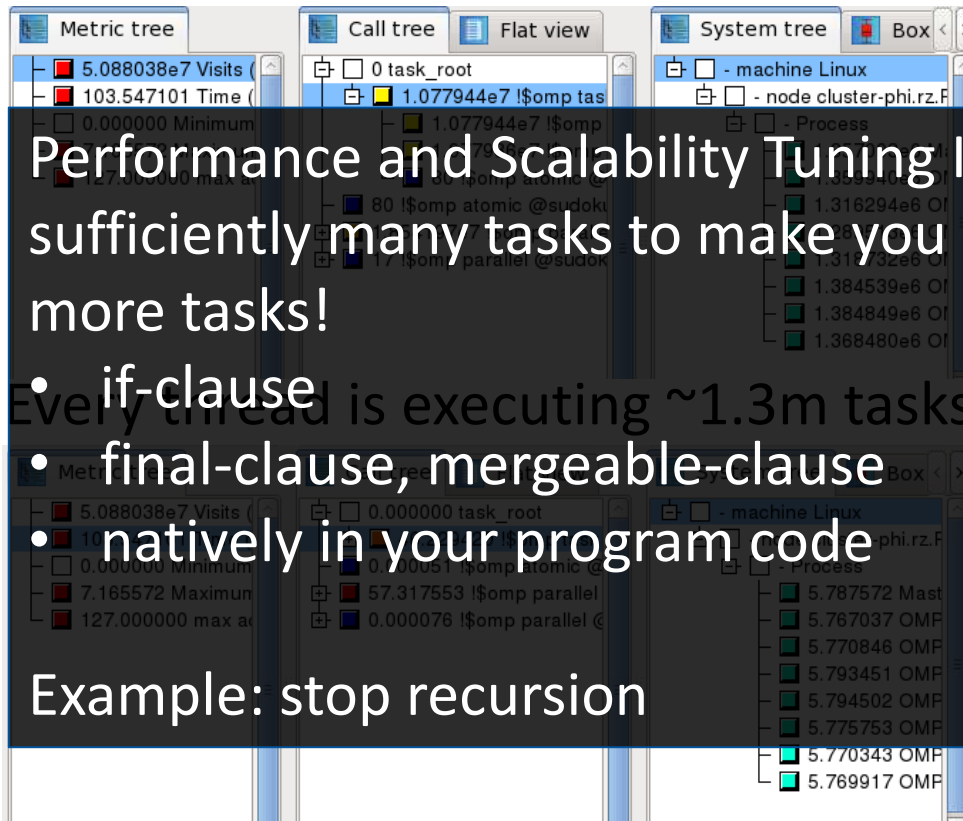
=> average duration of a task is ~4.4 μ s

Tracing gives more details:



Tasks get much smaller down the call-stack.

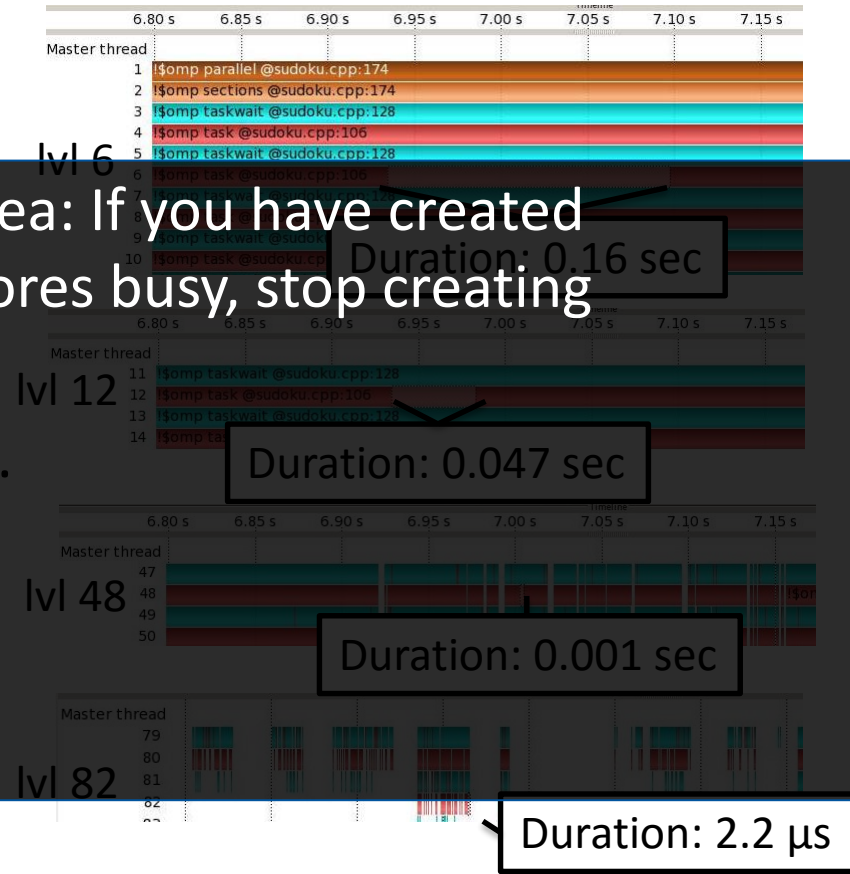
Event-based profiling gives a good overview :



... in ~5.7 seconds.


=> average duration of a task is ~4.4 μ s

Tracing gives more details:




Tasks get much smaller down the call-stack.

Sudoku on 2x Intel Xeon E5-2650 @2.0 GHz



- **Combines every element in a collection into a single result**

- employs an associative combiner function
- consequently, different ordering are possible




- **Combines every element in a collection into a single result**

- employs an associative combiner function

- consequently, different ordering are possible

- **Natural parallelization approach is to perform a tree reduction, delivering a speedup of $n / \lg n$:**



■ Reduction allows for logarithmic runtime implementation

Intel® Cilk™ Plus


```
cilk::reducer<sum_monoid> sum;  
cilk_for (int i = 1; i < 100; i++)  
    *sum += i;  
std::cout << "sum 1 .. 99 is " << *sum << std::endl;
```

OpenMP*

```
double sum = 0.0;  
#pragma omp parallel for reduction(+:sum)  
for (int i = 1; i < 100; i++)  
    sum += i;
```

- Produce all partial reductions of an input index set
- Function f needs to be associative

```
template<typename T, typename C>
void incl_scan(size_t n,
               const T in[],
               T out[],
               C f,
               T initial)
{
    for (size_t i=0; i<n; ++i) {
        initial = f(initial, in[i]);
        out[i] = initial;
    }
}
```



Simple example:

```
for (int i=1; i<N; i++)  
{  
    a[i] = f(a[i-1], b[i]);  
}
```

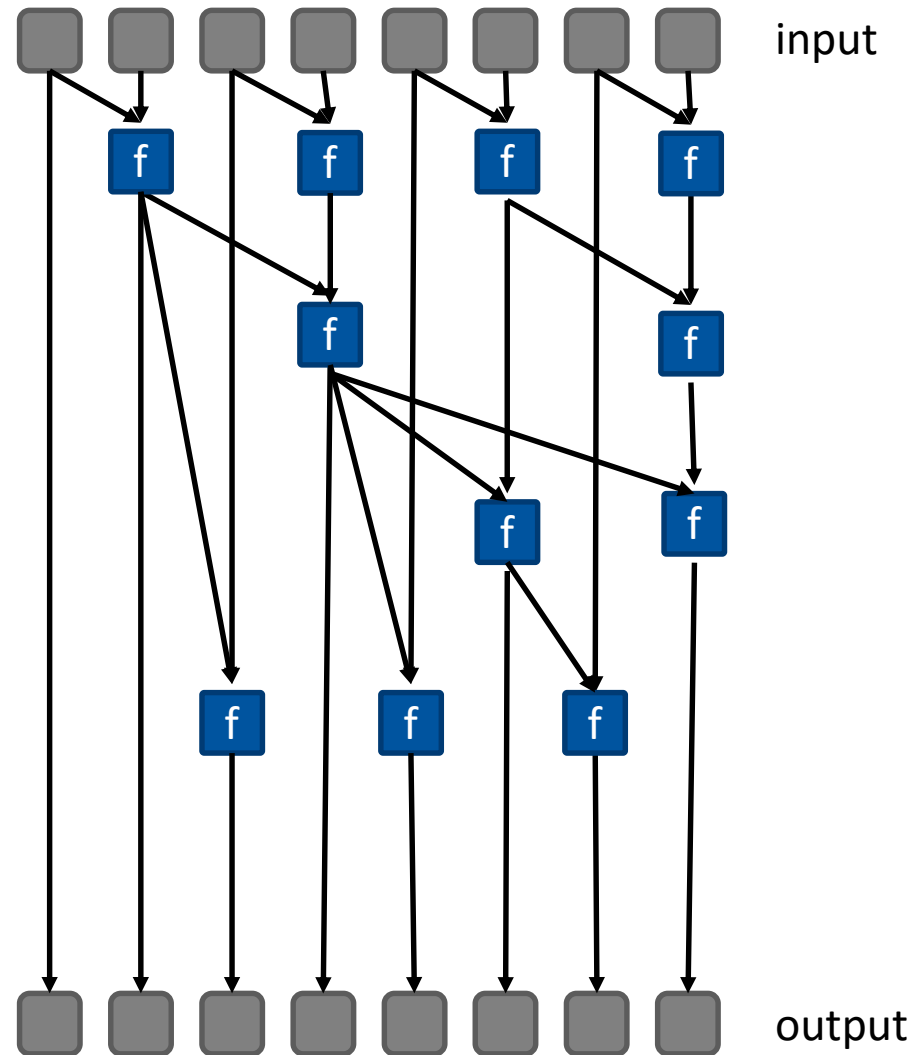
Associativity:

$$F(a, F(b, c)) = F(F(a, b), c)$$

Tree-wandering:

$$O(N * \log N)$$

The Scan Pattern: Implementation 1




The Scan Pattern: Implementation 2

Upward sweep:

- Get values L and R from left and right child
- Save L in local variable Mine
- Compute $Tmp = L + R$ and pass to parent

Downward sweep:

- Get value Tmp from parent
- Send Tmp to left child
- Send $Tmp + Mine$ to right child



- In *map*, all loop iterations are independent. In *recurrence*, there are some dependencies:


- 1D recurrence: computation in the element is associative

- has already been discussed with the *scan* pattern

- n-D recurrence: nested loop body

```
int i,j;  
for (i=1; i<N; ++i)  
  for (j=1; j<N; ++j)  
    b[i][j] = f(b[i-1][j],  
                b[i][j-1], a[i][j])
```

- Leslie Lamport's theorem: hyperplane separation can be found if the dependencies are constant offsets




Loop iteration space to hyperplane separation

```
int i,j;  
for (i=1; i<N; ++i)  
  for (j=1; j<N; ++j)  
    b[i][j] = f(b[i-1][j],  
                b[i][j-1], a[i][j])
```



OpenMP*

```
int i,j;  
#pragma omp parallel private(i,j)  
for (i=1; i<N; ++i)  
#pragma omp for  
  for (j=i+1; j<i+N; ++j)  
    b[i][j-i] = f(b[i-1][j-i],  
                  b[i][j-i-1], a[i][j])
```



■ Schedule

KW	Date	Type		Date	Type	
03	16.01.2017	V	Parallel Algorithms I	17.01.2017	V	Parallel Algorithms II
04	23.01.2017	V	GPGPU	24.01.2017	V	GPGPU/ Manycore Architectures
05	30.01.2017	Ü	Hybrid Programming/ Accelerators	31.01.2017	V	Energy Efficiency I
06	06.02.2017	V	Energy Efficiency II/ Summary/ Outlook	07.02.2017	F	Question Time

KW: week, V: lecture, Ü: exercise, F: Fragenstunde

■ Exams (February 13th and April 3rd 12:15 – 15:45)

→ 120 minutes, no additional materials allowed

Online Registration for Seminars and Practical Courses (Praktika) in Summer Term 2017

Who?

- Students of:
- Master Courses
 - Bachelor Informatik (~~Pro~~Seminar!)


Where?

www.graphics.rwth-aachen.de/apse

When?

13.01.2017 – 29.01.2017

Master Thesis: Representation of a User Jobmix as a Parameterization of Standard Benchmarks



■ Goal: Representation of the investigated programs by means of parametrization of standard benchmarks

- Characterization of the job mix by a set of metrics
- Comparison to common benchmarks with known characteristics
- Fitting a weighted mix of the benchmarks to the programs of the job mix


Loop iteration space to hyperplane separation

```
int i,j;  
for (i=1; i<N; ++i)  
  for (j=1; j<N; ++j)  
    b[i][j] = f(b[i-1][j],  
                b[i][j-1], a[i][j])
```



OpenMP*

```
int i,j;  
#pragma omp parallel private(i,j)  
for (i=1; i<N; ++i)  
#pragma omp for  
  for (j=i+1; j<i+N; ++j)  
    b[i][j-i] = f(b[i-1][j-i],  
                  b[i][j-i-1], a[i][j])
```




Case Study:

CG

■ Sparse Linear Algebra


- Sparse Linear Equation Systems occur in many scientific disciplines.
- Sparse matrix-vector multiplications (SpMxV) are the dominant part in many iterative solvers (like the CG) for such systems.
- number of non-zeros $\ll n \cdot n$



Beijing Botanical Garden

Oben Rechts: Original Gebäude
Unten Rechts: Modell
Unten Links: Matrix

(Quelle: Beijing Botanical Garden and University of Florida, Sparse Matrix Collection)




■ $A = \begin{pmatrix} 1 & 0 & 0 & 0 \\ 2 & 2 & 0 & 0 \\ 0 & 0 & 3 & 0 \\ 4 & 0 & 4 & 4 \end{pmatrix}$

```
for (i = 0; i < A.num_rows; i++){
    sum = 0.0;
    for (nz=A.row[i]; nz<A.row[i+1]; ++nz){
        sum+= A.value[nz]*x[A.index[nz]];
    }
    y[i] = sum;
}
```


$$\vec{y} = A * \vec{x}$$

- Format: compressed row storage
- store all values and columns in arrays (length nnz)
- store beginning of a new row in a third array (length n+1)

value:	1	2	2	3	4	4	4
index:	0	0	1	2	0	2	3
row:	0	1	3	4	7		



Hotspot analysis of the serial code:




Hotspots are:

1. matrix-vector multiplication
2. scaled vector additions
3. dot product

Tuning:



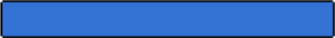
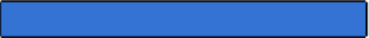
- parallelize all hotspots with a parallel for construct
- use a reduction for the dot-product
- activate thread binding



Hotspot analysis of naive parallel version:


Event Name
MEM_UNCORE_RETIRED.LOCAL_DRAM_AND_REMOTE_CACHE_HIT
MEM_UNCORE_RETIRED.REMOTE_DRAM

A lot of remote accesses occur in nearly all places.

	MEM_UNCORE_RETIRED.LOCAL_...	MEM_UNCORE_RETIRED.REMOTE...
void matvec(const int n, const int		
int i,j;		
#pragma omp parallel for private(j)	20,000	0
for(i=0; i<n; i++){	0	0
y[i]=0;	0	0
for(j=ptr[i]; j<ptr[i+1]; j	6,740,000 	3,720,000 
y[i]+=value[j]*x[index[17,580,000 	6,680,000 
}		
}		

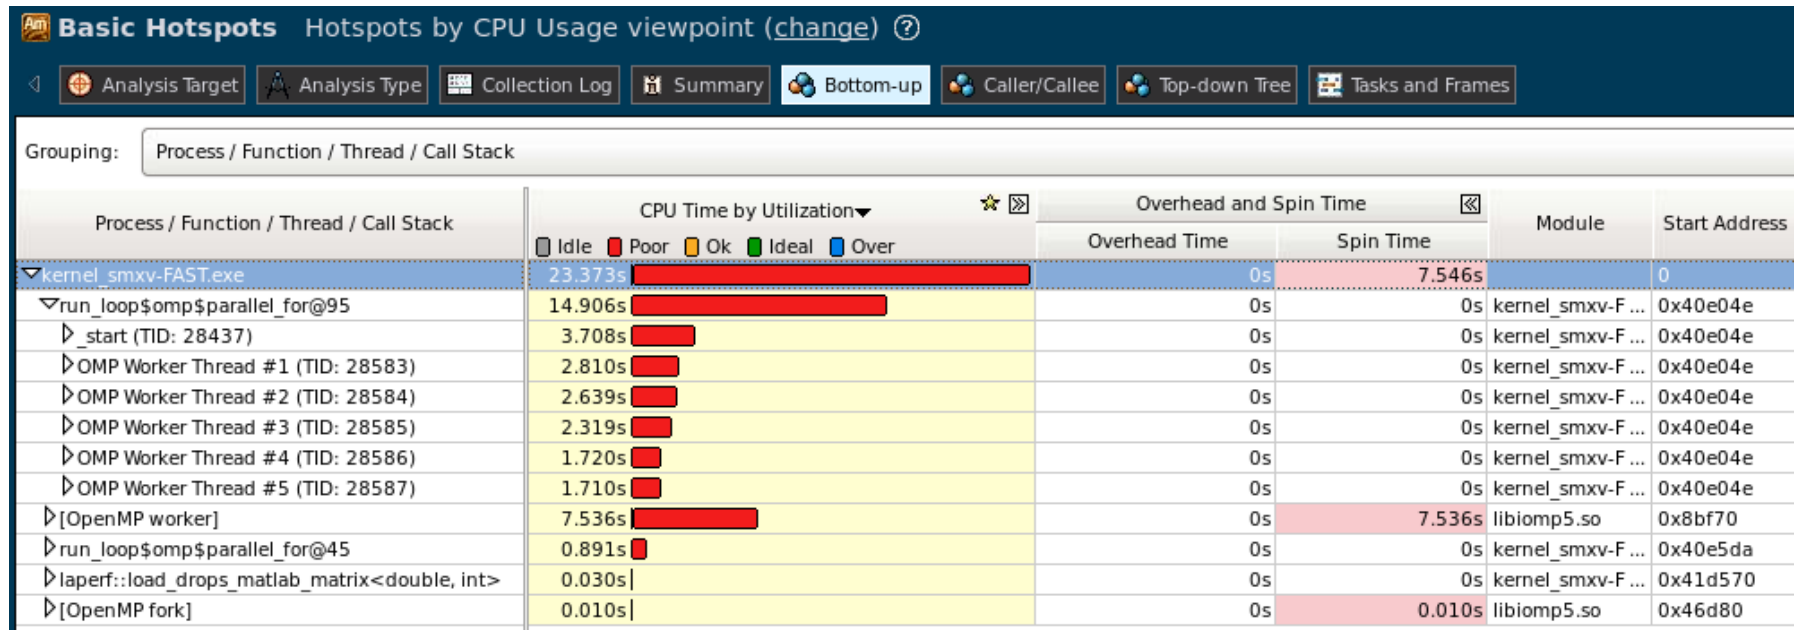
Tuning:

- Initialize the data in parallel
- Add parallel for constructs to all initialization loops




- Scalability improved a lot by this tuning on the large machine.

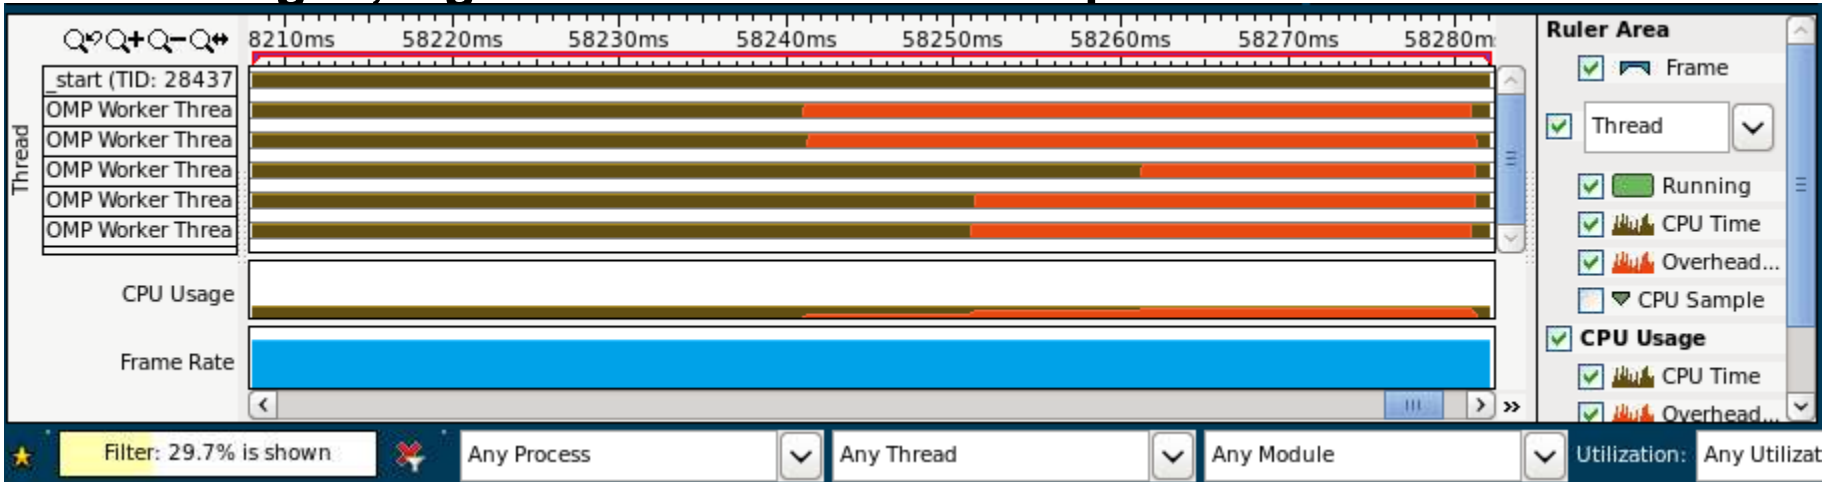
- Grouping execution time of parallel regions by threads helps to detect load imbalance.
- Significant portions of Spin Time also indicate load balance problems.
- Different loop schedules might help to avoid these problems.



- The Timeline can help to investigate the problem further.



- Zooming in, e.g. to one iteration is also possible.



■ Analyzing load imbalance in the concurrency view:

So.. Line	Source	CPU Time: Total by... Idle Poor Ok	Ove... and...
49	void matvec(const int n, const int nnz,		
50	int i,j;		
51	#pragma omp parallel for private(j)	22.462s	10.612s
52	for(i=0; i<n; i++){	0.050s	0s
53	y[i]=0;	0.060s	0s
54	for(j=ptr[i]; j<ptr[i+1]; j++){	1.741s	0s
55	y[i]+=value[j]*x[index[j]];	9.998s	0s

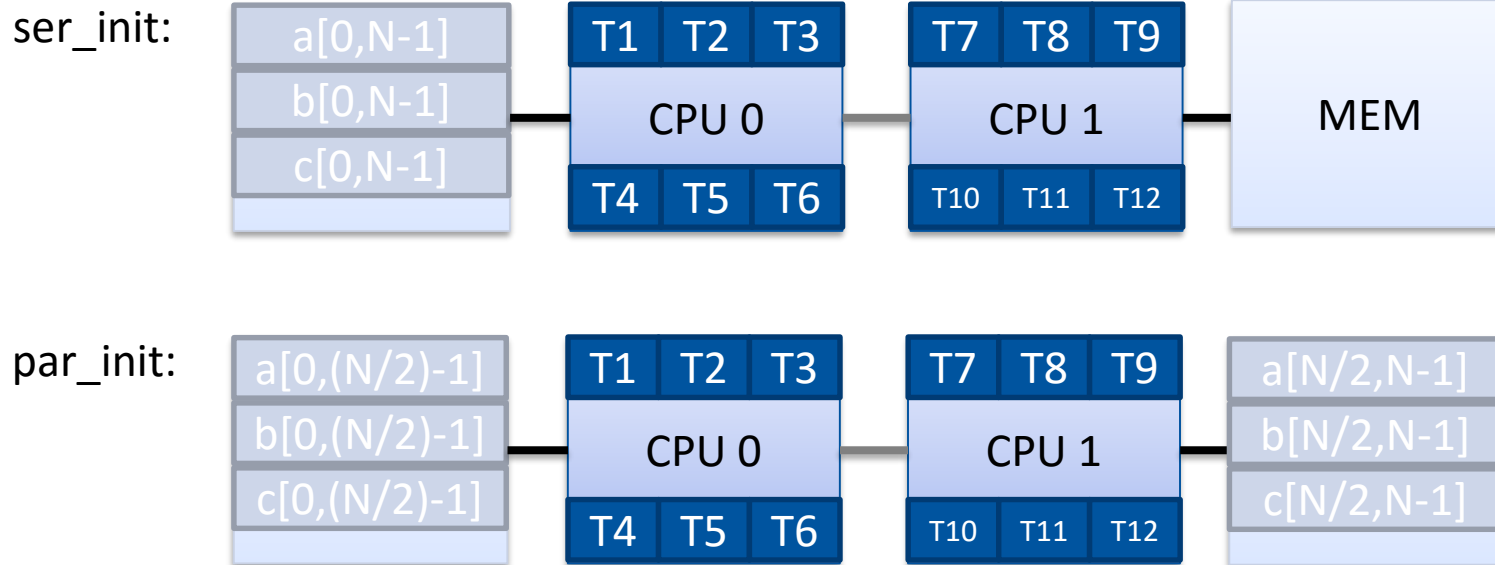
■ 10 seconds out of ~35 seconds are overhead time

■ other parallel regions which are called the same amount of time only produce 1 second of overhead

■ Stream example ($\vec{a} = \vec{b} + s * \vec{c}$) with and without parallel initialization.


→ 2 socket system with Xeon X5675 processors, 12 OpenMP threads

	copy	scale	add	triad
ser_init	18.8 GB/s	18.5 GB/s	18.1 GB/s	18.2 GB/s
par_init	41.3 GB/s	39.3 GB/s	40.3 GB/s	40.4 GB/s



■ Tuning:

→ pre-calculate a schedule for the matrix-vector multiplication, so that the non-zeros are distributed evenly instead of the rows




Patterns: Data Management

- A *pipeline* is a linear sequence of stages.


- **Parallelization approaches:**

- run different stages in parallel
- run multiple copies of stateless stages in parallel (reorder output)


- **Possibly need to manage buffering between stages.**




- Break an index set into sub-index sets
- *Geometric partitioning* is a special case for non-overlapping sub-index sets
- Does not reorganize data, only provides a different view on it
- Requirements:
 - Computation independent for elements
 - Regular data structure
- Irregularly structured data (e.g. graph) can also be partitioned



Original photo:



Blurred photo:



Original photo:



Blurred photo:




```
void blur(const picture &input, picture &output) {  
    // traverse the picture  
    #pragma omp parallel for shared(input) shared(output)  
    for (size_t y = 0; y < input.num_rows(); ++y) {  
        for (size_t x = 0; x < input.num_cols(); ++x) {  
            float pixel = 0.0;  
            // apply the stencil  
            for (size_t sy = STENCIL_WIDTH; sy < STENCIL_WIDTH; ++sy) {  
                for (size_t sx = STENCIL_WIDTH; sx < STENCIL_WIDTH; ++sx) {  
                    // wrap around picture to avoid bogus accesses  
                    size_t py = wrap_around(input.num_rows(), sy - y);  
                    size_t px = wrap_around(input.num_cols(), sx - x);  
                    pixel += gauss_coeffs[sy][sx] * input.get(px,py);  
                }  
            }  
            output.set(x,y,pixel);  
        }  
    }  
}
```

OpenMP*


- **To increase efficiency, the Map pattern can be fused with the Geometric Decomposition pattern**

- Useful for traditional OS-level threads
- Potentially having kazillions of threads executing only one instance of a function might be too fine-grained




- **Have multiple partitions to work on in parallel, invoke function sequentially within each partition**

Original photo:



Blurred photo:




```
void blur(const picture &input, picture &output) {  
    // traverse over the picture  
    #pragma omp parallel for shared(input) shared(output)  
    for (size_t y = 0; y < input.num_rows(); ++y) {  
        for (size_t x = 0; x < input.num_cols(); ++x) {  
            float pixel;  
            // "map" the gauss stencil for current (x,y) position  
            pixel = gauss_stencil(input.num_rows(), input.num_cols(), y, x);  
            output.set(x,y,pixel);  
        }  
    }  
}
```

OpenMP*

Case Study:

Karazuba

			$x^2 +$	$2x +$	3	b
			$x^2 +$	$4x +$	5	a
			<hr/>			
			$5x^2 +$	$10x +$	15	
		$4x^3 +$	$8x^2 +$	$12x$		
	$x^4 +$	$2x^3 +$	$3x^2$			
	<hr/>					
	$x^4 +$	$6x^3 +$	$16x^2 +$	$22x +$	15	c
Storage Scheme:			$b[2]$	$b[1]$	$b[0]$	b
			$a[2]$	$a[1]$	$a[0]$	a
			<hr/>			
						
	$c[4]$	$c[3]$	$c[2]$	$c[1]$	$c[0]$	c

■ $O(n^2)$ work:

vector initialization

```
void simple_mul( T c[], const T a[], const T b[], size_t n ) {  
    → c[0:2*n-1] = 0;  
    for (size_t i=0; i<n; ++i)  
        c[i:n] += a[i]*b[0:n];  
}
```

vector addition

■ Can we do any better?

→ Yes. There are several methods, but Karatsuba is the one with the nicest name!

■ Suppose polynomials a and b have degree n

→ let $K = x^{\lfloor n/2 \rfloor}$

$$a = a_1K + a_0$$

$$b = b_1K + b_0$$

Partition coefficients.

■ Compute:

$$t_0 = a_0 \cdot b_0$$

$$t_1 = (a_0 + a_1) \cdot (b_0 + b_1)$$

$$t_2 = a_1 \cdot b_1$$

3 half-sized multiplications.
Do these recursively.

■ Then

$$a \cdot b \equiv t_2K^2 + (t_1 - t_0 - t_2)K + t_0$$

Sum products, shifted by
multiples of K .

Save one (half)
multiplication.

```
void karatsuba( T c[], const T a[], const T b[], size_t n ) {
```

```
    size_t m = n/2;
```

```
    cilk_spawn karatsuba( c, a, b, m );
```

```
    cilk_spawn karatsuba( c+2*m, a+m, b+m, n-m );
```

```
    temp_space<T> s(4*(n-m));
```

```
    T *a_=s.data(), *b_=a_+(n-m), *t=b_+(n-m);
```

```
    a_[0:m] = a[0:m]+a[m:m];
```

```
    b_[0:m] = b[0:m]+b[m:m];
```

```
    karatsuba( t, a_, b_, n-m );
```

```
    cilk_sync;
```

```
    t[0:2*m-1] -= c[0:2*m-1] + c[2*m:2*m-1];
```

```
    c[2*m-1] = 0;
```

```
    c[m:2*m-1] += t[0:2*m-1];
```

```
}
```

$// t_0 = a_0 \times b_0$

$// t_2 = a_1 \times b_1$

$// a_+ = (a_0 + a_1)$

$// b_+ = (b_0 + b_1)$

$// t_1 = (a_0 + a_1) \times (b_0 + b_1)$

$// t = t_1 - t_0 - t_2$

$// c = t_2 K^2 + (t_1 - t_0 - t_2) K + t_0$

```
void karatsuba( T c[], const T a[], const T b[], size_t n ) {  
    if( n<=CutOff ) {  
        simple_mul( c, a, b, n );  
    } else {  
        size_t m = n/2;  
        cilk_spawn karatsuba( c, a, b, m );  
        cilk_spawn karatsuba( c+2*m, a+m, b+m, n-m );  
        temp_space<T> s(4*(n-m));  
        T *a_=s.data(), *b_=a_+(n-m), *t=b_+(n-m);  
        a_[0:m] = a[0:m]+a[m:m];  
        b_[0:m] = b[0:m]+b[m:m];  
        karatsuba( t, a_, b_, n-m );  
        cilk_sync;  
        t[0:2*m-1] -= c[0:2*m-1] + c[2*m:2*m-1];  
        c[2*m-1] = 0;  
        c[m:2*m-1] += t[0:2*m-1];  
    }  
}
```

$// t_0 = a_0 \times b_0$

$// t_2 = a_1 \times b_1$

$// a_- = (a_0 + a_1)$

$// b_- = (b_0 + b_1)$

$// t_1 = (a_0 + a_1) \times (b_0 + b_1)$

$// t = t_1 - t_0 - t_2$

$// c = t_2 K^2 + (t_1 - t_0 - t_2) K + t_0$


Case Study:

Mandelbrot set

w/ Master-Worker



- Set of complex numbers c for which

$z_0 = 0; z_{n+1} = z_n^2 + c \quad z, c \in \mathbb{C}$
does not diverge.

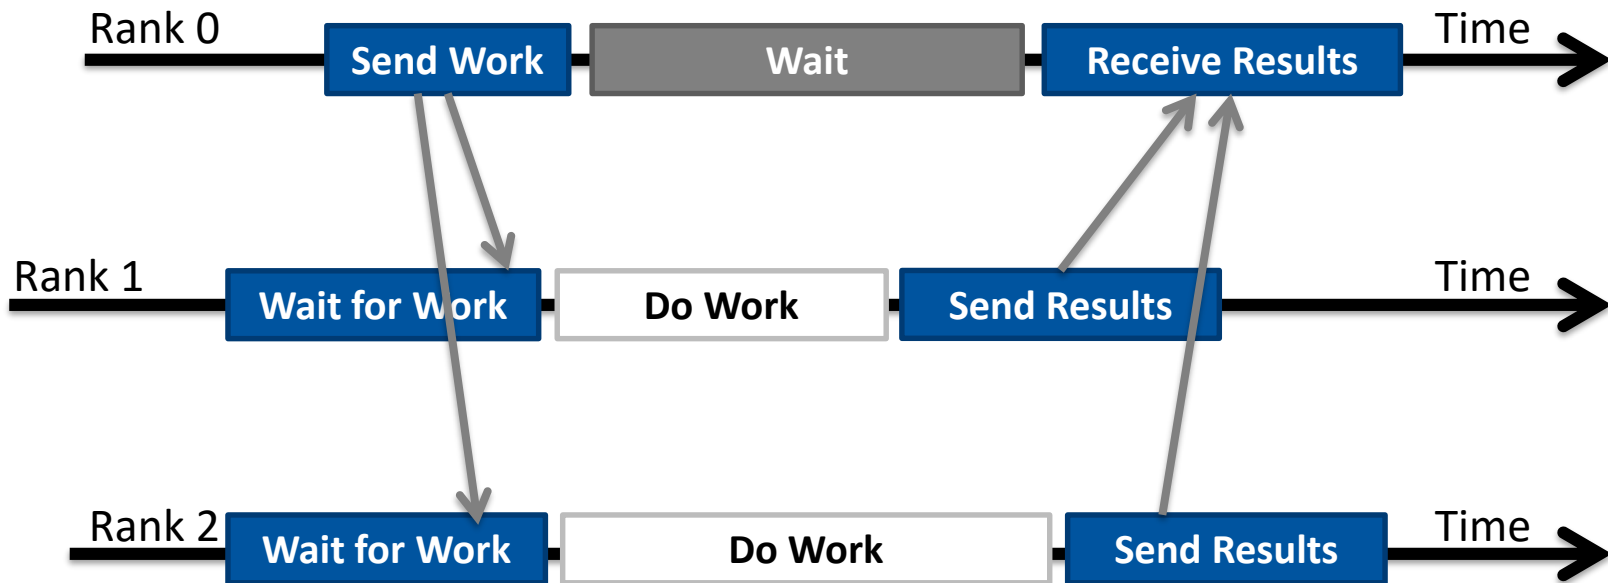


```
// For each image pixel (x, y):  
// (x, y) - scaled pixel cords  
c = x + i*y  
z = 0  
iteration = 0  
maxIters = 1000  
while (|z|^2 <= 2^2  
        && iteration < maxIters)  
{  
    z = z^2 + c  
    iteration = iteration + 1  
}  
if (iteration == maxIters)  
    color = black  
else  
    color = iteration  
plot(x, y, color)
```

Computation complexity mapped to ranks




- One process (master) manages the work.
 - Work is split into many relatively small work items.
- Many other processes (workers) compute over the work items:



- Above steps are repeated until all work items are processed.
- Sometimes called “bag of jobs” pattern.



■ The algorithm:

```
if (rank == 0) { // Master
    splitDomain;
    sendWorkItems;
    receivePartialResults;
    assembleResult;
    output;
} else { // Worker
    receiveWorkItems;
    processWorkItems;
    sendPartialResults;
}
```




Case Study: Mandelbrot Set Work Balancing

Computational complexity mapped to ranks



Case Study: Mandelbrot Set Results



Quiz

- **We wanted to make you think parallel: were these two lectures useful?**
 - a) Yes
 - b) No
 - c) Yes, at least I have learned how to build a house
 - d) -

- **Throughout this course we offered you to get experienced in using the HPC Cluster at RWTH Aachen University. Did you ever login?**
 - a) Yes, several times
 - b) Once to see how the environment looks like
 - c) Never
 - d) What is an HPC cluster?

- **In case you are interested in learning more about HPC: would you attend a dedicated lecture on Parallel and Data-centric Programming Models?**
 - a) Yes
 - b) Maybe if it fits my schedule
 - c) No interest whatsoever
 - d) –

Summary

- **Problem:** Amdahl's Law notes that scaling will be limited by the serial fraction of your program.
- ***Solution: go parallel “as far as you can”, exploit all dimensions of parallelism.***

- **Problem:** Locking, access to data (memory and communication), and overhead will strangle scaling.
- ***Solution: use programming approaches with good data locality and low overhead, and avoid locks.***

- **Problem:** Parallelism introduces new debugging challenges: deadlocks and race conditions.
- ***Solution: use structured programming strategies to avoid these by design, improving maintainability.***

Question Time...