

1. Why supercomputers?
2. Modern processors
3. Basic optimization techniques for serial code
4. Data access optimization
5. Parallel computers
6. Energy efficiency
7. Parallelization and optimization strategies
8. Parallel algorithms
9. Distributed-memory programming with MPI
10. Shared-memory programming with OpenMP
11. Hybrid programming (MPI + OpenMP)

## 12. Heterogeneous architectures (GPUs, Xeon Phis)

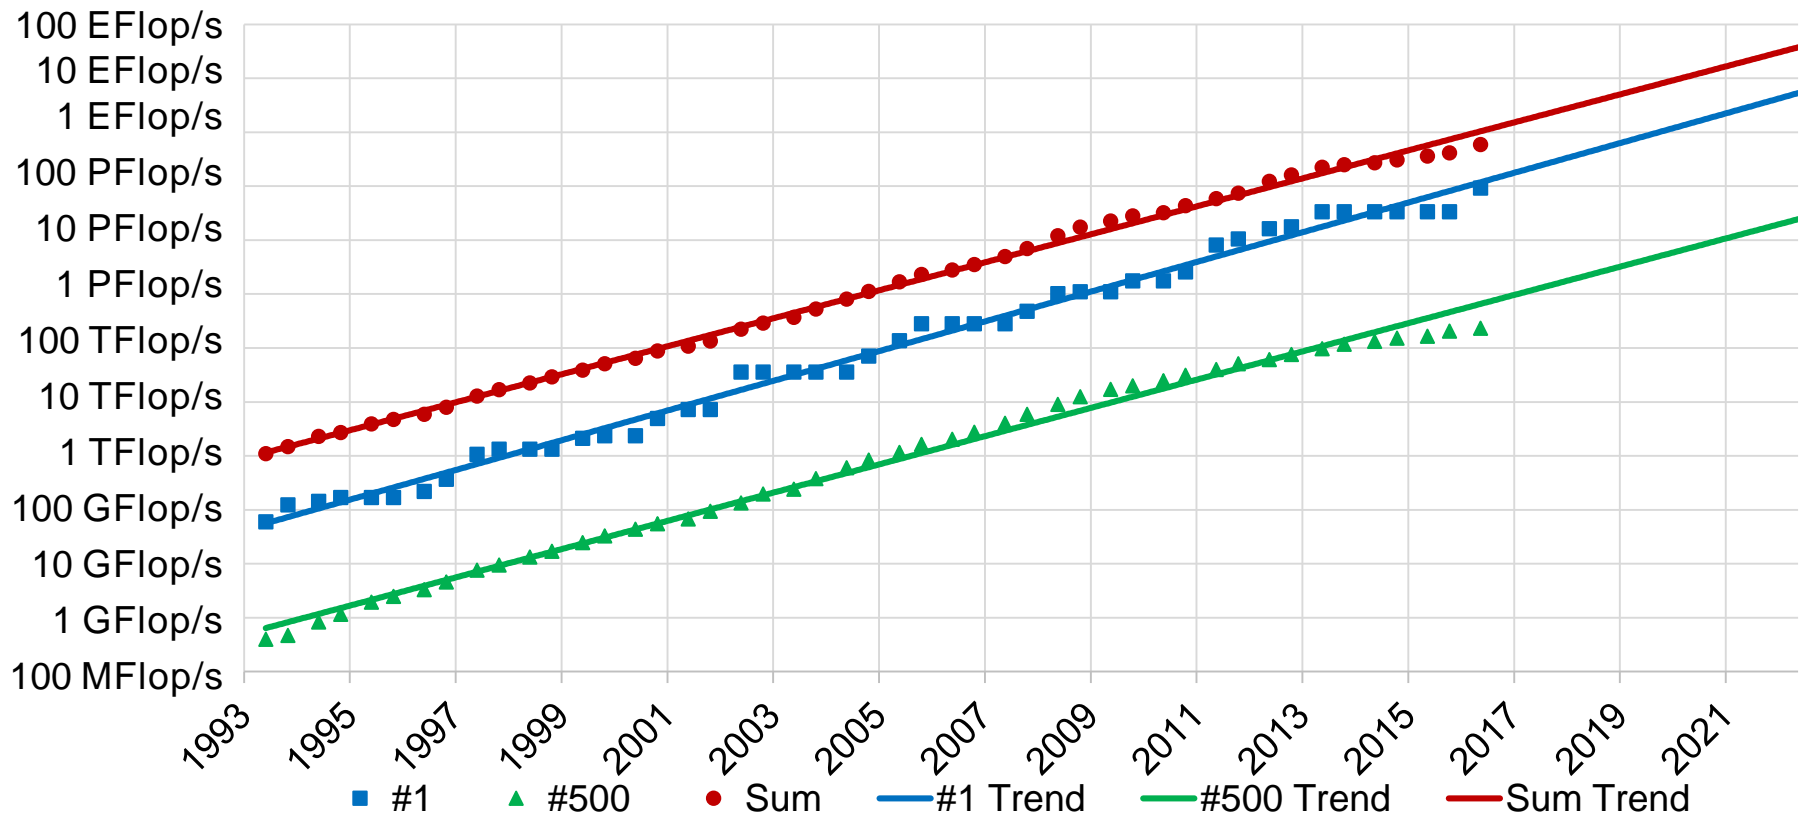
- Motivation
- RWTH Environment
- GPGPU – Basics
  - Overview
  - GPU Architecture

- Programming Model
- Execution Model
- Memory Model
- Misc on OpenACC
- Summary
- GPGPU – Advanced
  - Kernel optimizations
  - GPU-CPU interaction
- Intel Xeon Phi
  - Motivation
  - KNL Architecture
  - Vectorization
- Links

## ■ Why to care about accelerators?

→ Towards exa-flop computing (performance gain, but power constraints)

→ Accelerators provide good performance per watt ratio (first step)



Today, NVIDIA GPUs are the focus.

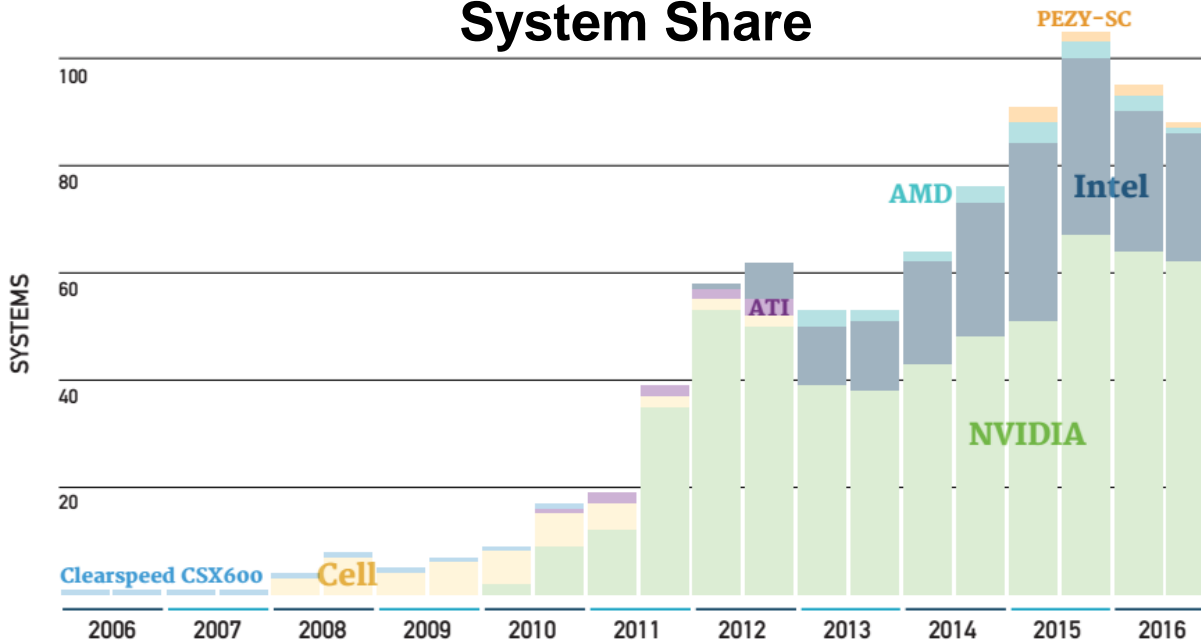
## ■ Accelerators/ co-processors

- GPGPUs (e.g. NVIDIA, AMD)
- Intel Many Integrated Core (MIC) Arch. (Intel Xeon Phi)
- FPGAs (e.g. Convey), ...

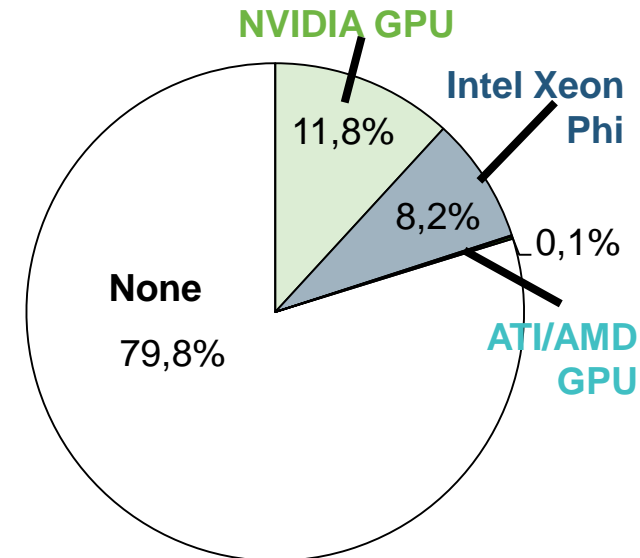
## ■ Heterogeneous Arch.

- Combination of commodity processors & accelerators

### System Share



### Performance Share



# Motivation – Green500 (11/2016)



Rank	MFLOPS/ W	Site*	Proc/ Accelerator	Total Power (kW)
1	9462.1	NVIDIA Corporation	NVIDIA Tesla P100	349.5
2	7453.5	Swiss National Supercomputing Centre (CSCS)	NVIDIA Tesla P100	1312
3	6673.8	Advanced Center for Computing and Communication, RIKEN	PEZY-SCnp	150.0
4	6051.3	National Supercomputing Center in Wuxi	[heterogeneous]	15371
5	5806.3	Fujitsu Technology Solutions GmbH	Intel Xeon Phi 7210	77
6	4985.7	Joint Center for Advanced High Performance Computing	Intel Xeon Phi 7250	2718.7
7	4688.0	DOE/SC/Argonne National Laboratory	Intel Xeon Phi 7230	1087
8	4112.1	Stanford Research Computing Center	Nvidia K80	190
9	4086.8	Academic Center for Computing and Media Studies (ACCMS), Kyoto University	Intel Xeon Phi 7250	748.1
10	3836.6	Thomas Jefferson National Accelerator Facility	Intel Xeon Phi 7230	111
33	2903.6	Scientific research institution	–	162

**Comparison to server with 2x Intel Sandy Bridge@ 2GHz**

- HPL: ~260 W
- Peak Performance: 256 GFLOPS

→ 985 MFLOPS/W

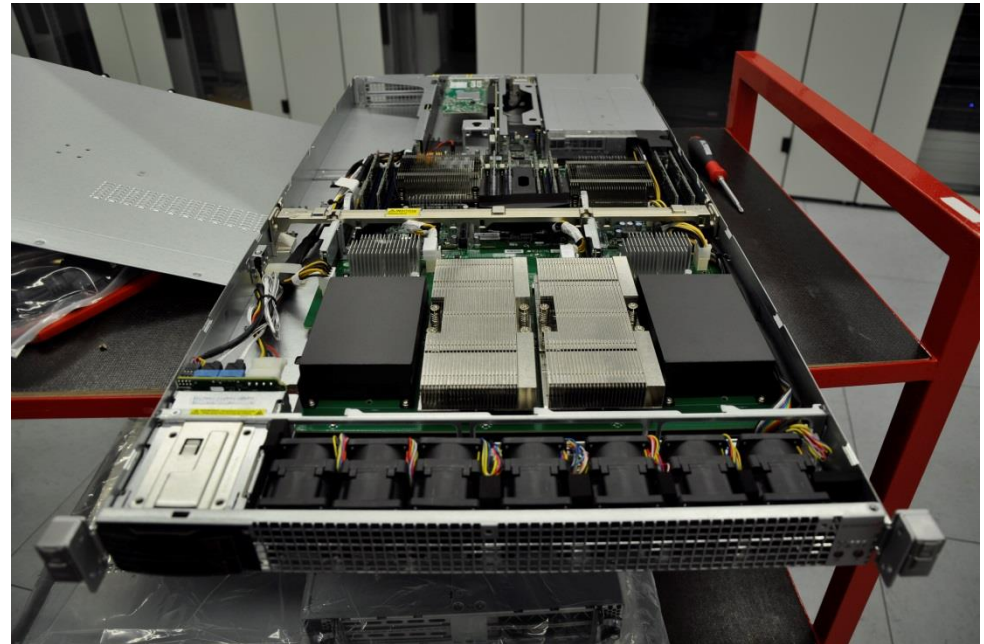
**performance per watt**

## ■ GPU cluster

- 28 nodes with each  
2 NVIDIA Quadro 6000 GPUs (Fermi)  
→ VR + HPC
- 2 nodes with each  
2 NVIDIA K20X GPUs (Kepler)
- 10 nodes with each  
2 NVIDIA P100 SXM2 (Pascal)  
with 16GB



aixCAVE, VR, RWTH Aachen, since June 2012



1. Why supercomputers?
2. Modern processors
3. Basic optimization techniques for serial code
4. Data access optimization
5. Parallel computers
6. Energy efficiency
7. Parallelization and optimization strategies
8. Parallel algorithms
9. Distributed-memory programming with MPI
10. Shared-memory programming with OpenMP
11. Hybrid programming (MPI + OpenMP)

## 12. Heterogeneous architectures (GPUs, Xeon Phis)

- Motivation
- RWTH Environment
- **GPGPU – Basics**
  - **Overview**
  - GPU Architecture

- Programming Model
- Execution Model
- Memory Model
- Misc on OpenACC
- Summary
- GPGPU – Advanced
  - Kernel optimizations
  - GPU-CPU interaction
- Intel Xeon Phi
  - Motivation
  - KNL Architecture
  - Vectorization
- Links

## ■ GPGPUs = **G**eneral **P**urpose **G**raphics **P**rocessing **U**nits

## ■ History – a very brief overview

→ '80s - '90s: Development is mainly driven by games

Fixed-function 3D graphics pipeline

Graphics APIs like OpenGL, DirectX popular

→ Since 2001: Programmable pixel and vertex shader in graphics pipeline

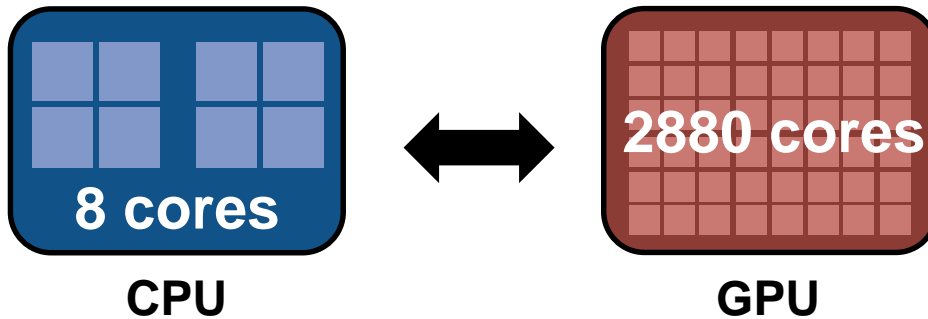
(adjustments in OpenGL, DirectX)

Researchers take notice of performance growth of GPUs: Tasks must be cast into native graphics operations

→ Since 2006: Vertex/pixel shader are replaced by a single processor unit

Support of programming language C, synchronization,...

→ “General purpose”



## ■ GPU-Threads

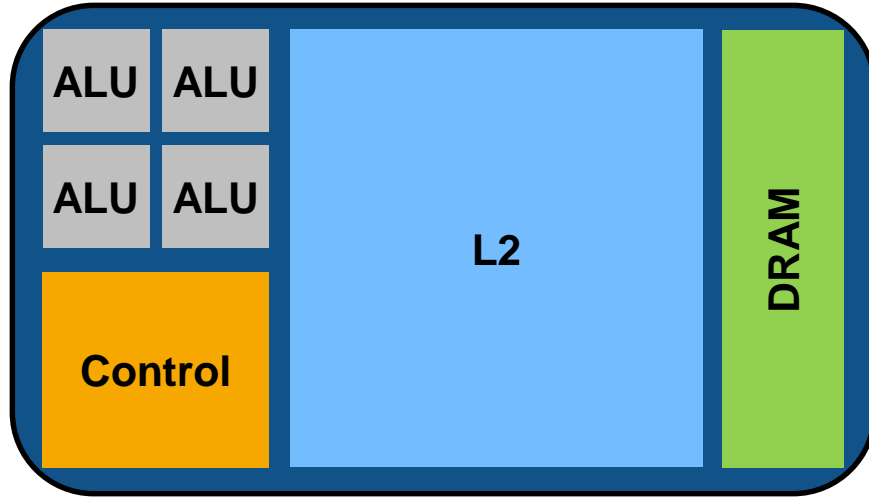
- Thousands (“few” on CPU)
- Light-weight, little creation overhead
- Fast switching

## ■ Massively Parallel Processors

## ■ Manycore Architecture



## ■ Different design



### CPU

- Optimized for **low latencies**
- Huge caches
- Control logic for out-of-order and speculative execution



### GPU

- Optimized for **data-parallel throughput**
- Architecture tolerant of memory latency
- More transistors dedicated to computation

# Why can accelerators deliver good performance watt ratio?



## 1. High (peak) performance

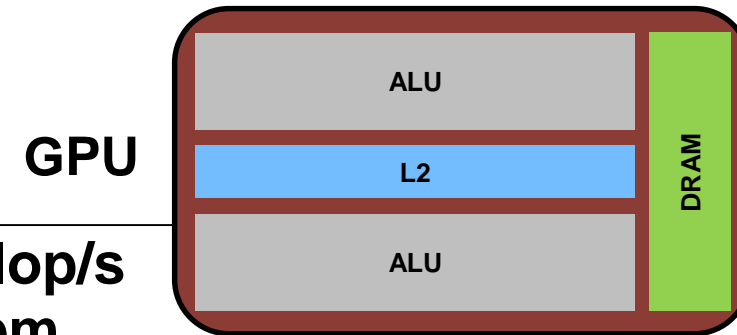
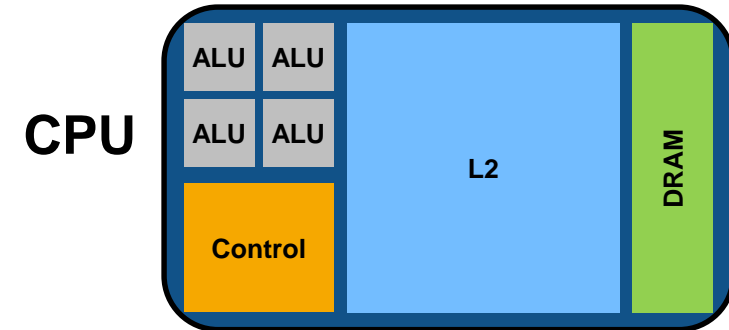
- More transistors for computation
- No control logic
- Small caches

## 2. Low power consumption

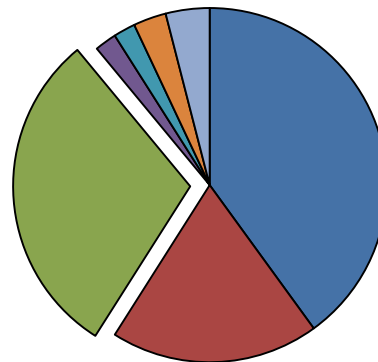
- Many low frequency cores

$$P \sim V^2 \cdot f$$

- No control logic



### Power use for 1 TFlop/s of an usual system



- Heat removal
- Power supply loss
- Control
- Disk
- Communication
- Memory
- Compute

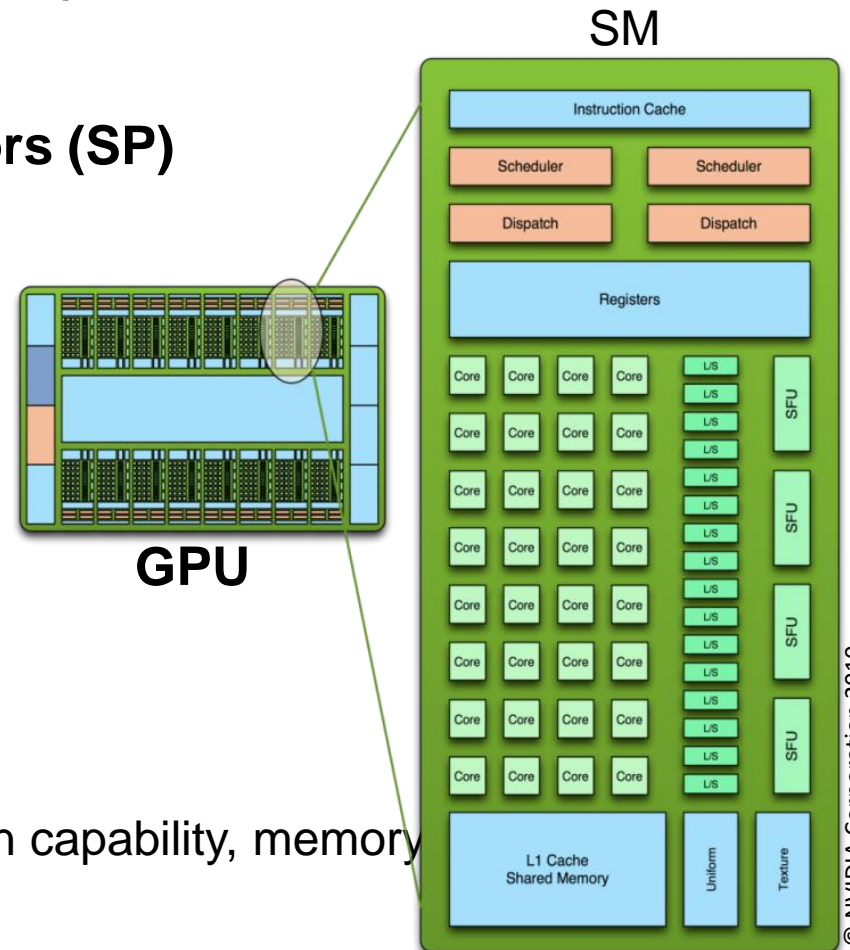
1. Why supercomputers?
2. Modern processors
3. Basic optimization techniques for serial code
4. Data access optimization
5. Parallel computers
6. Energy efficiency
7. Parallelization and optimization strategies
8. Parallel algorithms
9. Distributed-memory programming with MPI
10. Shared-memory programming with OpenMP
11. Hybrid programming (MPI + OpenMP)

## 12. Heterogeneous architectures (GPUs, Xeon Phis)

- Motivation
- RWTH Environment
- **GPGPU – Basics**
  - Overview
  - **GPU Architecture**

- Programming Model
- Execution Model
- Memory Model
- Misc on OpenACC
- Summary
- GPGPU – Advanced
  - Kernel optimizations
  - GPU-CPU interaction
- Intel Xeon Phi
  - Motivation
  - KNL Architecture
  - Vectorization
- Links

- 3 billion transistors
- 14-16 streaming **multiprocessors (SM)**
  - Each comprises 32 (SP) cores
- 448-512 **cores/ streaming processors (SP)**
  - i.a. Floating point & integer unit
- Memory hierarchy
- Peak performance
  - SP: 1.03 TFlops
  - DP: 515 GFlops
- ECC support
- Compute capability: 2.0
  - Defines features, e.g. double precision capability, memory



- 7.1 billion transistors
- 13-15 streaming multiprocessors extreme (SMX)

→ Each comprises 192 (SP) cores

- 2496-2880 cores
- Memory hierarchy

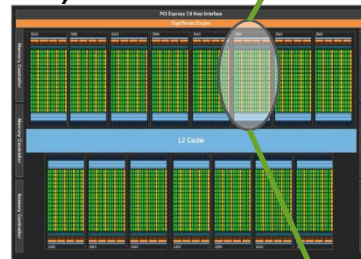
- Peak performance (K20) GPU


→ SP: 3.52 TFlops

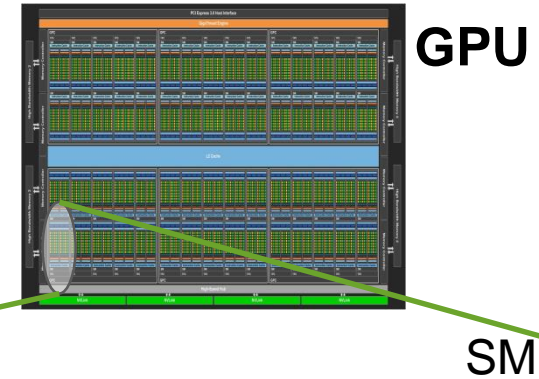
→ DP: 1.17 TFlops

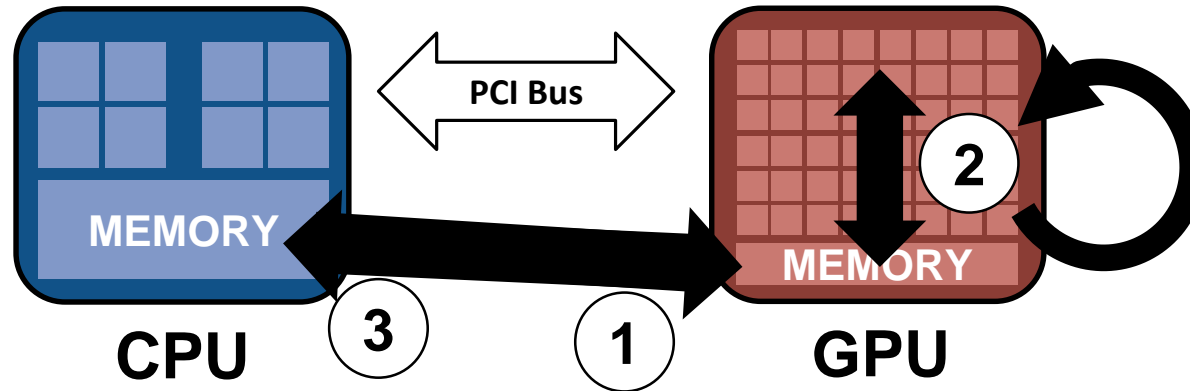
- ECC support
- Compute capability: 3.5

→ E.g. dynamic parallelism = possibility to launch dynamically new work from GPU



- **15.3 billion transistors**
  - **56-60 streaming multiprocessors (SM)**
    - Each comprises 64 (SP) cores
    - Divided into 2 processing blocks
  - **3584-3840 cores**
  - **Memory hierarchy**
  - **Peak performance (P100)**
    - SP: 9.52 Tflops
    - DP: 4.76 Tflops
  - **ECC support, NVLink**
  - **Compute capability: 6.0**
    - E.g. atomicAdd()  
on 64-bit float
- 





## ■ Weak memory model

- Host + device memory = separate entities
- No coherence between host + device
  - Data transfers needed

## ■ Host-directed execution model

- Copy input data from CPU mem. to device mem.
- Execute the device program
- Copy results from device mem. to CPU mem.

1. Why supercomputers?
2. Modern processors
3. Basic optimization techniques for serial code
4. Data access optimization
5. Parallel computers
6. Energy efficiency
7. Parallelization and optimization strategies
8. Parallel algorithms
9. Distributed-memory programming with MPI
10. Shared-memory programming with OpenMP
11. Hybrid programming (MPI + OpenMP)

## 12. Heterogeneous architectures (GPUs, Xeon Phis)

- Motivation
- RWTH Environment
- **GPGPU – Basics**
  - Overview
  - GPU Architecture

- **Programming Model**
- Execution Model
- Memory Model
- Misc on OpenACC
- Summary
- GPGPU – Advanced
  - Kernel optimizations
  - GPU-CPU interaction
- Intel Xeon Phi
  - Motivation
  - KNL Architecture
  - Vectorization
- Links



**Application**

**Libraries**

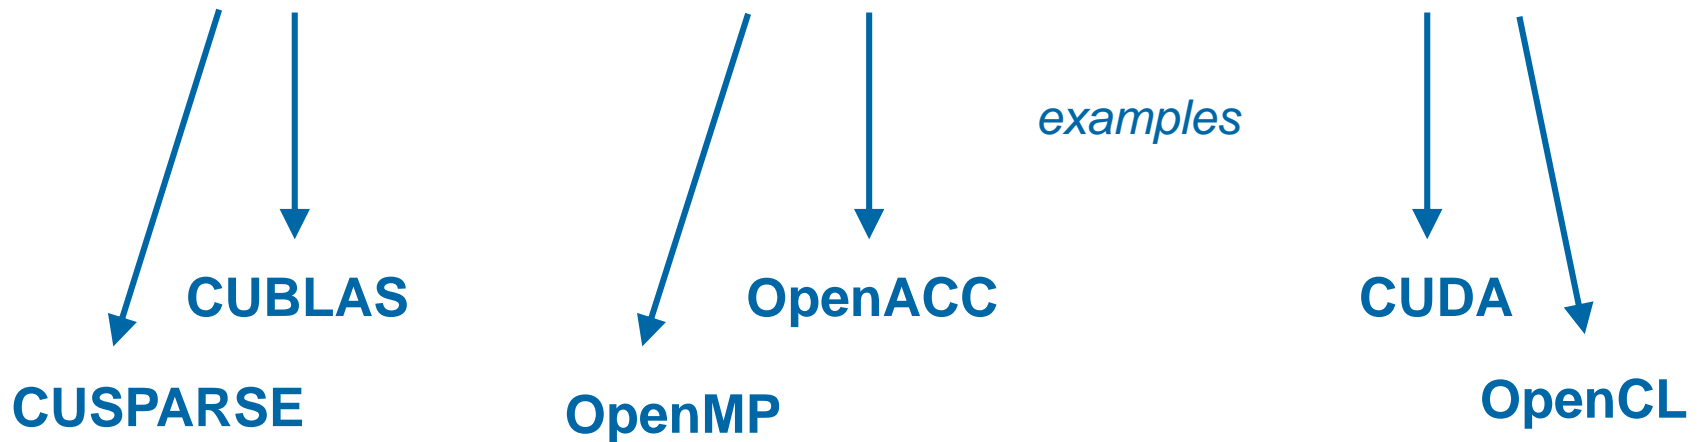
**Directives**

**Programming  
Languages**

“Drop-in” acceleration

High-level programming

Low-level programming



## ■ **CUDA (Compute Unified Device Architecture)**

→ C/C++ (NVIDIA): architecture + programming language, NVIDIA GPUs

→ Fortran (PGI): NVIDIA's CUDA for Fortran, NVIDIA GPUs

## ■ **OpenCL**

→ C (Khronos Group): open standard, portable, CPU/GPU/...

## ■ **OpenACC**

→ C/Fortran (PGI, Cray, CAPS, NVIDIA): Directive-based accelerator programming, industry standard published in Nov. 2011 (NVIDIA GPUs)

## ■ **OpenMP**

→ C/C++, Fortran: Directive-based programming for hosts and accelerators, standard, portable, published in July 2013, implementations for Xeon Phi

■ ...

## = Compute Unified Device Architecture

### ■ CUDA C/C++ from NVIDIA

**CUDA Fortran** available by PGI.

- Based on industry standard C/C++ (extensions & restrictions)
- Driver API (low level), Runtime API (higher level)

### ■ Brief timeline

- Nov'06: Introduction of CUDA, G80 GPU architecture
- Jun'07: CUDA Toolkit 1.0
- Jun'08: GT200 GPU architecture
- March'10: Fermi GPU architecture
- Nov'12: Kepler K20(X) GPU architecture
- July'16: Pascal GPU architecture

**CUDA Toolkit**  
Developer toolkit

Already installed in RWTH Cluster environment:  
`module load cuda`

## ■ Download + install CUDA Toolkit

(cf. “Links” section)

## ■ Compiling

```
module load cuda
```

*# on our cluster*

```
nvcc -arch=sm_20 saxpy.cu
```

*# see provided Makefile*

→ **nvcc**: NVIDIA’s compiler for C/C++ GPU code


→ **-arch=sm\_20**: Set compute capability 2.0

→ Sets certain architecture features, e.g. enabling double precision floating point operations

## ■ Nowadays: GPU APIs (like CUDA, OpenCL) often used

- May be difficult to program (as/but more flexibility)
- Verbose/ may complicate software design

## ■ Directive-based programming model delegates responsibility for low-level GPU programming tasks to compiler

- Data movement
  - Kernel execution
  - “Awareness” of particular GPU type
  - ...
- 
- A large right-facing curly bracket groups the four items on the left. To the right of the bracket, two items are listed, indicating that these tasks are delegated to the compiler or runtime.
- Many tasks can be done by compiler/ runtime
  - User-directed programming

## ➔ OpenACC or OpenMP 4.x device constructs

## ■ Open industry standard

→ Portability

Here, PGI's OpenACC compiler is used for examples. Details are compiler dependent.

## ■ Introduced by CAPS, Cray, NVIDIA, PGI (Nov. 2011)

## ■ Support

→ C, C++ and Fortran

→ NVIDIA GPUs, AMD GPUs, x86 Multicore, Intel MIC  
(now/near future)

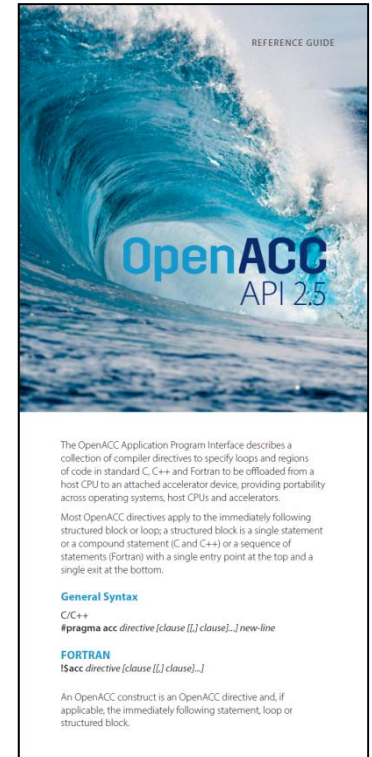
## ■ Timeline

→ Nov'11: Specification 1.0

→ Jun'13: Specification 2.0

→ Nov'14: Proposal for complex data management

→ Nov'15: Specification 2.5



Source: [www.openacc.org](http://www.openacc.org)

## ■ Syntax

### C

```
#pragma acc directive-name [clauses]
```

### Fortran

```
!$acc directive-name [clauses]
```

## ■ Iterative development process

### ➔ Compiler feedback helpful

- ➔ Whether an accelerator kernel could be generated
- ➔ Which loop schedule is used
- ➔ Where/which data is copied

```
pgcc -acc -ta=nvidia,cc20,7.0 -Minfo=accel saxpy.c
```

- **pgcc** C PGI compiler (**pgf90** for Fortran)
- **-acc** Tells compiler to recognize OpenACC directives
- **-ta=nvidia** Specifies the target architecture → here: NVIDIA GPUs
- **cc20** Optional. Specifies target compute capability 2.0
- **7.0** Optional. Uses CUDA Toolkit 7.0 for code generation
- **-Minfo=accel** Optional. Compiler feedback for accelerator code

The PGI tool **pgaccelinfo** prints the minimal needed compiler flags for the usage of OpenACC on the current hardware (see *Development Tips*).



```
void saxpyCPU(int n, float a, float *x, float *y) {  
    for (int i = 0; i < n; ++i)  
        y[i] = a*x[i] + y[i];  
}
```

SAXPY = Single-precision real Alpha X Plus Y

$$\vec{y} = \alpha \cdot \vec{x} + \vec{y}$$

```
int main(int argc, const char* argv[]) {  
    int n = 10240; float a = 2.0f;  
    float *x = (float*) malloc(n * sizeof(float));  
    float *y = (float*) malloc(n * sizeof(float));  
  
    // Initialize x, y  
    for(int i=0; i<n; ++i){  
        x[i]=i;  
        y[i]=5.0*i-1.0;  
    }  
  
    // Invoke serial SAXPY kernel  
    saxpyCPU(n, a, x, y);  
  
    free(x); free(y);  
    return 0;  
}
```

```
void saxpyOpenACC(int n, float a, float *x, float *y) {  
#pragma acc kernels loop gang vector(128)  
    for (int i = 0; i < n; ++i)  
        y[i] = a*x[i] + y[i];  
}  
  
int main(int argc, const char* argv[]) {  
    int n = 10240; float a = 2.0f;  
    float *x = (float*) malloc(n * sizeof(float));  
    float *y = (float*) malloc(n * sizeof(float));  
  
    // Initialize x, y  
    for(int i=0; i<n; ++i){  
        x[i]=i;  
        y[i]=5.0*i-1.0;  
    }  
  
    // Invoke serial SAXPY kernel  
    saxpyOpenACC(n, a, x, y);  
  
    free(x); free(y);  
    return 0;  
}
```

```
__global__ void saxpyCUDA(int n, float a,
    float *x, float *y) {
    int i = blockIdx.x * blockDim.x +
        threadIdx.x;
    if (i < n){
        y[i] = a*x[i] + y[i];
    }
}

int main(int argc, char* argv[]) {
    int n = 10240; float a = 2.0f;
    float* h_x,*h_y; // Pointer to CPU memory
    h_x = (float*) malloc(n* sizeof(float));
    h_y = (float*) malloc(n* sizeof(float));
    // Initialize h_x, h_y
    for(int i=0; i<n; ++i){
        h_x[i]=i;
        h_y[i]=5.0*i-1.0;
    }
    float *d_x,*d_y; // Pointer to GPU memory
    cudaMalloc(&d_x, n*sizeof(float));
    cudaMalloc(&d_y, n*sizeof(float));

    cudaMemcpy(d_x, h_x, n * sizeof(float),
        cudaMemcpyHostToDevice);
    cudaMemcpy(d_y, h_y, n * sizeof(float),
        cudaMemcpyHostToDevice);

    // Invoke parallel SAXPY kernel
    dim3 threadsPerBlock(128);
    dim3 blocksPerGrid(n/threadsPerBlock.x);
    saxpyCUDA<<<blocksPerGrid,
        threadsPerBlock>>>(n, 2.0, d_x, d_y);

    cudaMemcpy(h_y, d_y, n * sizeof(float),
        cudaMemcpyDeviceToHost);

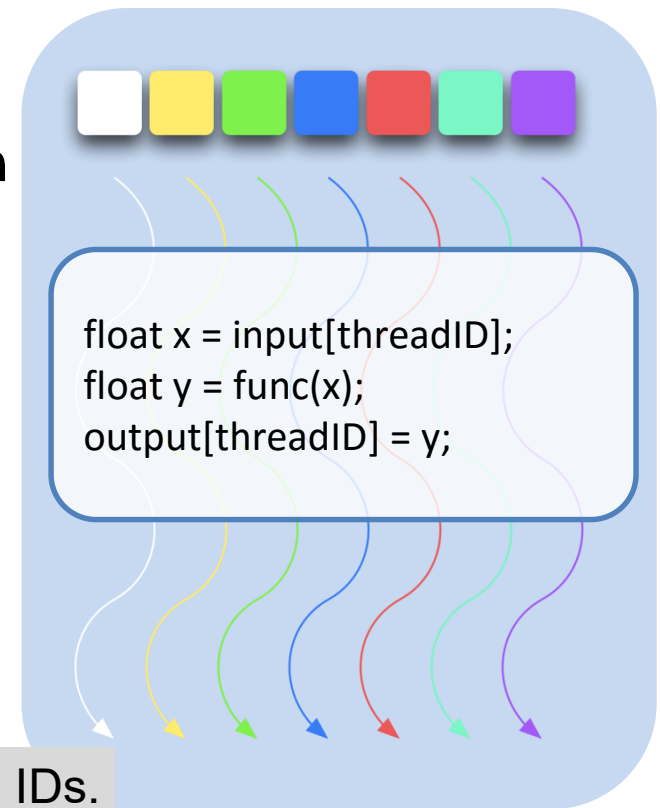
    cudaFree(d_x); cudaFree(d_y);
    free(h_x); free(h_y);
    return 0;
}
```

## ■ Definitions

- **Host**: CPU, executes functions
- **Device**: usually GPU, executes kernels

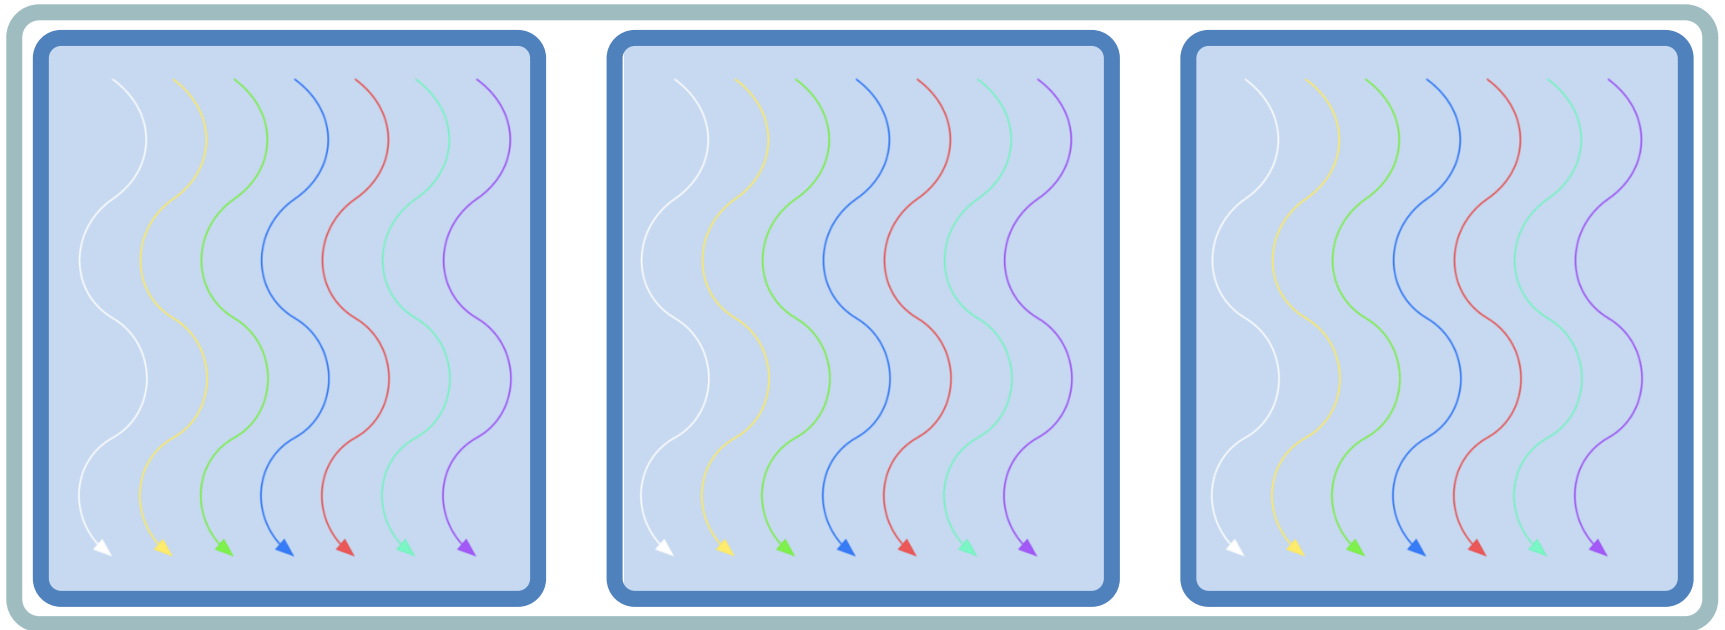
## ■ Parallel portion of application executed on

- Kernel is executed as array of **threads**
- All threads execute the same code
- Threads are identified by **IDs**
  - Select input/output data
  - Control decisions



© NVIDIA Corporation 2010

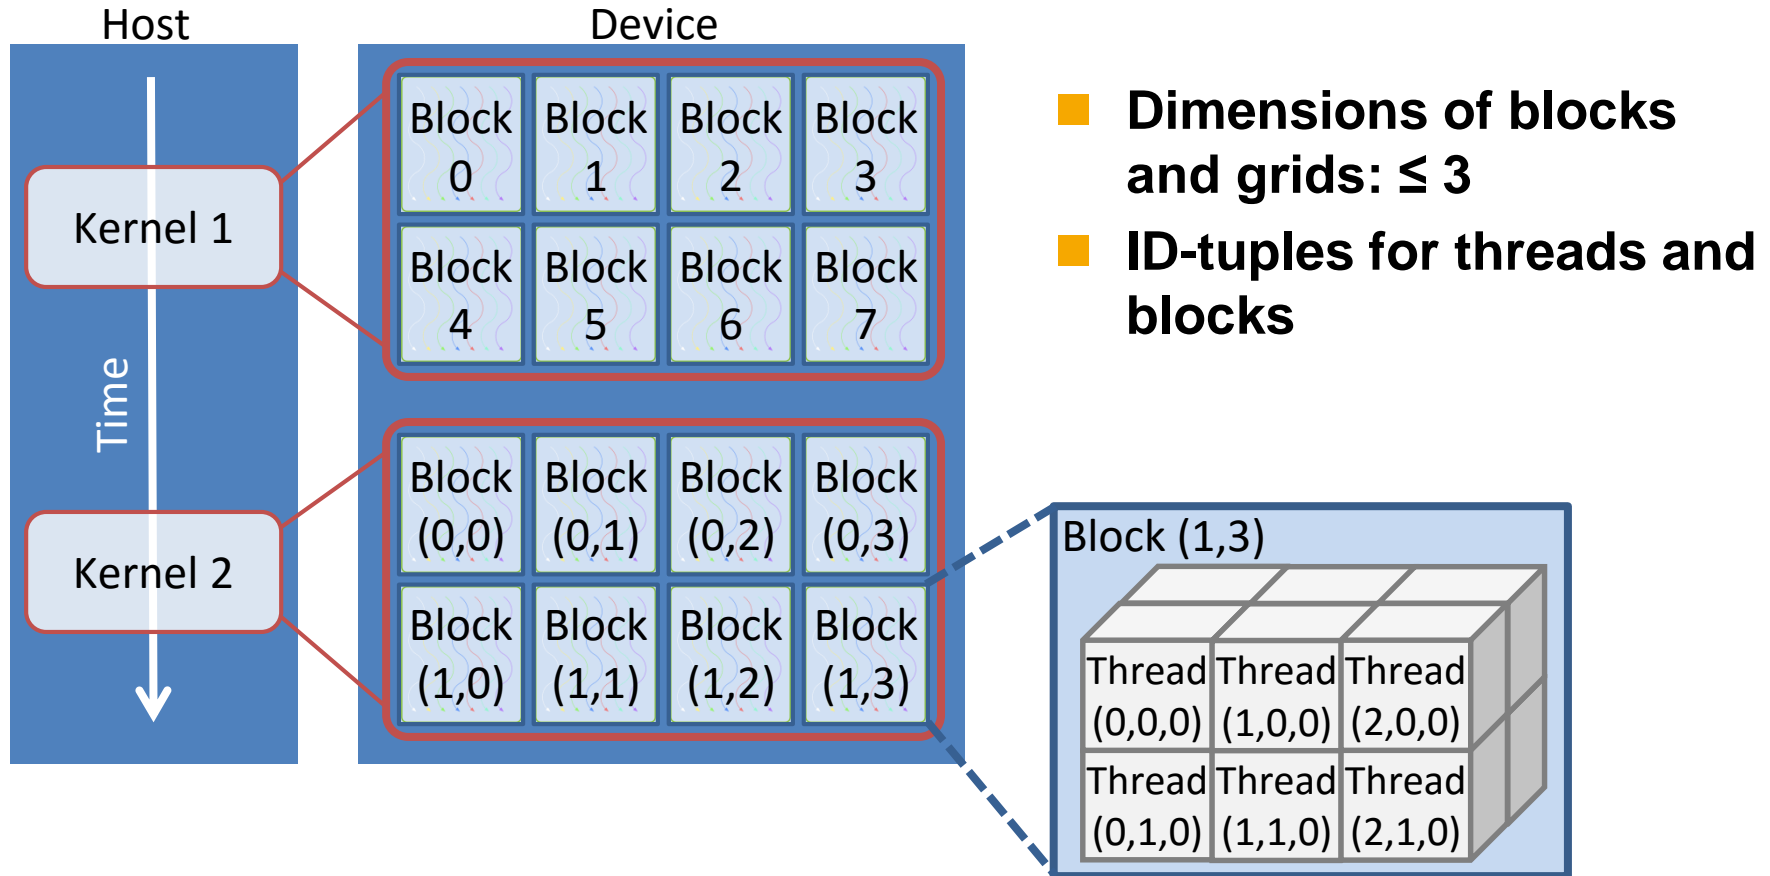
With OpenACC, you don't have to bother with thread IDs.



Based on: NVIDIA Corporation 2010

- Threads are grouped into **blocks**
- Blocks are grouped into a **grid**

- Kernel is executed as a grid of blocks of threads



# Example SAXPY: Kernel usage



## CUDA

```
__global__ void saxpyCUDA(int n, float a,
    float *x, float *y)
{
    int i = blockIdx.x * blockDim.x +
        threadIdx.x;
    if (i < n){
        y[i] = a*x[i] + y[i];
    }
}

int main(int argc, char* argv[])
{
    [...]
    // Invoke parallel SAXPY kernel
    dim3 threadsPerBlock(128);
    dim3 blocksPerGrid(n/threadsPerBlock.x);
    saxpyCUDA<<<blocksPerGrid,
        threadsPerBlock>>>(n,2.0,d_x,d_y);
    [...]
}
```

## OpenACC

```
void saxpyOpenACC (int n, float a, float
    *x, float *y)
{
    #pragma acc kernels or #pragma acc parallel
    #pragma acc loop gang vector(128)
        for (int i=0; i < n; ++i) {
            y[i] = a*x[i] + y[i];
        }
}

int main(int argc, char* argv[])
{
    [...]
    // Invoke parallel SAXPY kernel

    saxpyOpenACC(n,2.0,d_x,d_y);

    [...]
}
```

optional  
for  
kernels

## ■ Kernel code

- Function qualifiers: `__global__`, `__device__`, `__host__`
- Built-in variables: `gridDim`: contains dimensions of grid (type `dim3`)  
`blockDim`: contains dimensions of block (type `dim3`)  
`blockIdx`: contains block index within grid (type `uint3`)  
`threadIdx`: contains thread index within block (type `uint3`)
- Compute unique IDs, e.g. global 1D Idx:  
`gIdx = blockIdx.x * blockDim.x + threadIdx.x`

## ■ Kernel usage

- Compiling with `nvcc` (creating PTX code)
- ▶ Kernel arguments can be passed directly to the kernel
- ▶ Kernel invocation with *execution configuration* (chevron syntax):  
`func<<<dimGrid, dimBlock>>> (parameter)`



## ■ Offload region

→ Region maps to a CUDA kernel function

### C/C++

```
#pragma acc parallel [clauses]
```

### Fortran

```
!$acc parallel [clauses]
```

```
!$acc end parallel
```

- User responsible for finding parallelism (loops)
- **acc loop** needed for work-sharing
- No automatic sync between several loops

### C/C++

```
#pragma acc kernels [clauses]
```

### Fortran

```
!$acc kernels [clauses]
```

```
!$acc end kernels
```

- Compiler responsible for finding parallelism (loops)
- **acc loop** directive only for tuning needed
- Automatic sync between loops within kernels region

## ■ Clauses for compute constructs (`parallel`, `kernels`)

	C/C++, Fortran
→ If <i>condition</i> true, acc version is executed.....	<code>if(condition)</code>
→ Executes async, see Tuning slides.....	<code>async [(int-expr)]</code>
→ Define number of parallel gangs <sub>(parallel only)</sub> .....	<code>num_gangs(int-expr)</code>
→ Define number of workers within gang <sub>(parallel only)</sub> .....	<code>num_workers(int-expr)</code>
→ Define length for vector operations <sub>(parallel only)</sub> .....	<code>vector_length(int-expr)</code>
→ Reduction with <i>op</i> at end of region <sub>(parallel only)</sub> .....	<code>reduction(op:list)</code>
→ H2D-copy at region start + D2H at region end .....	<code>copy(list)</code>
→ Only H2D-copy at region start .....	<code>copyin(list)</code>
→ Only D2H-copy at region end .....	<code>copyout(list)</code>
→ Allocates data on device, no copy to/from host .....	<code>create(list)</code>
→ Data is already on device .....	<code>present(list)</code>
→ Test whether data on device. If not, transfer.....	<code>present_or_*(list)</code>
→ See Tuning slides.....	<code>deviceptr(list)</code>
→ Copy of each <i>list</i> -item for each parallel gang <sub>(parallel only)</sub>	<code>private(list)</code>
→ As <code>private</code> + copy initialization from host <sub>(parallel only)</sub> .	<code>firstprivate(list)</code>

OpenACC 2.0 also: `wait`, `device_type`, `default(none)`

## ■ Share work of loops

- Loop work gets distributed among threads on GPU (in certain schedule)

### C/C++

```
#pragma acc loop [clauses]
```

### Fortran

```
!$acc loop [clauses]
```

**kernels** loop defines loop schedule by int-expr in gang, worker or vector (instead of num\_gangs etc with parallel)

## ■ Loop clauses

- Distributes work into thread blocks .....
  - Distributes work into warps .....
  - Distributes work into threads within warp/ thread block .....
  - Executes loop sequentially on the device.....
  - Collapse *n* tightly nested loops.....
  - Says independent loop iterations (kernels loop only)....
  - Reduction with *op*.....
  - Private copy for each loop iteration.....
- | C/C++, Fortran                       | Loop schedule |
|--------------------------------------|---------------|
| gang                                 | }             |
| worker                               |               |
| vector                               |               |
| seq                                  |               |
| collapse( <i>n</i> )                 |               |
| independent                          |               |
| reduction( <i>op</i> : <i>list</i> ) |               |
| private( <i>list</i> )               |               |

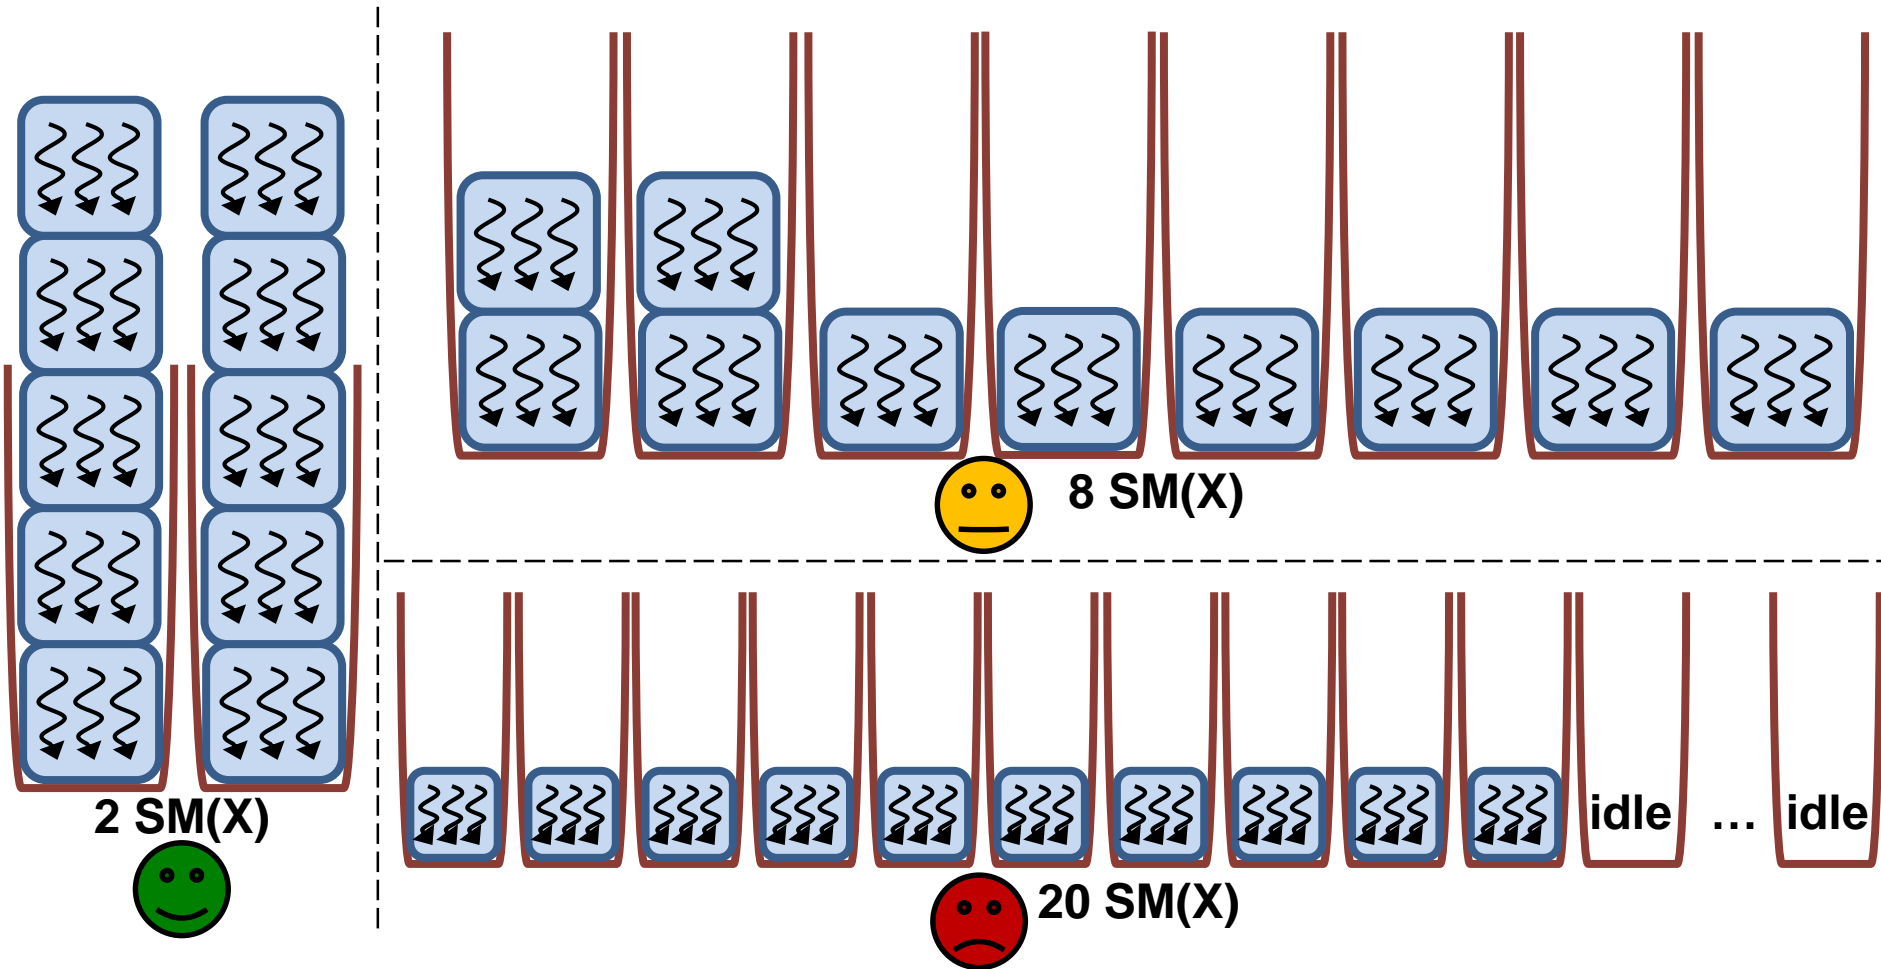
OpenACC 2.0 also: auto, tile, device\_type

background OpenACC

## ■ Why blocks?

- **Cooperation** of threads within a block possible
  - Synchronization (barrier)
  - Share data/ results using shared memory
- **Scalability**
  - Fast communication between  $n$  threads is not feasible when  $n$  large
    - No global synchronization on GPU possible (only by completing one kernel and starting another one from the CPU)
  - But: blocks are executed independently
  - Blocks can be distributed across arbitrary number of multiprocessors
    - In any order, concurrently, sequentially

■ Assume: 10 thread blocks



1. Why supercomputers?
2. Modern processors
3. Basic optimization techniques for serial code
4. Data access optimization
5. Parallel computers
6. Energy efficiency
7. Parallelization and optimization strategies
8. Parallel algorithms
9. Distributed-memory programming with MPI
10. Shared-memory programming with OpenMP
11. Hybrid programming (MPI + OpenMP)

## 12. Heterogeneous architectures (GPUs, Xeon Phis)

- Motivation
- RWTH Environment
- **GPGPU – Basics**
  - Overview
  - GPU Architecture

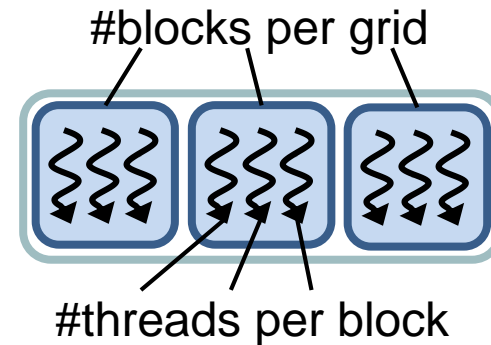
- Programming Model
- **Execution Model**
- Memory Model
- Misc on OpenACC
- Summary
- GPGPU – Advanced
  - Kernel optimizations
  - GPU-CPU interaction
- Intel Xeon Phi
  - Motivation
  - KNL Architecture
  - Vectorization
- Links

## ■ Host-directed execution model

- Main program runs on host
- Certain code regions run on device

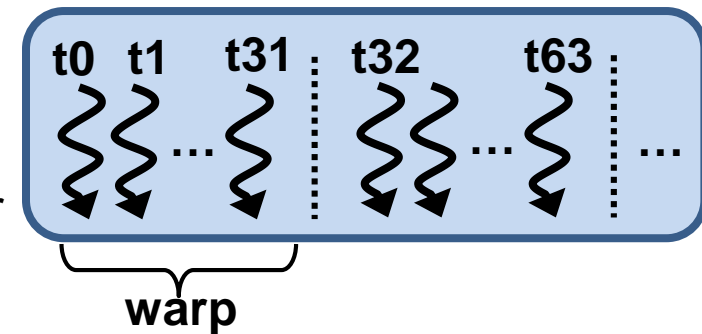
## ■ Launch configuration

- blocks per grid, threads per block

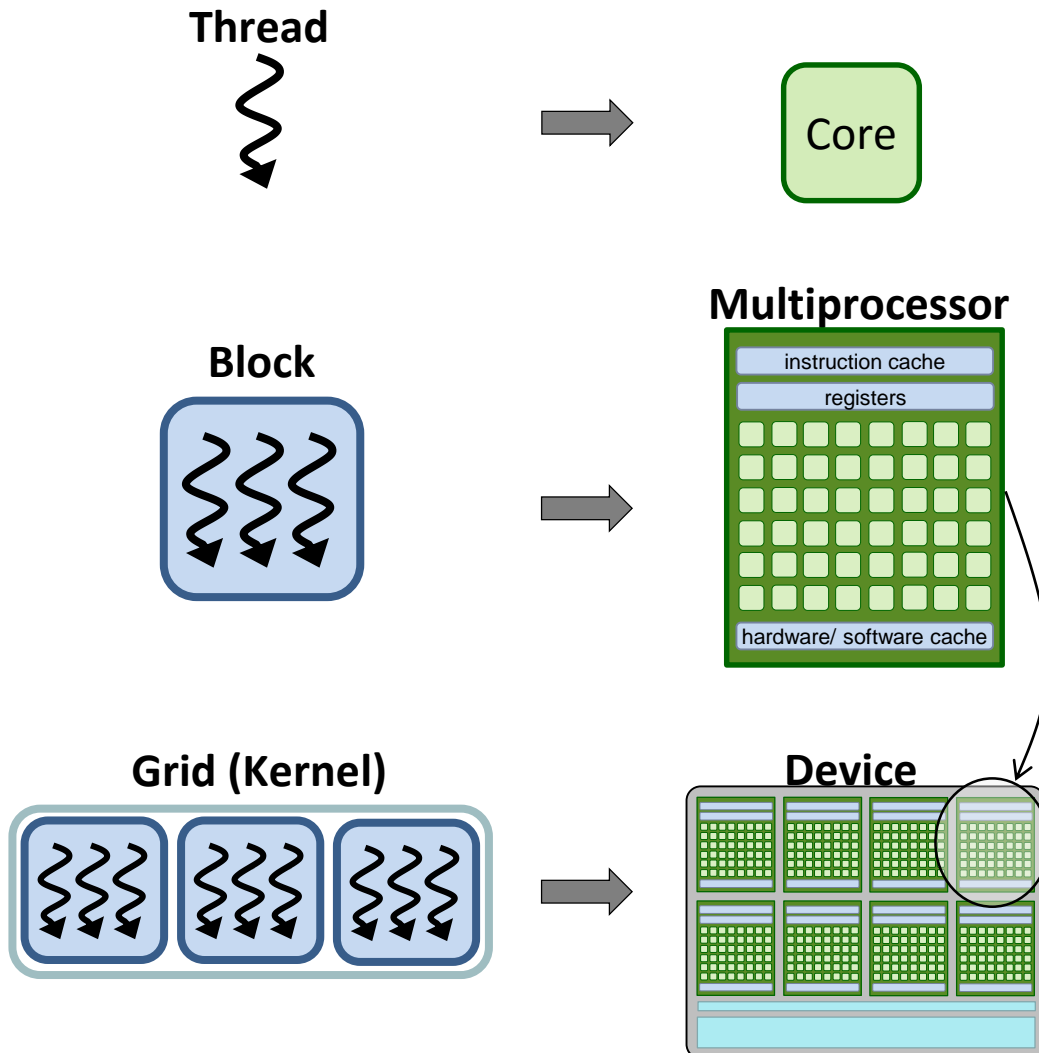


## ■ Warps

- Threads execute as groups of 32
- Threads in warp share same program counter



## → Single instruction multiple threads (SIMT)



→ Each thread is executed by a core

→ Each block is executed on a SM(X)

→ Several concurrent blocks can reside on a SM(X) depending on shared resources

→ Each kernel is executed on a device



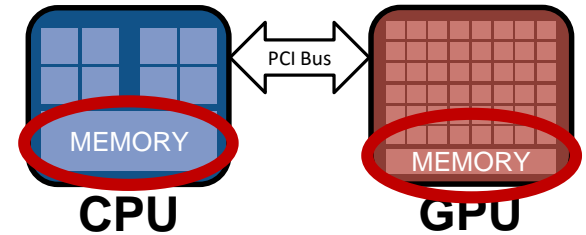
1. Why supercomputers?
2. Modern processors
3. Basic optimization techniques for serial code
4. Data access optimization
5. Parallel computers
6. Energy efficiency
7. Parallelization and optimization strategies
8. Parallel algorithms
9. Distributed-memory programming with MPI
10. Shared-memory programming with OpenMP
11. Hybrid programming (MPI + OpenMP)

## 12. Heterogeneous architectures (GPUs, Xeon Phis)

- Motivation
- RWTH Environment
- **GPGPU – Basics**
  - Overview
  - GPU Architecture

- Programming Model
- Execution Model
- **Memory Model**
- Misc on OpenACC
- Summary
- GPGPU – Advanced
  - Kernel optimizations
  - GPU-CPU interaction
- Intel Xeon Phi
  - Motivation
  - KNL Architecture
  - Vectorization
- Links

- **Host + device memory = separate entities**
- **No coherence between host + device**
  - Data synchronization/transfer



- **Host**
  - (De-)Allocates device memory (global, constant, texture)
  - Triggers data transfer
- **Device**
  - Works on device memory (hierarchy)

## ■ Thread

→ *Registers*

**Fermi**: 63 per thread

**K20**: 255 per thread

→ *Local* memory

## ■ Block

→ *Shared* memory

**Fermi**: 64KB configurable, on-chip

16KB shared + 48KB L1 OR

48KB shared + 16KB L1

**K20**: 64KB configurable, on-chip

16KB shared + 48KB L1 OR

48KB shared + 16KB L1 OR

32KB shared + 32KB L1

## ■ Grid/ application

→ *Constant* memory

read-only; off-chip; cached

→ *Global* memory

several GB; off-chip

**Fermi/K20**: L2 cache

## ■ Caches

→ L1

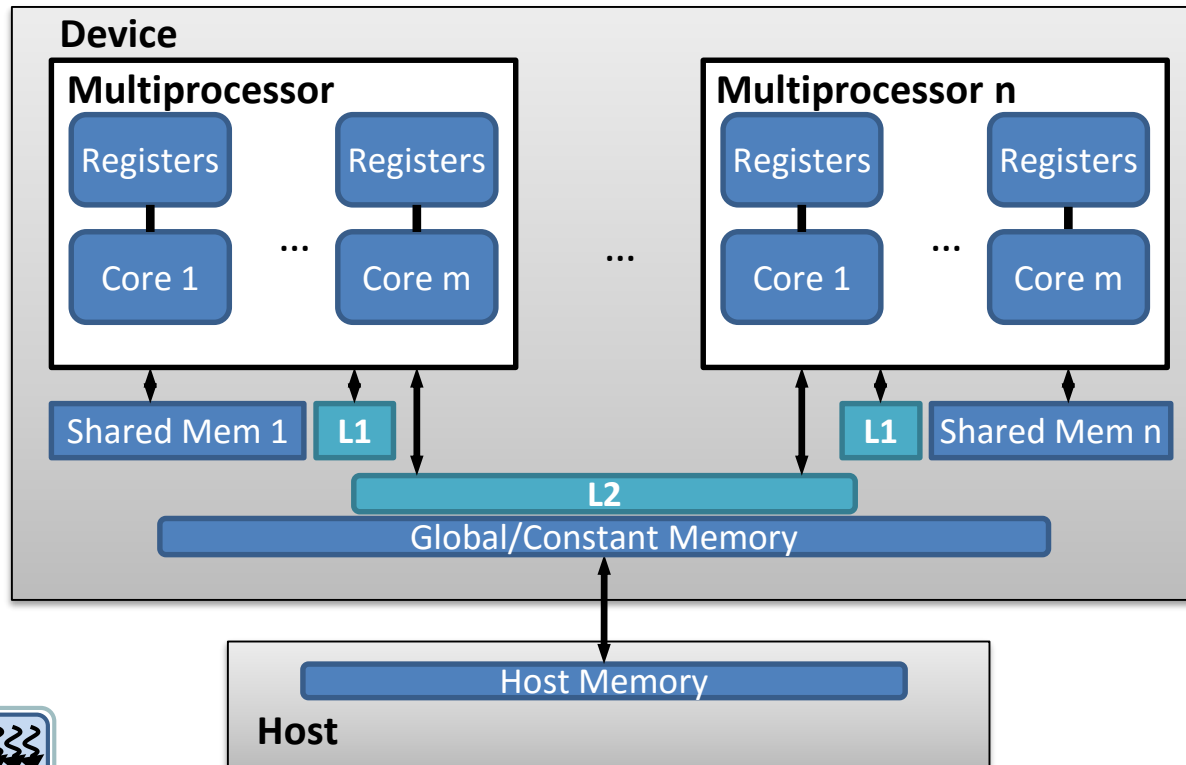
**Fermi**: configurable 16/48KB

**K20**: configurable 16/32/ 48 KB

→ L2

**Fermi**: 768KB

**K20**: 1536KB



# Example SAXPY: Memory



## CUDA

```
int main(int argc, char* argv[]) {  
    float* h_x, *h_y; // host pointer  
    // Allocate and initialize h_x and h_y  
  
    float *d_x, *d_y; // device pointer  
    cudaMalloc(&d_x, n*sizeof(float));  
    cudaMalloc(&d_y, n*sizeof(float));  
  
    cudaMemcpy(d_x, h_x, n * sizeof(float),  
               cudaMemcpyHostToDevice);  
    cudaMemcpy(d_y, h_y, n * sizeof(float),  
               cudaMemcpyHostToDevice);  
  
    // Invoke parallel SAXPY kernel  
  
    cudaMemcpy(h_y, d_y, n * sizeof(float),  
               cudaMemcpyDeviceToHost);  
  
    cudaFree(d_x); cudaFree(d_y);  
    free(h_x); free(h_y); return 0;  
}
```

## OpenACC

```
int main(int argc, char* argv[]) {  
    float* h_x, *h_y; // host pointer  
    // Allocate and initialize h_x and h_y  
  
    #pragma acc data copyin(x[0:n]  
                           copy(y[0:n]))  
    {  
        #pragma acc kernels  
            copyin(x[0:n]) copy(y[0:n])  
        #pragma acc loop gang vector(128)  
        // Ir for (int i=0; i < n; ++i) {  
            y[i] = a*x[i] + y[i];  
        }  
    }  
  
    free(h_x); free(h_y); return 0;  
}
```

may be  
optional

# Example SAXPY – CUDA (all)



```
__global__ void saxpyCUDA(int n, float a,
float *x, float *y) {
    int i = blockIdx.x * blockDim.x +
threadIdx.x;
    if (i < n){
        y[i] = a*x[i] + y[i];
    }
}
```

```
int main(int argc, char* argv[]) {
    int n = 10240; float a = 2.0f;
    float* h_x,*h_y; // Pointer to CPU memory
    h_x = (float*) malloc(n* sizeof(float));
    h_y = (float*) malloc(n* sizeof(float));
    // Initialize h_x, h_y
    for(int i=0; i<n; ++i){
        h_x[i]=i;
        h_y[i]=5.0*i-1.0;
    }
}
```

```
float *d_x,*d_y; // Pointers to GPU memory
```

**1. Allocate data on GPU + transfer data to CPU**

```
cudaMemcpy(d_x, h_x, n * sizeof(float),
cudaMemcpyHostToDevice);
cudaMemcpy(d_y, h_y, n * sizeof(float),
cudaMemcpyHostToDevice);
```

```
// Invoke parallel SAXPY kernel
```

```
dim3 threadsPerBlock(128);
dim3 blocksPerGrid(n/128);
saxpyCUDA<<<blocksPerGrid,
threadsPerBlock>>>(n, 2.0, d_x, d_y);
```

**2. Launch kernel**

```
cudaMemcpy(h_y, d_y, n * sizeof(float),
cudaMemcpyDeviceToHost);
```

**3. Transfer data to CPU + free data on GPU**

```
cudaFree(d_x); cudaFree(d_y);
```

```
free(h_x); free(h_y);
```

```
return 0;
```

```
}
```

## ■ Variable type qualifiers

`__device__`, `__shared__`, `__constant__`

## ■ Memory management

`cudaMalloc(pointerToGPUMem, size)`

`cudaFree(pointerToGPUMem)`

## ■ Memory transfer (synchronous)

`cudaMemcpy(dest, src, size, direction)`

direction:  
`cudaMemcpyHostToDevice`  
`cudaMemcpyDeviceToHost`  
`cudaMemcpyDeviceToDevice`  
`cudaMemcpyDefault` (with UVA)

## ■ Data region

→ Decouples data movement from offload regions

### C/C++

```
#pragma acc data [clauses]
```

### Fortran

```
!$acc data [clauses]
```

```
!$acc end data
```

## ■ Data clauses

→ Triggers data movement of denoted arrays

→ If *cond* true, move data to accelerator.....

→ H2D-copy at region start + D2H at region end .....

→ Only H2D-copy at region start .....

→ Only D2H-copy at region end .....

→ Allocates data on device, no copy to/from host .....

→ Data is already on device .....

→ Test whether data on device. If not, transfer.....

→ See Tuning slides.....

### C/C++, Fortran

```
if (cond)
```

```
copy(list)
```

```
copyin(list)
```

```
copyout(list)
```

```
create(list)
```

```
present(list)
```

```
present_or_*(list)
```

```
deviceptr(list)
```

1. Why supercomputers?
2. Modern processors
3. Basic optimization techniques for serial code
4. Data access optimization
5. Parallel computers
6. Energy efficiency
7. Parallelization and optimization strategies
8. Parallel algorithms
9. Distributed-memory programming with MPI
10. Shared-memory programming with OpenMP
11. Hybrid programming (MPI + OpenMP)

## 12. Heterogeneous architectures (GPUs, Xeon Phis)

- Motivation
- RWTH Environment
- **GPGPU – Basics**
  - Overview
  - GPU Architecture

- Programming Model
- Execution Model
- Memory Model
- **Misc on OpenACC**
- Summary
- GPGPU – Advanced
  - Kernel optimizations
  - GPU-CPU interaction
- Intel Xeon Phi
  - Motivation
  - KNL Architecture
  - Vectorization
- Links



## ■ Combined directives

```
#pragma acc kernels loop
for (int i=0; i<n; ++i) { /*...*/}

#pragma acc parallel loop
for (int i=0; i<n; ++i) { /*...*/}
```

## ■ Reductions

```
#pragma acc parallel
#pragma acc loop reduction (+:sum)
for (int i=0; i<n; ++i) {
    sum += i;
}
```

Also possible:  
\*, max, min,  
&, |, ^, &&, ||

PGI compiler can often recognize reductions on its own. See compiler feedback, e.g.:  
**Sum reduction generated for var**

```
template <unsigned int blockSize>
__global__ void reduce6(int *g_idata, int *g_odata, unsigned int n)
{
    extern __shared__ int sdata[];
    unsigned int tid = threadIdx.x;
    unsigned int i = blockIdx.x*(blockSize*2) + tid;
    unsigned int gridSize = blockSize*2*gridDim.x;
    sdata[tid] = 0;
    while (i < n) { sdata[tid] += g_idata[i] + g_idata[i+blockSize]; i += gridSize; }
    __syncthreads();
    if (blockSize >= 512) { if (tid < 256) { sdata[tid] += sdata[tid + 256]; }
    __syncthreads(); }
    if (blockSize >= 256) { if (tid < 128) { sdata[tid] += sdata[tid + 128]; }
    __syncthreads(); }
    if (blockSize >= 128) { if (tid < 64) { sdata[tid] += sdata[tid + 64]; }
    __syncthreads(); }
    if (tid < 32) {
        if (blockSize >= 64) sdata[tid] += sdata[tid + 32];
        if (blockSize >= 32) sdata[tid] += sdata[tid + 16];
        if (blockSize >= 16) sdata[tid] += sdata[tid + 8];
        if (blockSize >= 8) sdata[tid] += sdata[tid + 4];
        if (blockSize >= 4) sdata[tid] += sdata[tid + 2];
        if (blockSize >= 2) sdata[tid] += sdata[tid + 1];
    }
    if (tid == 0) g_odata[blockIdx.x] = sdata[0];
}
```

Source: NVIDIA, "Optimizing Parallel Reduction in CUDA"

[Introduction to HPC](#)

Prof. Matthias Müller | IT Center der RWTH Aachen University

Efficient reduction implementations with CUDA are hard! See the NVIDIA SDK for hints.

## ■ Data clauses can be used on data, kernels or parallel

→ `copy`, `copyin`, `copyout`, `present`, `present_or_copy`,  
`create`, `deviceptr`

## ■ Array shaping

→ Compiler sometimes cannot determine size of arrays

→ Specify explicitly using data clauses and array “shape”

### C/C++

```
#pragma acc data copyin(a[0:length]) copyout(b[s/2:s/2])
```

### Fortran

```
!$acc data copyin(a(0:length-1)) copyout(b(s/2:s))
```

[lower bound: size]

[lower bound: upper bound]

```
#pragma acc data copy(x[0:n])
{
    for (t=0; t<T; t++){
        // Modify x on device (e.g. in subroutine
        // w/o data copying)

        #pragma acc update host(x[0:n])

        // Modify x on host

        #pragma acc update device(x[0:n])
    }
}
```

## ■ Update executable directive

- Move data from GPU to host, or host to GPU
- Used to update existing data after it has changed in its corresponding copy

### C/C++

```
#pragma acc update host|device [clauses]
```

### Fortran

```
!$acc update host|device [clauses]
```

OpenACC 2.0:  
update self is  
preferred over  
update host.  
wait clause also  
possible.

- Data movement can be conditional or asynchronous

## ■ Update clauses

- *list* variables are copied from acc to host.....
- *list* variables are copied from host to acc.....
- If *cond* true, move data to accelerator.....
- Executes async, see Tuning slides.....

### C/C++, Fortran

```
host(list)  
device(list)  
if(cond)  
async[(int-expr)]
```

## ■ Structured data lifetime

→ Since OpenACC 1.0

```
#pragma acc data copyin(x[0:n]) \  
                create(y[0:n])  
  
{  
    // data lifetime  
}
```

## ■ Unstructured data lifetime

→ Since OpenACC 2.0

→ **enter data**: allocates (+ copies) data to device memory

→ **exit data**: deallocates (+ copies) data from device memory

```
class Matrix {  
    Matrix() {  
        v = new double[n];  
        #pragma acc enter data create(v[0:n])  
    }  
    ~Matrix() {  
        #pragma acc exit data delete(v[0:n])  
        delete[] v;  
    }  
private:  
    double* v;  
}
```

Also possible:  
**copyin**

Also possible:  
**copyout**

## ■ Enter data construct

- Allocation (& copy) of scalars and (sub-)arrays in(to) device memory
- Remain there until end of program or corresponding exit data call

### C/C++

```
#pragma acc enter data clause-list
```

### Fortran

```
!$acc enter data clause-list
```

### C/C++, Fortran

```
copyin(list)
```

```
create(list)
```

```
present_or_copyin(list)
```

```
present_or_create(list)
```

- Specific clauses for copy or just creation.....

## ■ Exit data construct

- (Copies data to host memory &) deletes data from device memory

### C/C++

```
#pragma acc exit data clause-list
```

### Fortran

```
!$acc exit data clause-list
```

### C/C++, Fortran

```
copyout(list)
```

```
delete(list)
```

- Specific clauses for copy or deletion.....

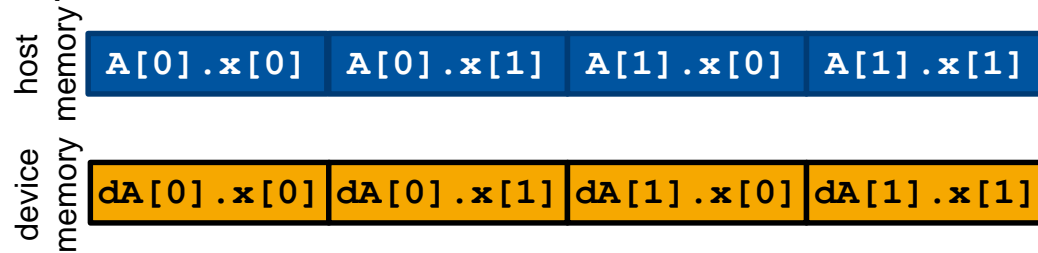
## ■ Clauses valid for both: **if**, **async**, **wait**

## ■ Support of a “flat” object model

→ Primitive types

→ Composite types w/o allocatable/pointer members

```
struct {  
  int x[2]; // static size 2  
} *A; // dynamic size 2  
#pragma acc data copy(A[0:2])
```

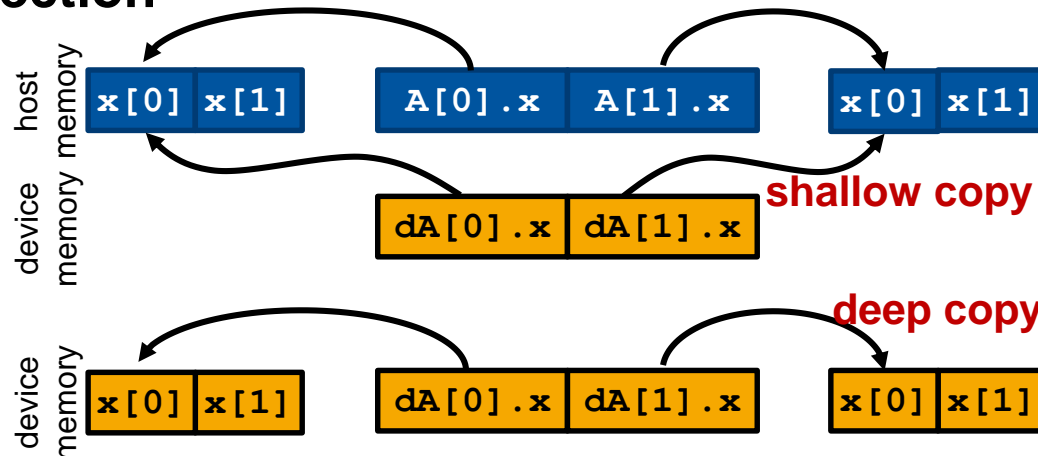


## ■ Challenges with pointer indirection

→ Non-contiguous transfers

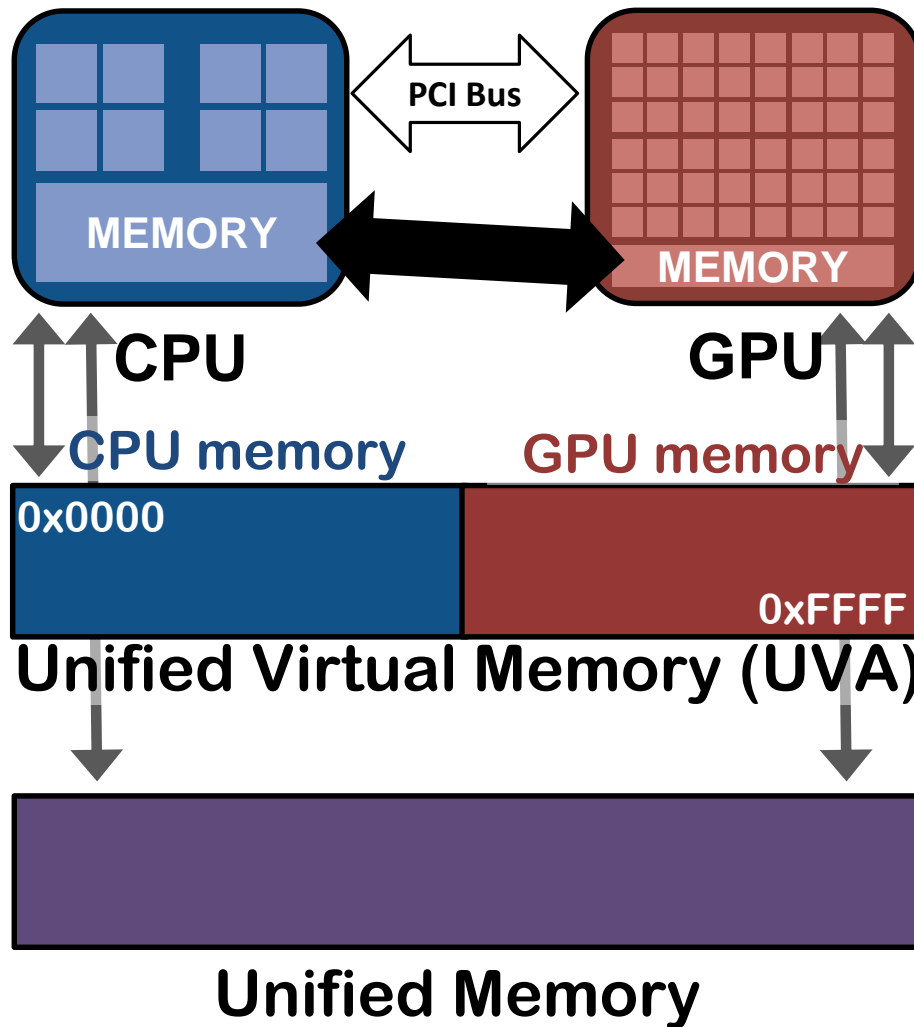
→ Pointer translation

```
struct {  
  int *x; // dynamic size 2  
} *A; // dynamic size 2  
#pragma acc data copy(A[0:2])
```



## ■ Manual deep copy is possible with data API, but usually not practical





- Explicit data transfers
- Single virtual address space for all CPU and GPU memory
- Did not get rid of the required explicit memory copying
- Depends on UVA
- “Managed memory” accessible by single pointer for CPU/GPU
- System automatically migrates corresponding data

# Example SAXPY: Managed Memory



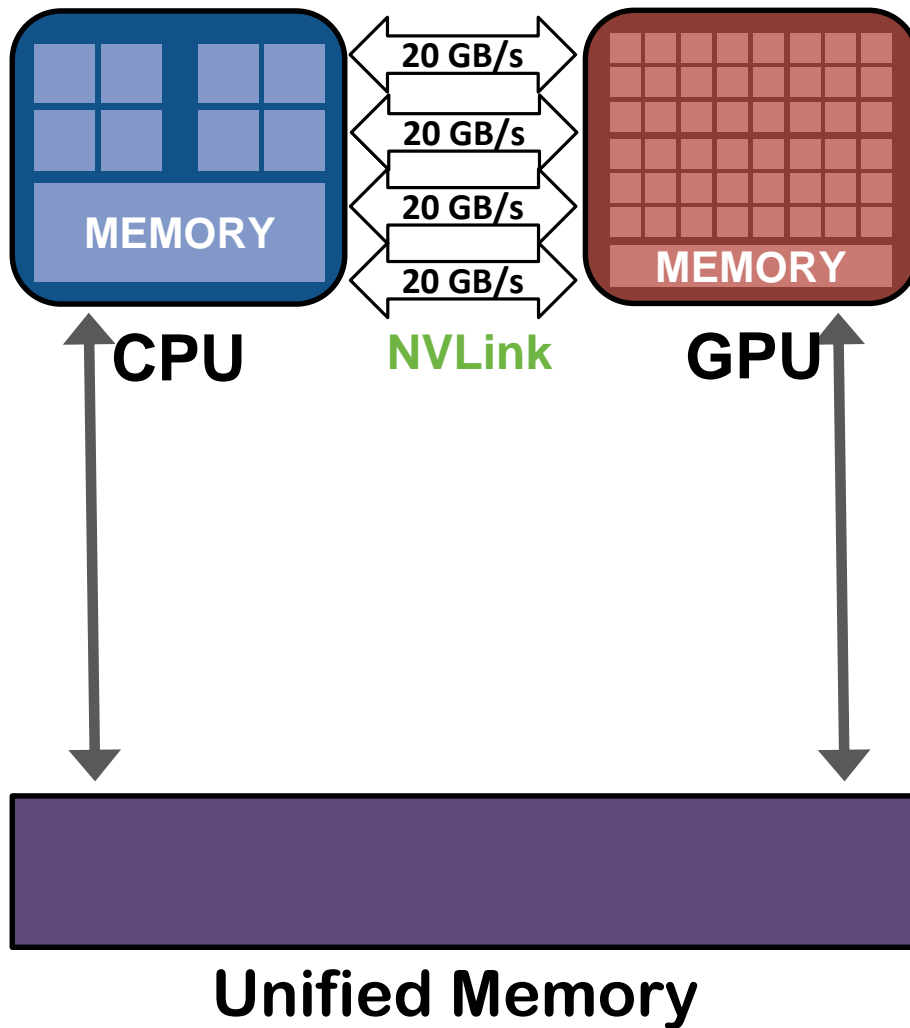
CUDA

OpenACC

```
pgcc -acc -ta=tesla:managed saxypy.c
```

All data clauses  
will be ignored

```
int main(int argc, char* argv[]) {  
    float* h_x, *h_y; // host pointer  
    cudaMallocManaged(&h_x, n*sizeof(float));  
    cudaMallocManaged(&h_y, n*sizeof(float));  
    // Initialize h_x and h_y  
  
    // Invoke parallel SAXPY kernel  
    dim3 threadsPerBlock(128);  
    dim3 blocksPerGrid(n/threadsPerBlock.x);  
    saxpyCUDA<<<blocksPerGrid,  
        threadsPerBlock>>>(n, 2.0, h_x, h_y);  
  
    cudaFree(h_x); cudaFree(h_y);  
    return 0;  
}
```



- **Interconnect technology by NVIDIA (2016)**

- 20 GB/s per link

- **Use cases**

- CPU-GPU

- GPU-GPU

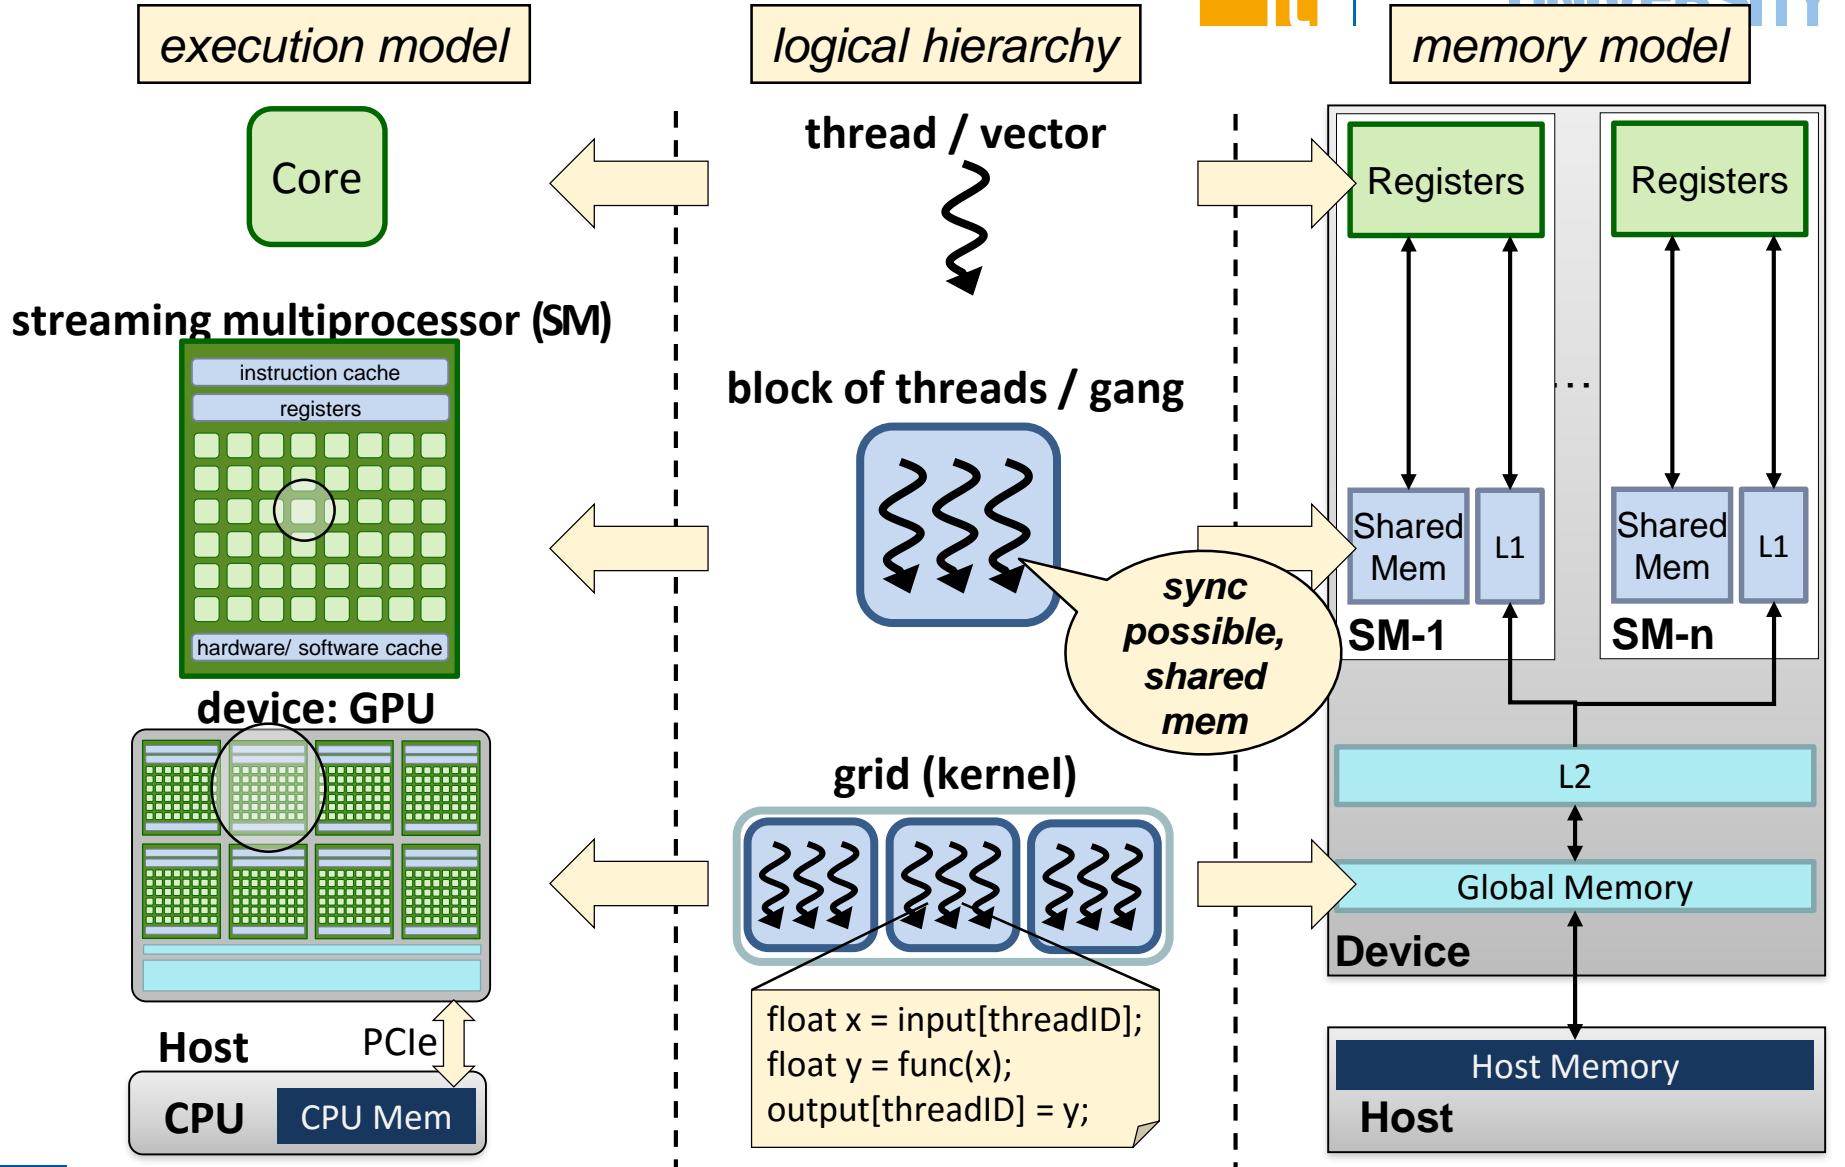
1. Why supercomputers?
2. Modern processors
3. Basic optimization techniques for serial code
4. Data access optimization
5. Parallel computers
6. Energy efficiency
7. Parallelization and optimization strategies
8. Parallel algorithms
9. Distributed-memory programming with MPI
10. Shared-memory programming with OpenMP
11. Hybrid programming (MPI + OpenMP)

## 12. Heterogeneous architectures (GPUs, Xeon Phis)

- Motivation
- RWTH Environment
- **GPGPU – Basics**
  - Overview
  - GPU Architecture

- Programming Model
- Execution Model
- Memory Model
- Misc on OpenACC
- **Summary**
- GPGPU – Advanced
  - Kernel optimizations
  - GPU-CPU interaction
- Intel Xeon Phi
  - Motivation
  - KNL Architecture
  - Vectorization
- Links

# Summary



## ■ How does a GPU look like?

- Why do GPUs deliver a good performance per Watt ratio?
- What is the difference to CPUs?
- How does the memory hierarchy look like?
- How can the logical programming hierarchy be mapped to the execution model?

## ■ Which models can be used to program a GPU?

- How to handle offloading of regions?
- How to handle data management?
- What are the main differences?

- **What speedup would you expect when porting your code from a modern CPU node to the latest high-end GPU?**
  - a) 0.5x
  - b) 2x
  - c) 10x
  - d) 100x

### ■ Order the following interconnects by the bandwidth:

- Network fabrics: InfiniBand EDR (4x)
- CPU/ GPU interconnects: PCIe 3.0 x16
- Memory: 2-sockets each with CPU RAM DDR3 (3 channels), GPU GDDR5 (NVIDIA K20x)

Answers (low to high):

- a) InfiniBand, PCIe, CPU RAM, GPU DDR5
- b) PCIe, InfiniBand, CPU RAM, GPU DDR5
- c) CPU RAM, PCIe, InfiniBand, GPU DDR5
- d) PCIe, CPU RAM, GPU DDR5, InfiniBand



## Order 1

→ Network

→ CPU

→ Memory

GPU

Answer

Interconnect	Bandwidth
100 Gigabit Ethernet	100 Gbit/s
Infiniband EDR (4x or 12x)	100 Gbit/s or 290 Gbit/s
RAM DDR3 (3 channels) per socket	32 GB/s = 256 Gbit/s
Stream Benchmark (Intel Westmere, 2 sockets)	40 GB/s = 320 Gbit/s
GPU GDDR5 (NVIDIA K20x)	250 GB/s = 2000 Gbit/s
GPU Stream Benchmark	180 GB/s = 1440 Gbit/s
PCIe 3.0 x16 (GPU-CPU)	32 GB/s = 256 Gbit/s
NVIDIA NVLink (single link)	20 GB/s = 160 Gbit/s

- a) InfiniBand, PCIe, CPU RAM, GPU DDR5
- b) PCIe, InfiniBand, CPU RAM, GPU DDR5
- c) CPU RAM, PCIe, InfiniBand, GPU DDR5
- d) PCIe, CPU RAM, GPU DDR5, InfiniBand