

1. Why supercomputers?
2. Modern processors
3. Basic optimization techniques for serial code
4. Data access optimization
5. Parallel computers
6. Energy efficiency
7. Parallelization and optimization strategies
8. Parallel algorithms
9. Distributed-memory programming with MPI
10. Shared-memory programming with OpenMP
11. Hybrid programming (MPI + OpenMP)

12. Heterogeneous architectures (GPUs, Xeon Phis)

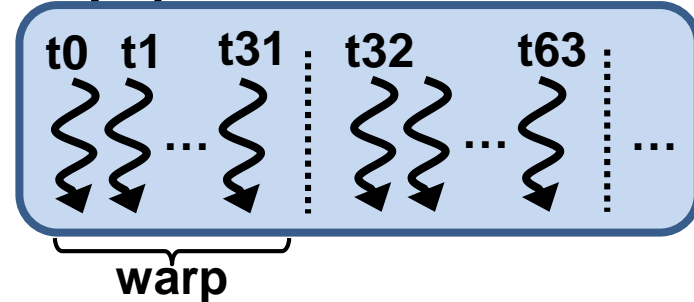
- Motivation
- RWTH Environment
- GPGPU – Basics
 - Overview
 - GPU Architecture

- Programming Model
- Execution Model
- Memory Model
- Misc on OpenACC
- Summary
- **GPGPU – Advanced**
 - **Kernel optimizations**
 - GPU-CPU interaction
- Intel Xeon Phi
 - Motivation
 - KNL Architecture
 - Vectorization
- Links

■ Threads execute as groups of 32 threads (warps)

→ Threads in warp share same program counter

→ SIMT architecture



■ Diverging threads within a warp

→ Threads within a single warp take different paths

→ Different execution paths within warp are serialized

```
if (threadIdx.x > 2)
{...} else {...}
```

→ Same code path for threads within warp

→ Threads execute synchronously

→ Granularity of branches: multiple of warp size

```
if (threadIdx.x / WARP_SIZE > 2) {...} else {...}
```

Different code for different warps: no impact on performance

■ Kernel executes grid of blocks of threads

→ Launch configuration

```
myKernel<<<numBlocks, numThreads>>>(...)
```

```
#pragma acc kernels loop gang(numBlocks) vector(numThreads)
```

```
#pragma acc parallel num_gangs(numBlocks) vector_length(numThreads)
```

■ Possible mapping to CUDA

terminology (GPUs) *compiler dependent*

→ gang = block

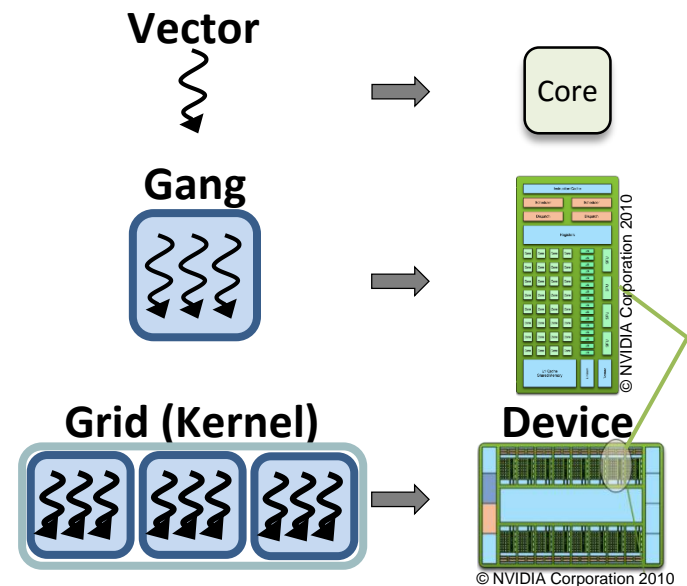
→ worker = warp

→ vector = threads

→ Within block (if omitting worker)

→ Within warp (if specifying worker)

■ Execution Model



■ How many threads/blocks to launch?

- OpenACC runtime tries to find good values automatically
 - Performance portability across different device types possible
 - Own tuning may deliver better performance on a certain hardware
- In CUDA codes, the developer has to specify both sizes

■ Hardware operation

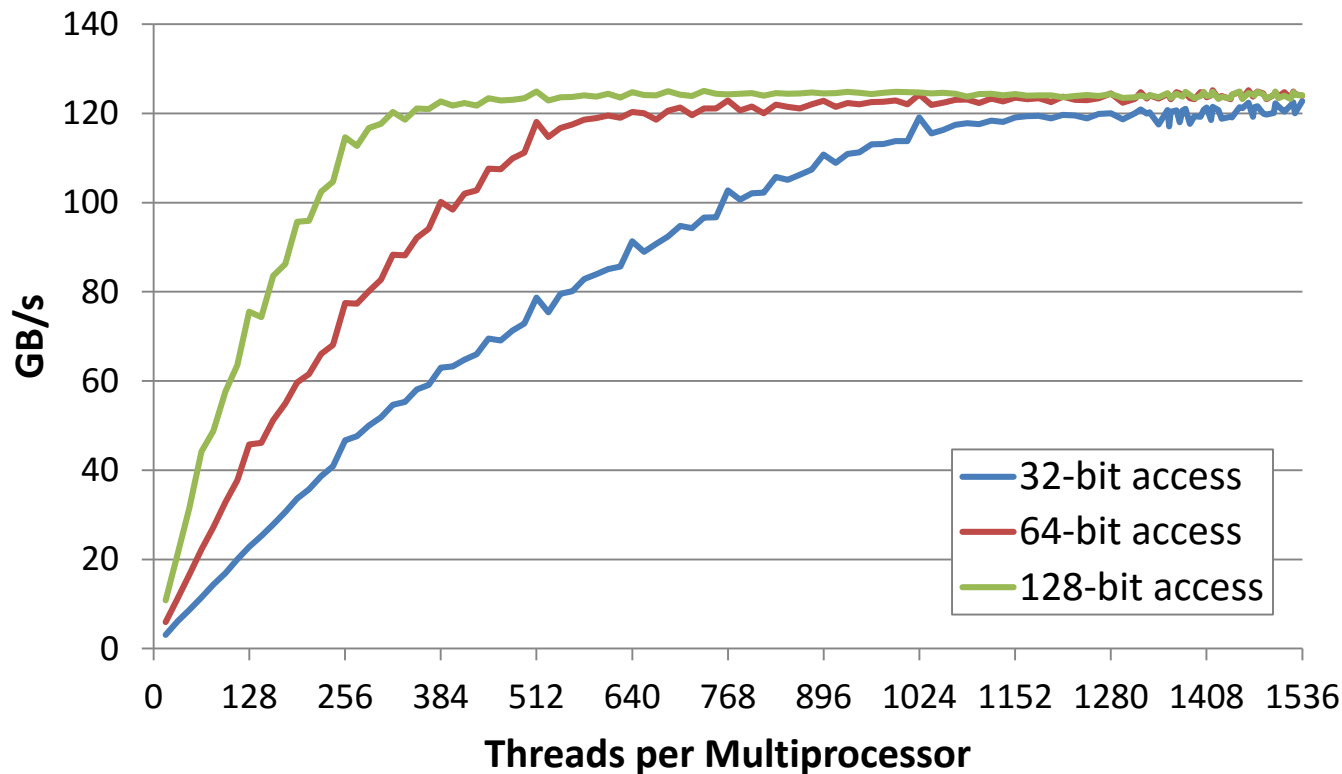
- Instructions are issued in order
- Thread execution stalls when one of the operands isn't ready
- Latency is hidden by switching threads
 - GMEM latency: 400-800 cycles
 - Arithmetic latency: 18-22 cycles

→ Need enough threads to hide latency

■ Maximizing global memory throughput

→ Example program: increment an array of ~67M elements

Impact of launch configuration

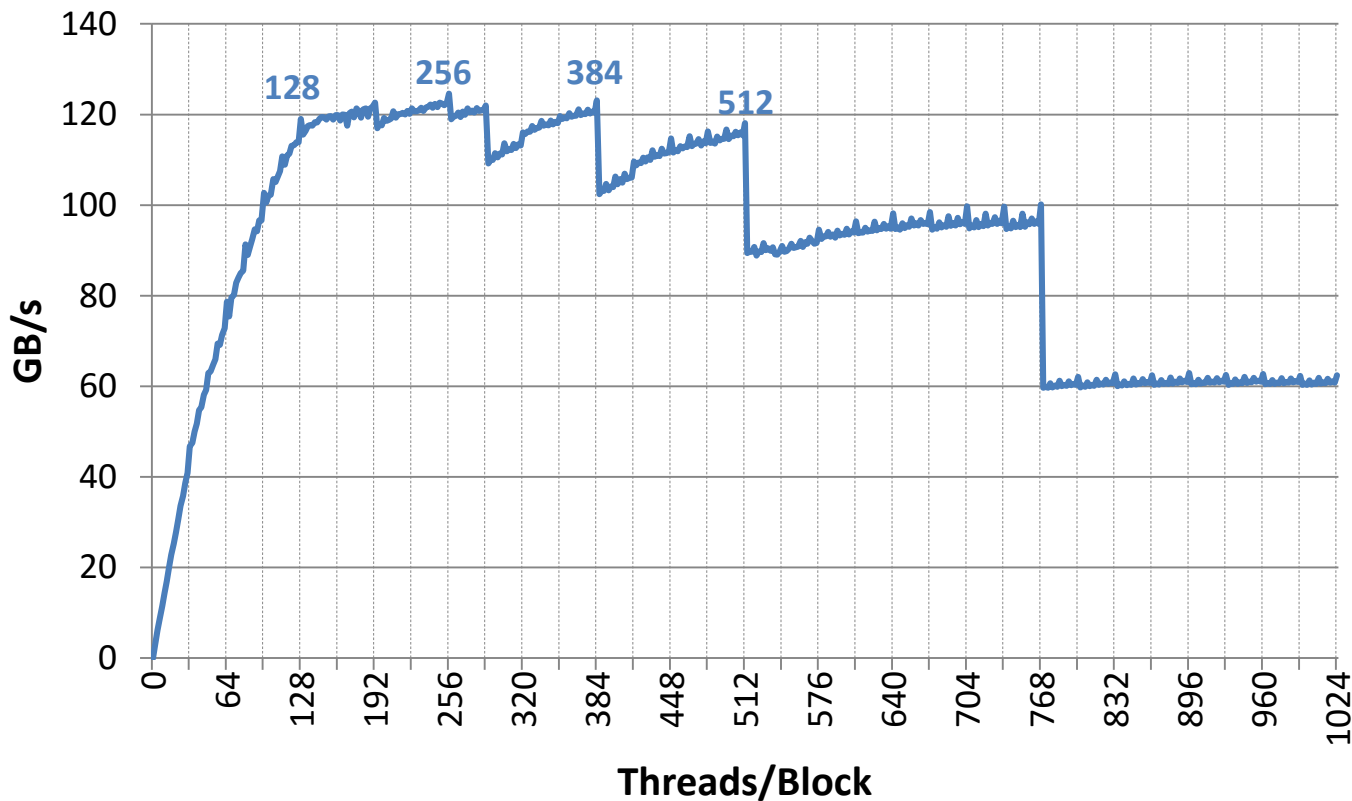


- NVIDIA Tesla C2050 (Fermi)
- ECC off
- Bandwidth: 144 GB/s

■ Maximizing global memory throughput

→ Example program: increment an array of ~67M elements

Impact of launch configuration (32-bit words)



- NVIDIA Tesla C2050 (Fermi)
- ECC off
- Bandwidth: 144 GB/s

■ Threads per block: Multiple of warp size (32)

- Threads are always spanned in warps onto resources
- Not used threads are marked as inactive
- Starting point: 128-256 threads/block

■ Blocks per grid heuristics

- `#blocks > #SM`
 - MPs have at least one block to execute
- `#blocks / #SM > 2`
 - MP can concurrently execute up to 8 blocks
 - Blocks that aren't waiting in barrier keep hardware busy
 - Subject to resource availability (registers, smem)
- `#blocks > 100` to scale to future devices
- Most obvious: `#blocks * #threads = #problem-size`
 - Multiple elements per thread may amortize setup costs of simple kernels:
$$\#blocks * \#threads < \#problem-size$$



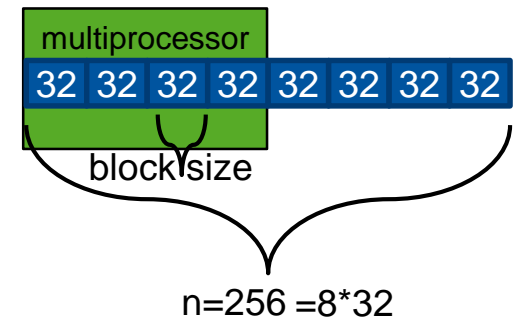
**Launch MANY
threads to keep
GPU busy!**

- **Main concept: hiding latencies by switching threads**
 - Strive for high occupancy
- **Occupancy (per SM) = $\frac{\text{active warps}}{\text{max. supported active warps}}$**
- **Some reasons for low occupancy**
 - Too few blocks launched on GPU (e.g., problem too small)
 - Unbalanced workloads
 - Device capabilities (warps, blocks)
 - Limit of shared resources (registers, shared memory)
- **Poor instruction issue efficiency (not able to hide latencies)**
 - BUT: does not necessarily mean low performance
 - E.g., several simultaneous kernel launches can increase utilization

Loop schedule (in examples)



```
int main(int argc, const char* argv[]) {
    int n = 256; int blocks = 4; int bsize = 32;
    float *x = (float*) malloc(n * sizeof(float));
    float *y = (float*) malloc(n * sizeof(float));
    // Define scalars a, b & initialize x, y
    // Run SAXPY TWICE
#pragma acc data copyin(x[0:n])
{
#pragma acc parallel copy(y[0:n]) present(x[0:n]) \
    num_gangs(blocks) vector_length(bsize)
    all do the same
    e.g. int tmp = 5;
#pragma acc loop gang
    workshare (w/o barrier)
#pragma acc loop vector
    workshare (w/ barrier)
    for (int i = 0; i < n; ++i){
        y[i] = a*x[i] + y[i];
    } // Change y on host & do second SAXPY on device
}
free(x); free(y); return 0; }
```



Note: Loop schedules are implementation dependent. This representation is one (common) example.

Loop schedule (in OpenACC examples)



```
#pragma acc parallel vector_length(256)
```

```
#pragma acc loop gang vector
```

```
    for (int i = 0; i < n; ++i){  
        // do something  
    }
```

Distributes loop to n threads on GPU.

→ 256 threads per thread block

→ usually $\text{ceil}(n/256)$ blocks in grid

```
#pragma acc parallel vector_length(256)
```

```
#pragma acc loop gang vector
```

```
    for (int i = 0; i < n; ++i){  
        for (int j = 0; j < m; ++j){  
            // do something  
        }  
    }
```

Distributes outer loop to n threads on GPU.

Each thread executes inner loop sequentially.

→ 256 threads per thread block

→ usually $\text{ceil}(n/256)$ blocks in grid

```
#pragma acc parallel vector_length(256)
```

```
#pragma acc loop gang
```

```
    for (int i = 0; i < n; ++i){
```

```
#pragma acc loop vector
```

```
        for (int j = 0; j < m; ++j){  
            // do something  
        }
```

```
    }
```

Distributes outer loop to GPU

multiprocessors (block-wise). Distributes

inner loop to threads within thread blocks.

→ 256 threads per thread block

→ usually n blocks in grid

■ Local memory/ Registers

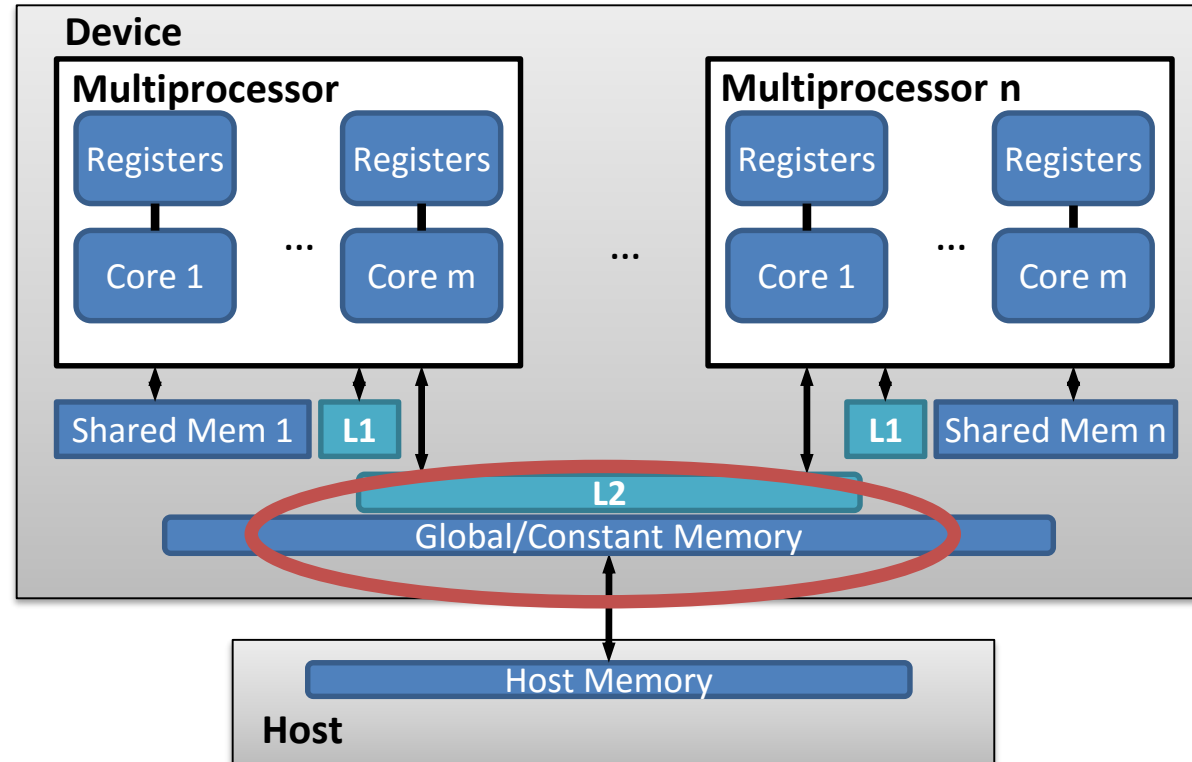
■ Shared memory/ L1

- Very low latency
(~100x than gmem)
- Bandwidth (aggregate):
1+ TB/s (Fermi),
2.5 TB/s (Kepler)

■ L2

■ Global memory

- High latency
(400-800 cycles)
- Bandwidth: 144 GB/s (Fermi),
250 GB/s (Kepler)



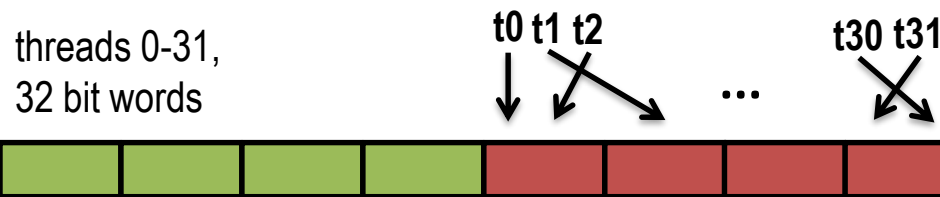
- **Stores:** Invalidate L1, write-back for L2

- **Loads:**

	Caching	Non-caching
Enabling by	default (Fermi)	Compiler-option CUDA: <code>-Xptxas-dlcm=cg</code> OpenACC PGI: <code>-Mx,180,8</code>
Attempt to hit:	L1 → L2 → gmem	L2 → gmem (no L1: invalidate line if it's already in L1)
Load granularity	128-byte line	32-byte line

Kepler uses L1 cache only for thread-private data → L2 cache with 128-byte cache line.

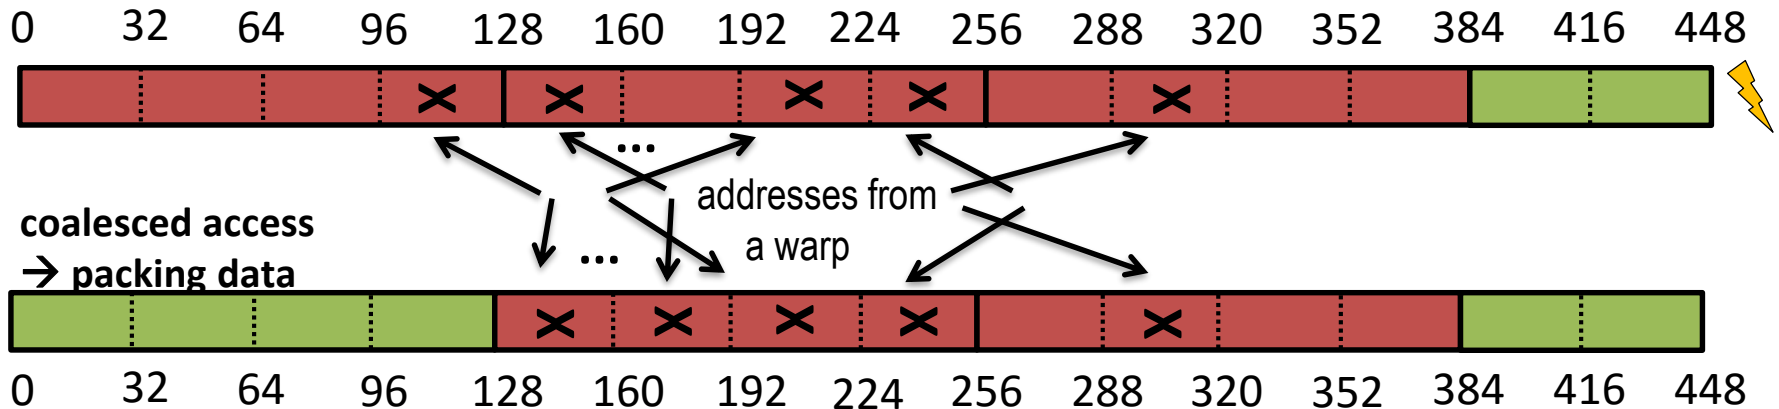
- Threads in a warp provide memory addresses
- Determine which lines/segments are needed
- Request the needed lines/segments



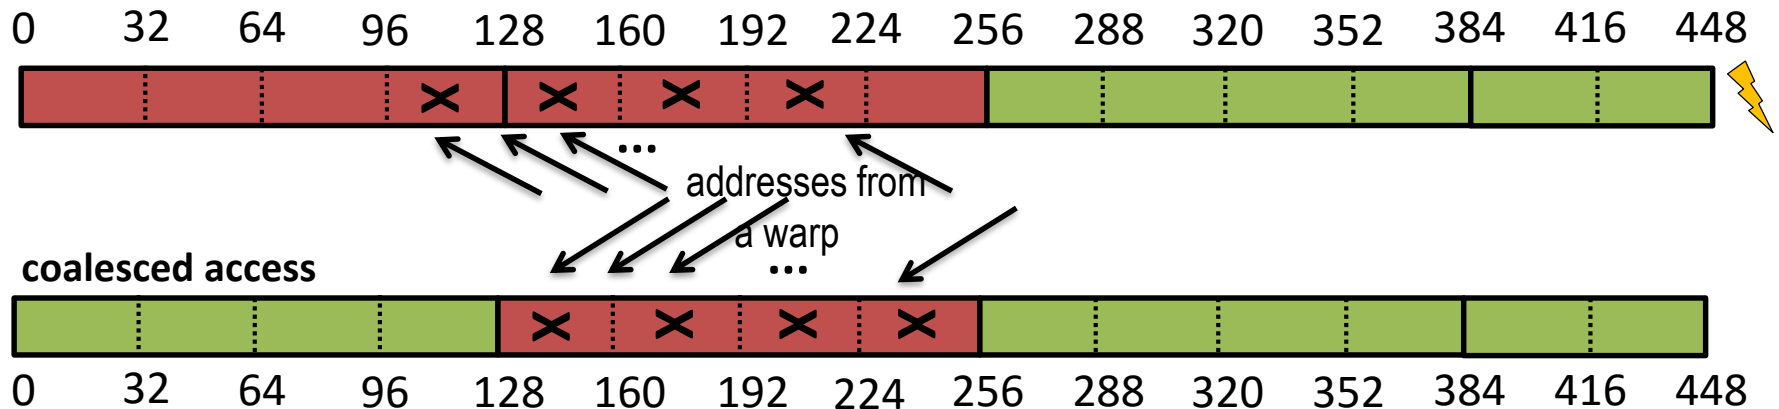
Gmem access per warp (32 threads)

Range of accesses

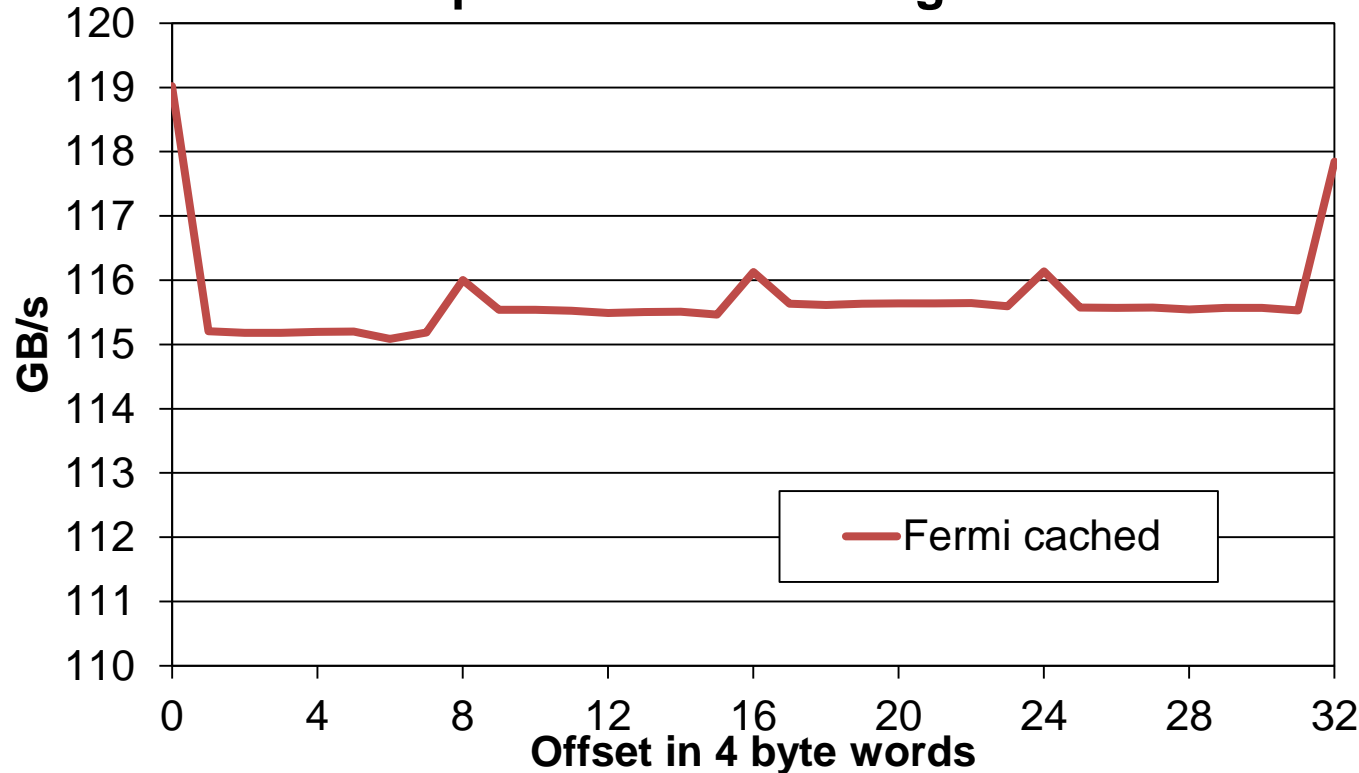
caching loads



Address alignment



Impact of Address Alignment



Example program:

- Copy ~67 MB of floats
- NVIDIA Tesla C2050 (Fermi)
- ECC off
- 256 threads/block

➔ Misaligned accesses can drop memory throughput

■ Example: AoS vs. SoA

Array of Structures (AoS)

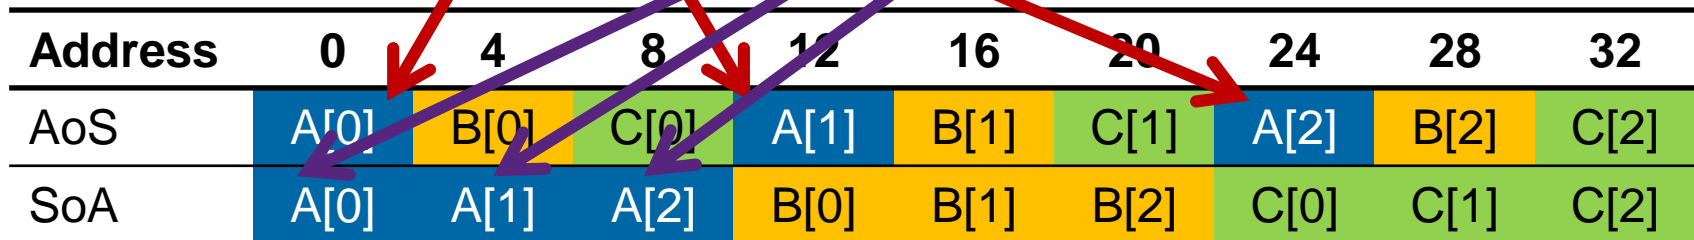
```
struct myStruct_t {  
    float a;  
    float b;  
    int c;  
}  
myStruct_t myData[];
```

Structure of Arrays (SoA)

```
struct {  
    float a[];  
    float b[];  
    int c[];  
}
```

```
#pragma acc kernels  
for (int i=0; i<n; i++) {  
    ... myData[i].a / myData.a[i] ...  
}
```

Address	0	4	8	12	16	20	24	28	32
AoS	A[0]	B[0]	C[0]	A[1]	B[1]	C[1]	A[2]	B[2]	C[2]
SoA	A[0]	A[1]	A[2]	B[0]	B[1]	B[2]	C[0]	C[1]	C[2]



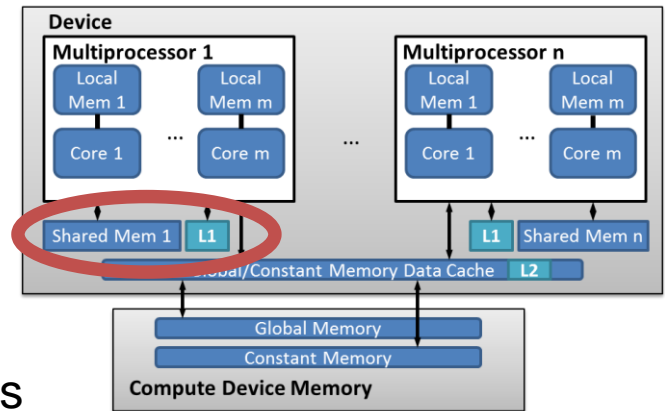
■ Smem per MP (10s of KB)

→ Inter-thread communication within a block

→ Synchronization to avoid race conditions

■ Low-latency, high-throughput memory

→ Cache data in smem to reduce gmem accesses

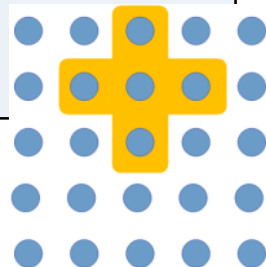


```
__global__ void myKernel(...) {  
    // Compute thread IDs x, y &  
    // mapping to row, col  
    __shared__ double as[BS][BS];  
    as[x][y] = a[row][col];  
    __syncthreads();  
    if (/* valid */) {  
        b[row][col] = (as[x][y-1] +  
            as[x][y+1] + as[x-1][y] +  
            as[x+1][y]) / 4;  
    }  
}
```

CUDA

```
#pragma acc kernels  
#pragma acc loop gang vector  
    for (int i = 1; i < N-1; ++i){  
#pragma acc loop gang vector  
    for (int j = 1; j < N-1; ++j){  
#pragma acc cache(a[i-1:3][j-1:3])  
        b[i][j] = (a[i][j-1] +  
            a[i][j+1] + a[i-1][j] +  
            a[i+1][j]) / 4;  
    }  
}
```

OpenACC



- Per thread block

- Static shared memory

 - Kernel code: `__shared__ int myarray[64];`

- Dynamic shared memory

 - Host code (specify shared memory size as third argument):

 - `myKernel<<<1, n, n*sizeof(int)>>>(...);`

 - Kernel code: `extern __shared__ int myarray[];`

- Synchronize threads (within thread block)

 - `__syncthreads();`

Cache construct

→ Prioritizes data for placement in the highest level of data cache on GPU

C/C++

```
#pragma acc cache (list)
```

Fortran

```
!$acc cache (list)
```

Sometimes the PGI compiler ignores the `cache` construct.
→ See compiler feedback

→ Use at the beginning of the loop or in front of the loop

Tile clause

→ Strip-mines each loop in a loop nest according to the values in *expr-list*

C/C++

```
#pragma acc loop [<schedule>] tile (expr-list)
```

Fortran

```
!$acc loop [<schedule>] tile (expr-list)
```

1. entry applies to innermost loop

- **Which impact does the PCIe have?**
- **What is branch divergence?**
 - Which performance impact does it have?
- **What can be a good launch configuration and why?**
 - How can the launch configuration be specified?
- **What can be done to saturate the bus?**
 - What is coalescing?
 - How can it be achieved?
 - What impact does AoS and SoA have?

1. Why supercomputers?
2. Modern processors
3. Basic optimization techniques for serial code
4. Data access optimization
5. Parallel computers
6. Energy efficiency
7. Parallelization and optimization strategies
8. Parallel algorithms
9. Distributed-memory programming with MPI
10. Shared-memory programming with OpenMP
11. Hybrid programming (MPI + OpenMP)

12. Heterogeneous architectures (GPUs, Xeon Phis)

- Motivation
- RWTH Environment
- GPGPU – Basics
 - Overview
 - GPU Architecture

- Programming Model
- Execution Model
- Memory Model
- Misc on OpenACC
- Summary
- GPGPU – Advanced
 - Kernel optimizations
 - GPU-CPU interaction
- **Intel Xeon Phi**
 - **Motivation**
 - KNL Architecture
 - Vectorization
- Links

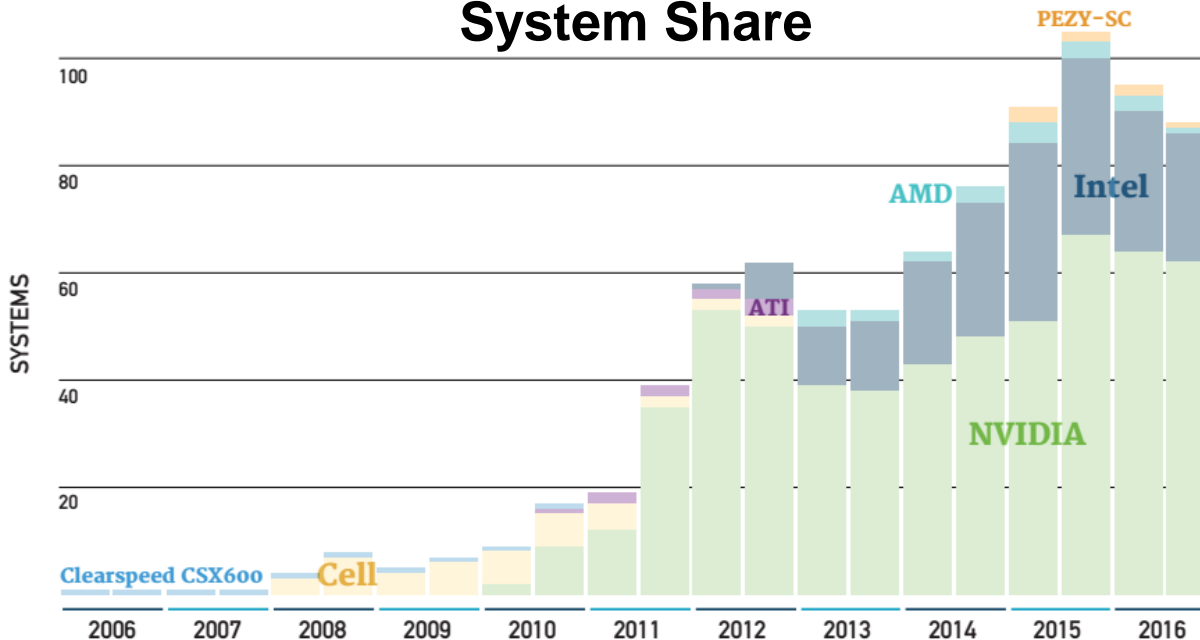
■ Accelerators/ co-processors

- GPGPUs (e.g. NVIDIA, AMD)
- Intel Many Integrated Core (MIC) Arch. (Intel Xeon Phi)
- FPGAs (e.g. Convey), ...

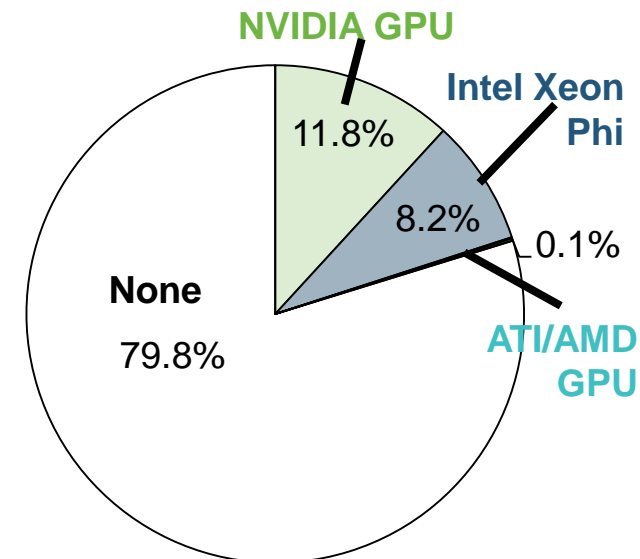
■ Heterogeneous Arch.

- Combination of commodity processors & accelerators

System Share



Performance Share



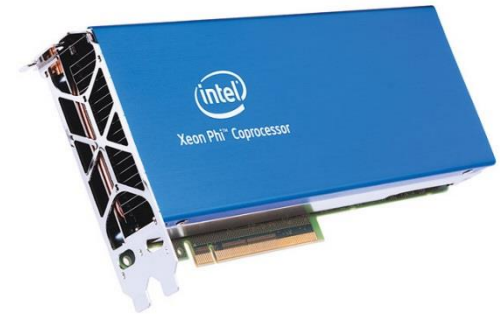
■ Knight's Ferry (engineering sample) (2010)

- 32 x86-based cores, 128 threads
- 1.2 GHz, 512-bit SIMD
- 1-2 GB GDDR5



■ Knight's Corner (2013)

- 61 x86-based cores, 240 threads
- 1.2 GHz, 512-bit SIMD, ~ 1 TFLOPS peak
- 8 or 16 GB GDDR5



■ Knight's Landing (2016)

- 72 x86-based cores, 288 threads
- 1.5 GHz, 2x 512-bit SIMD, ~ 3.5 TFLOPS peak
- 16 GB MCDRAM + 384 GB DDR4

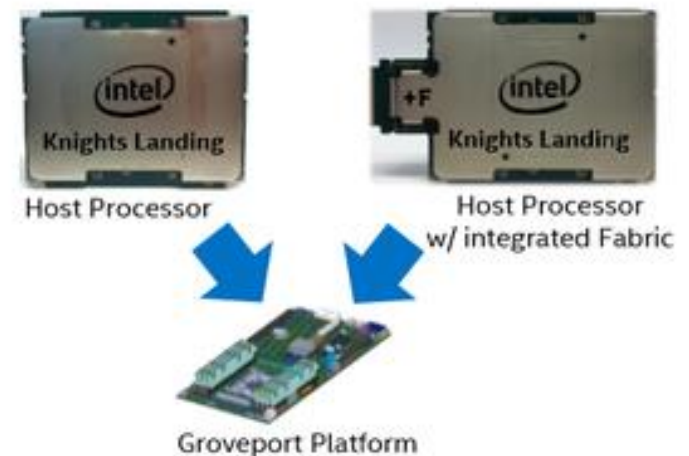


■ Xeon Phi Processor

- binary compatible with x86-based Servers
- use MPI, OpenMP, ...
- no need for Offloading
- no PCI bottleneck
- OPA Fabric on-package available

■ Important for applications

- to use SIMD operations
- to express a lot of parallelism
- make efficient use of the MCDRAM



1. Why supercomputers?
2. Modern processors
3. Basic optimization techniques for serial code
4. Data access optimization
5. Parallel computers
6. Energy efficiency
7. Parallelization and optimization strategies
8. Parallel algorithms
9. Distributed-memory programming with MPI
10. Shared-memory programming with OpenMP
11. Hybrid programming (MPI + OpenMP)

12. Heterogeneous architectures (GPUs, Xeon Phis)

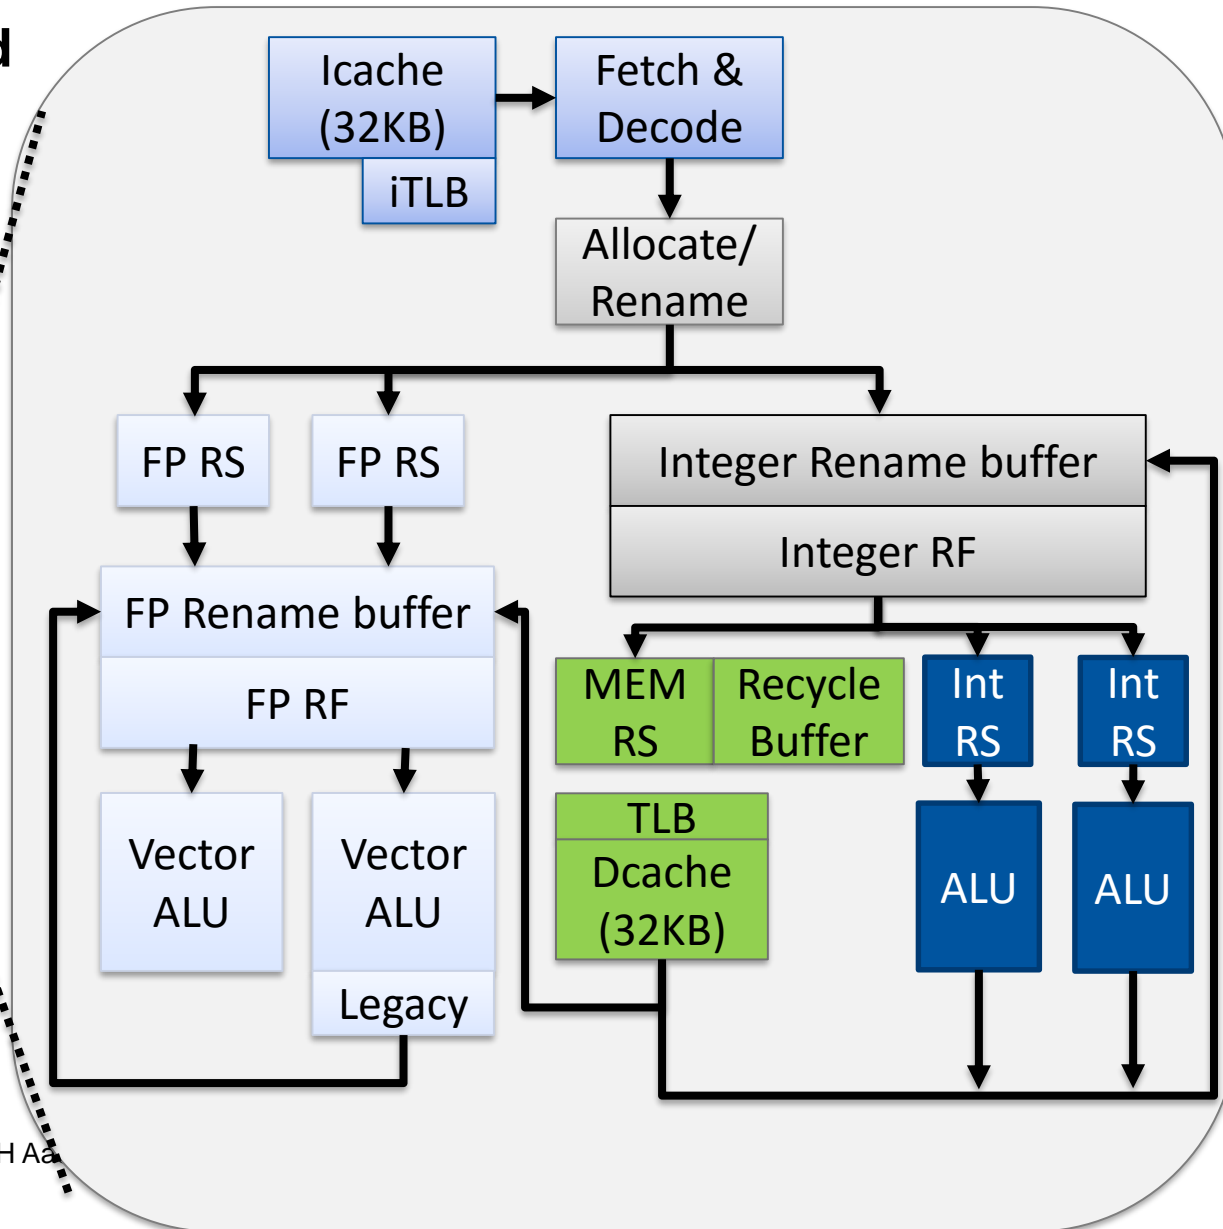
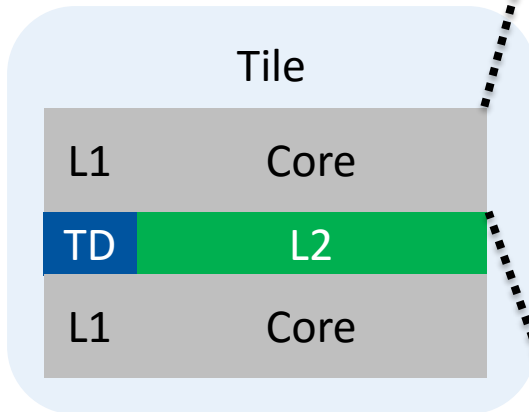
- Motivation
- RWTH Environment
- GPGPU – Basics
 - Overview
 - GPU Architecture

- Programming Model
- Execution Model
- Memory Model
- Misc on OpenACC
- Summary
- GPGPU – Advanced
 - Kernel optimizations
 - GPU-CPU interaction
- **Intel Xeon Phi**
 - Motivation
 - **KNL Architecture**
 - Vectorization
- Links

Knights Landing Tile



- Intel Atom Core based
- Out-of-order execution
- 2 x 512-bit VPU's per core
- 2 cores per tile
- 1 MB L2 cache



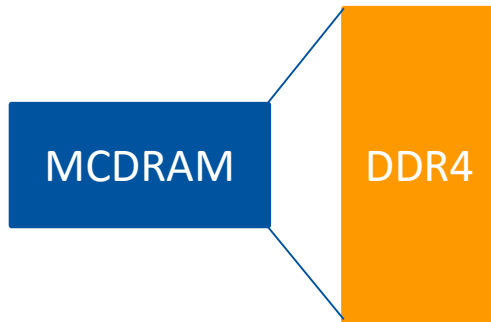
Knights Landing Overview



- up to 36 tiles
- 2D mesh interconnect
- 4 MCDRAM on-chip
- DDR4 off-chip



Cache mode



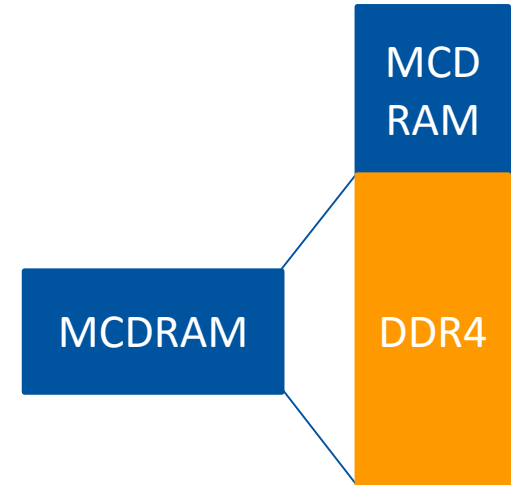
- automatic caching to MCDRAM
- easy to use
- everything is cached
- caching overhead

Flat mode



- MCDRAM shown as NUMA node
- explicit allocation needed
- user decides which data is in MCDRAM

Hybrid mode

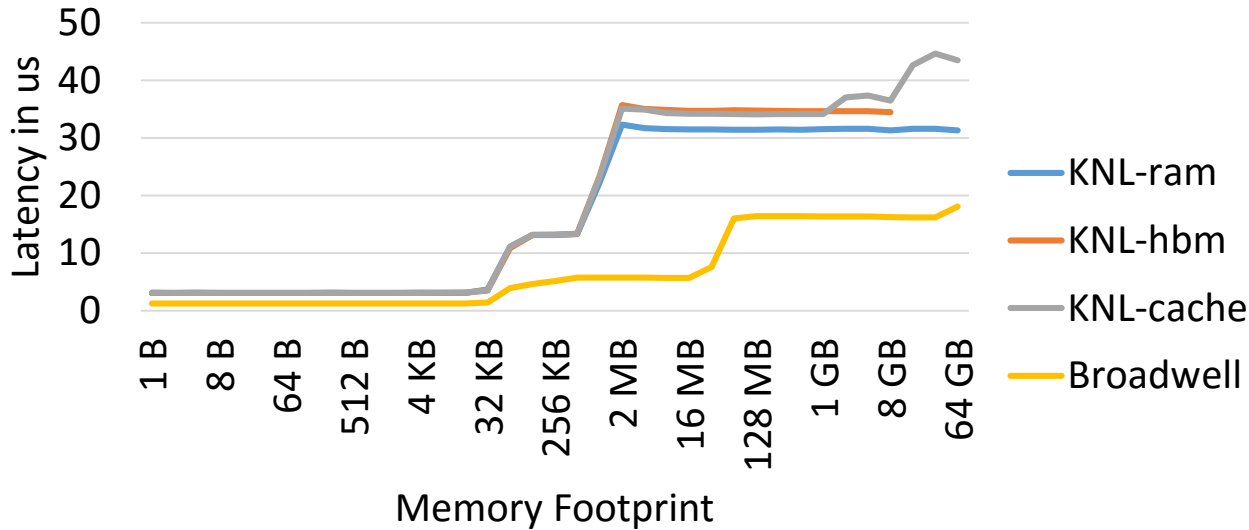


- mixture of both modes
- 8/8 GB or 4/12 GB splitting

numactl - command line tool to investigate NUMA topologies

- `$ numactl --hardware` - prints information on NUMA nodes in the system
- `$ numactl --show` - prints information on available resources for the process

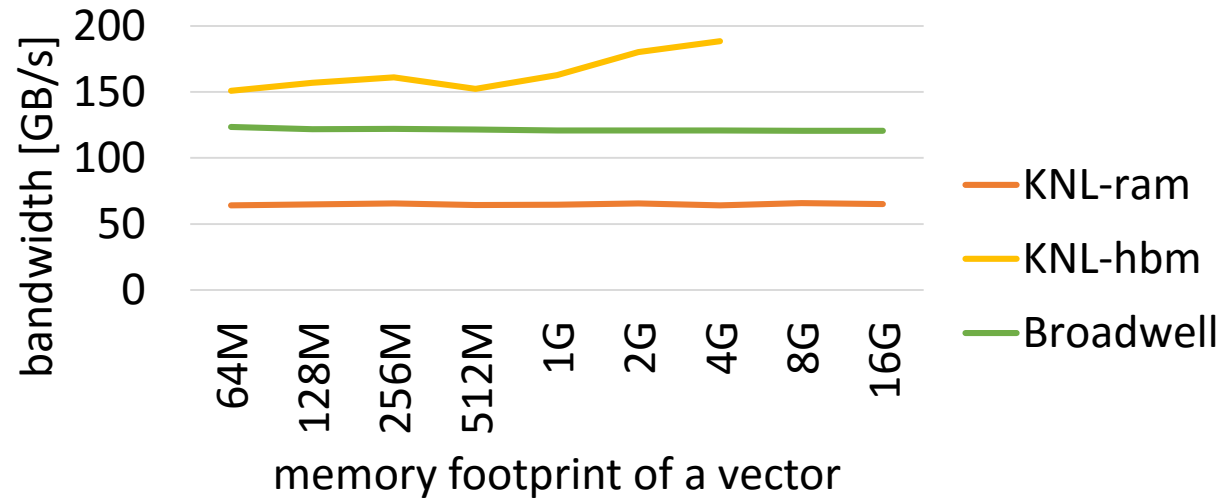
```
$ numactl --hardware
available: 2 nodes (0-1)
node 0 cpus: 0 .. 255
node 0 size: 98178 MB
node 0 free: 94138 MB
node 1 cpus:
node 1 size: 16384 MB
node 1 free: 15932 MB
node distances:
node 0 1
0: 10 31
1: 31 10
```



latency of MCDRAM is slightly higher than DDR4

bandwidth of MCDRAM higher than DDR4

close to 200 GB/s



numactl - command line tool to investigate and handle NUMA under Linux

- `$ numactl --cpunodebind 0,1,2 ./a.out`

- only use cores of NUMA node 0-2 to execute a.out

- `$ numactl --physcpubind 0-17 ./a.out`

- only use cores 0-17 to execute a.out

- `$ numactl --membind 0,3 ./a.out`

- only use memory of NUMA node 0 and 3 to execute a.out

libnuma - library for NUMA control
(include numa.h and link -lnuma)

- `void *numa_alloc_local(size_t size);`

→ allocate memory on the local NUMA node

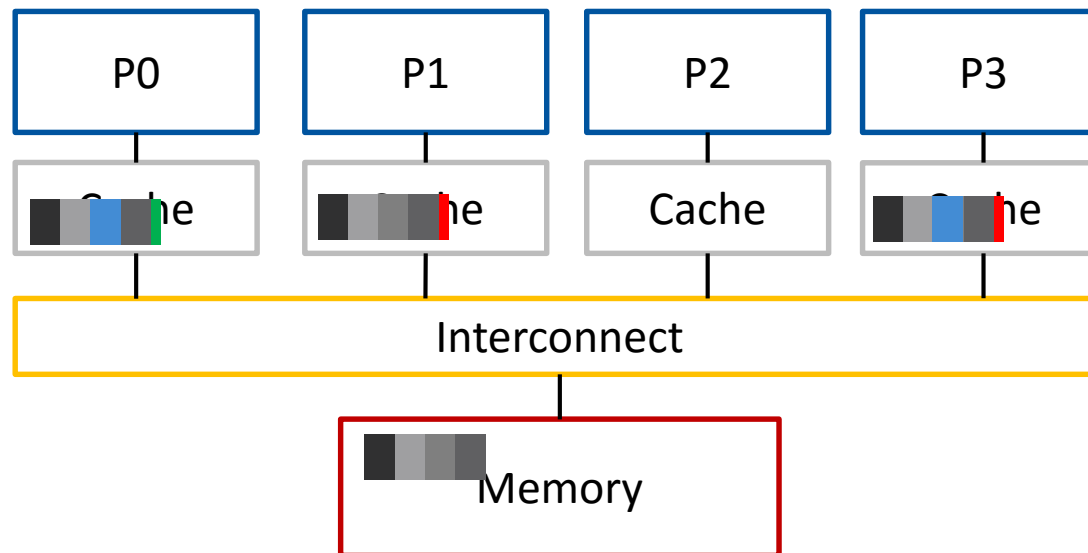
- `void *numa_alloc_onnode(size_t size, int node);`

→ allocate memory on NUMA node node

- `int numa_move_pages(int pid, unsigned long count, void **pages, const int *nodes, int *status, int flags);`

→ migrate memory pages at runtime to different NUMA nodes

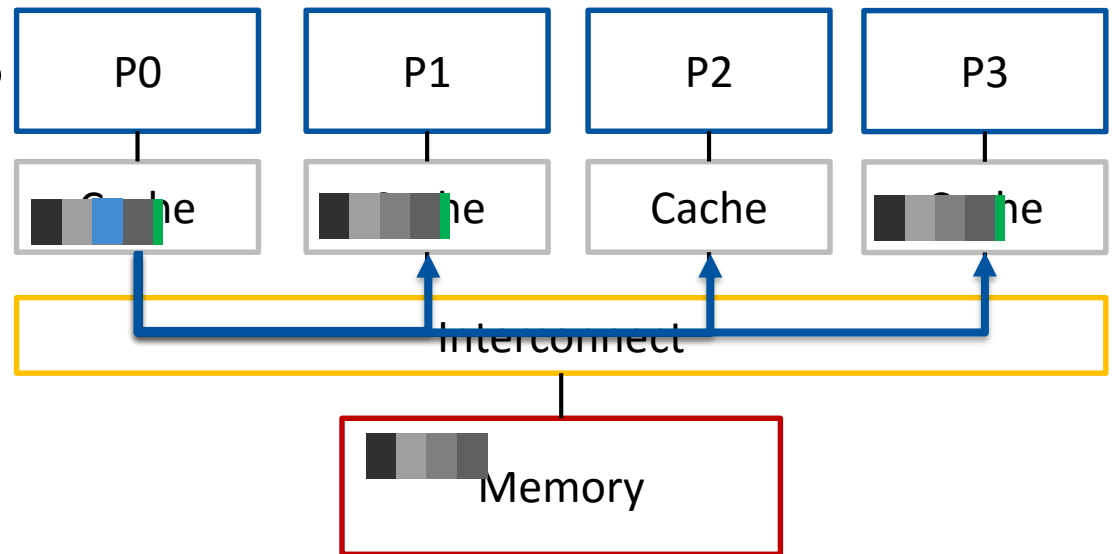
1. Cache-line in Memory
2. read by P1 -> copy CL in Cache of P1
3. read by P3 -> copy CL in Cache of P3
4. read by P0 -> copy CL in Cache of P0
5. write by P0 -> P0 changes local CL; other CLs marked invalid
6. read by P3 -> local line is invalid, get copy from P0



Task: Inform every cache holding the cache line to invalidate it

Solution 1: Snooping

- send write information to processes
- always check for write information and check the local cache to invalidate cache lines



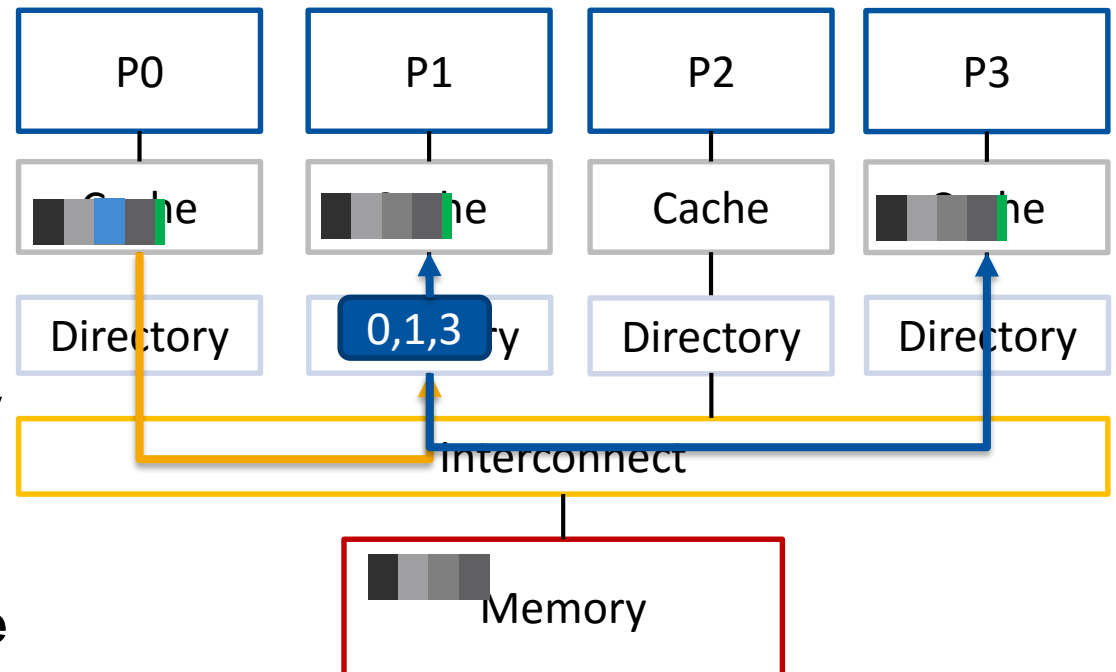
Pro: Fast, all messages at once

Con: Lots of traffic for many processes

Task: Inform every cache holding the cache line to invalidate it

Solution 2: Directories

- every cache line belongs to a home directory
- store all Ps with a copy
- on a write inform the home agent
- home agent informs only Ps with a CL



Pro: reduces traffic on large machines

Con: Slower, one more hop

Quadrant mode

1. miss in local L2\$
2. send request to TD
3. forward request to memory
4. send cache-line to requesting core

TD: random distribution

DDR: same quadrant as TD



Sub NUMA Clusters:

1. miss in local L2\$
2. send request to TD
3. forward request to memory
4. send cache-line to requesting core

TD: quadrants shown as
NUMA-nodes

DDR: same quadrant as TD

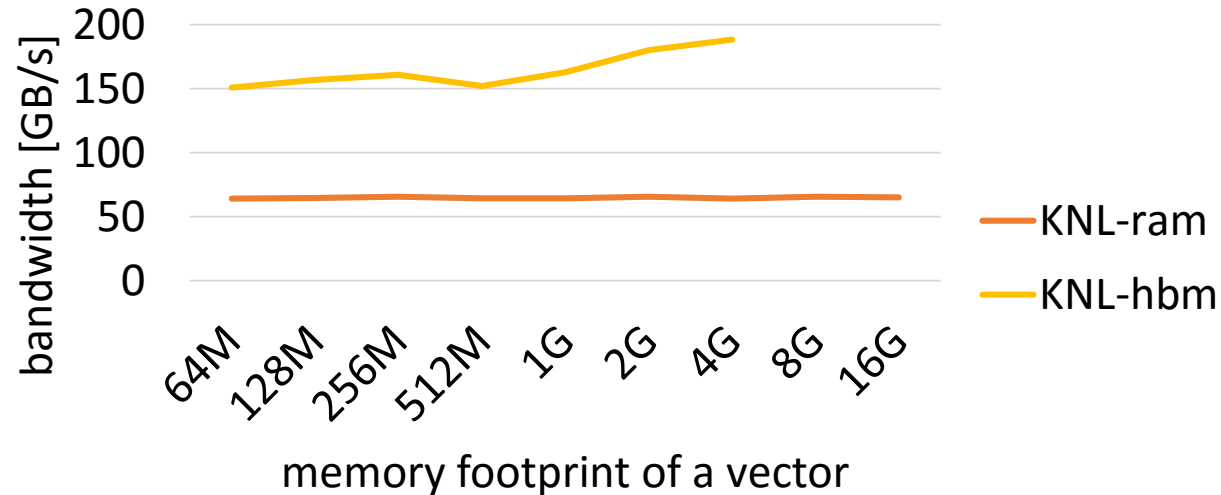


Sub NUMA Clustering

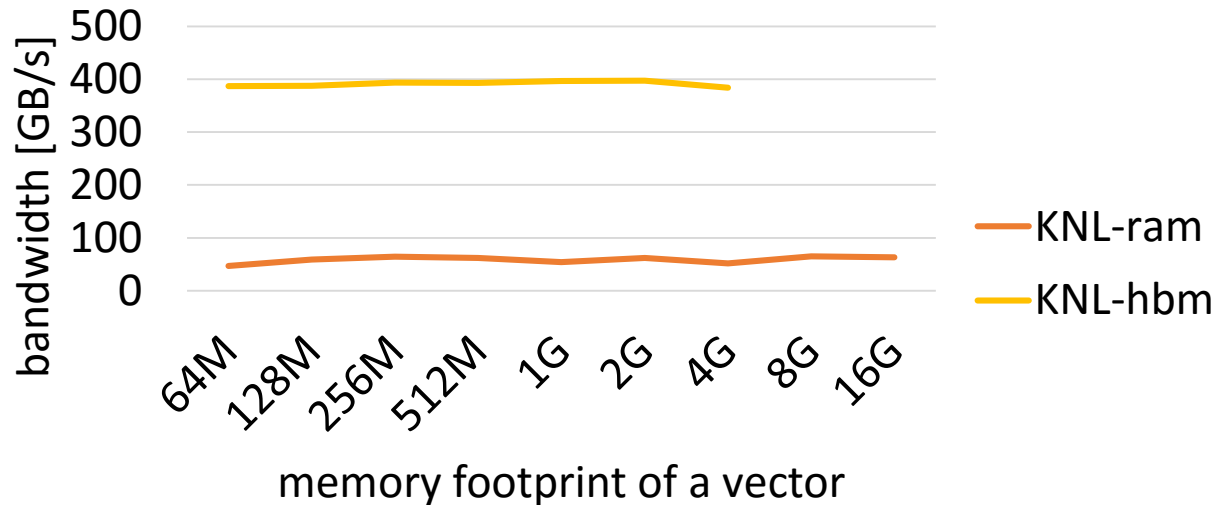


```
$ numactl --hardware
available: 8 nodes (0-7)
node 0 cpus: 0-15,64-79, 128-143, 192-207
node 0 size: 24450 MB
node 0 free: 21956 MB
node 1 cpus: 16-31, 80-95, 144-159, 208-223
node 1 size: 24576 MB
node 1 free: 23266 MB
node 2 cpus: 32-47, 96-111, 160-175, 224-239
node 2 size: 24576 MB
node 2 free: 23306 MB
node 3 cpus: 48-63, 112-127, 176-191, 240-255
node 3 size: 24576 MB
node 3 free: 23590 MB
node 4 cpus:
node 4 size: 4096 MB
node 4 free: 3981 MB
node 5 cpus:
node 5 size: 4096 MB
node 5 free: 3984 MB
node 6 cpus:
node 6 size: 4096 MB
node 6 free: 3981 MB
node 7 cpus:
node 7 size: 4096 MB
node 7 free: 3984 MB
node distances:
node 0 1 2 3 4 5 6 7
0: 10 21 21 21 31 41 41 41
1: 21 10 21 21 41 31 41 41
2: 21 21 10 21 41 41 31 41
3: 21 21 21 10 41 41 41 31
4: 31 41 41 41 10 41 41 41
5: 41 31 41 41 41 10 41 41
6: 41 41 31 41 41 41 10 41
7: 41 41 41 31 41 41 41 10
```

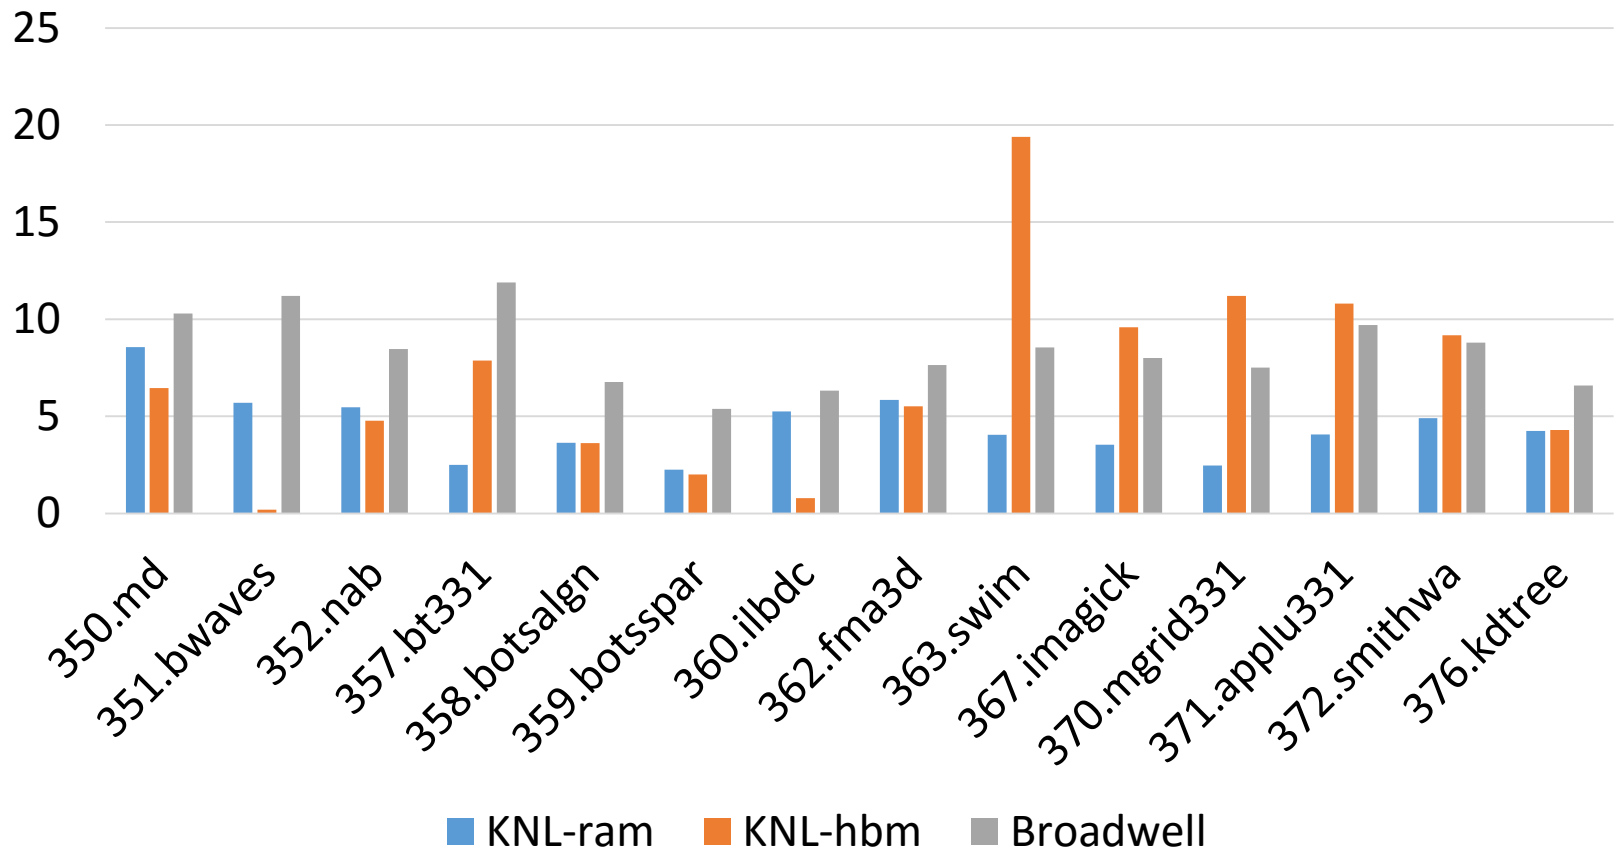
Quadrant mode



SNC mode



- SPEC-OMP benchmarks on KNL and Broadwell
- no code optimization done so far



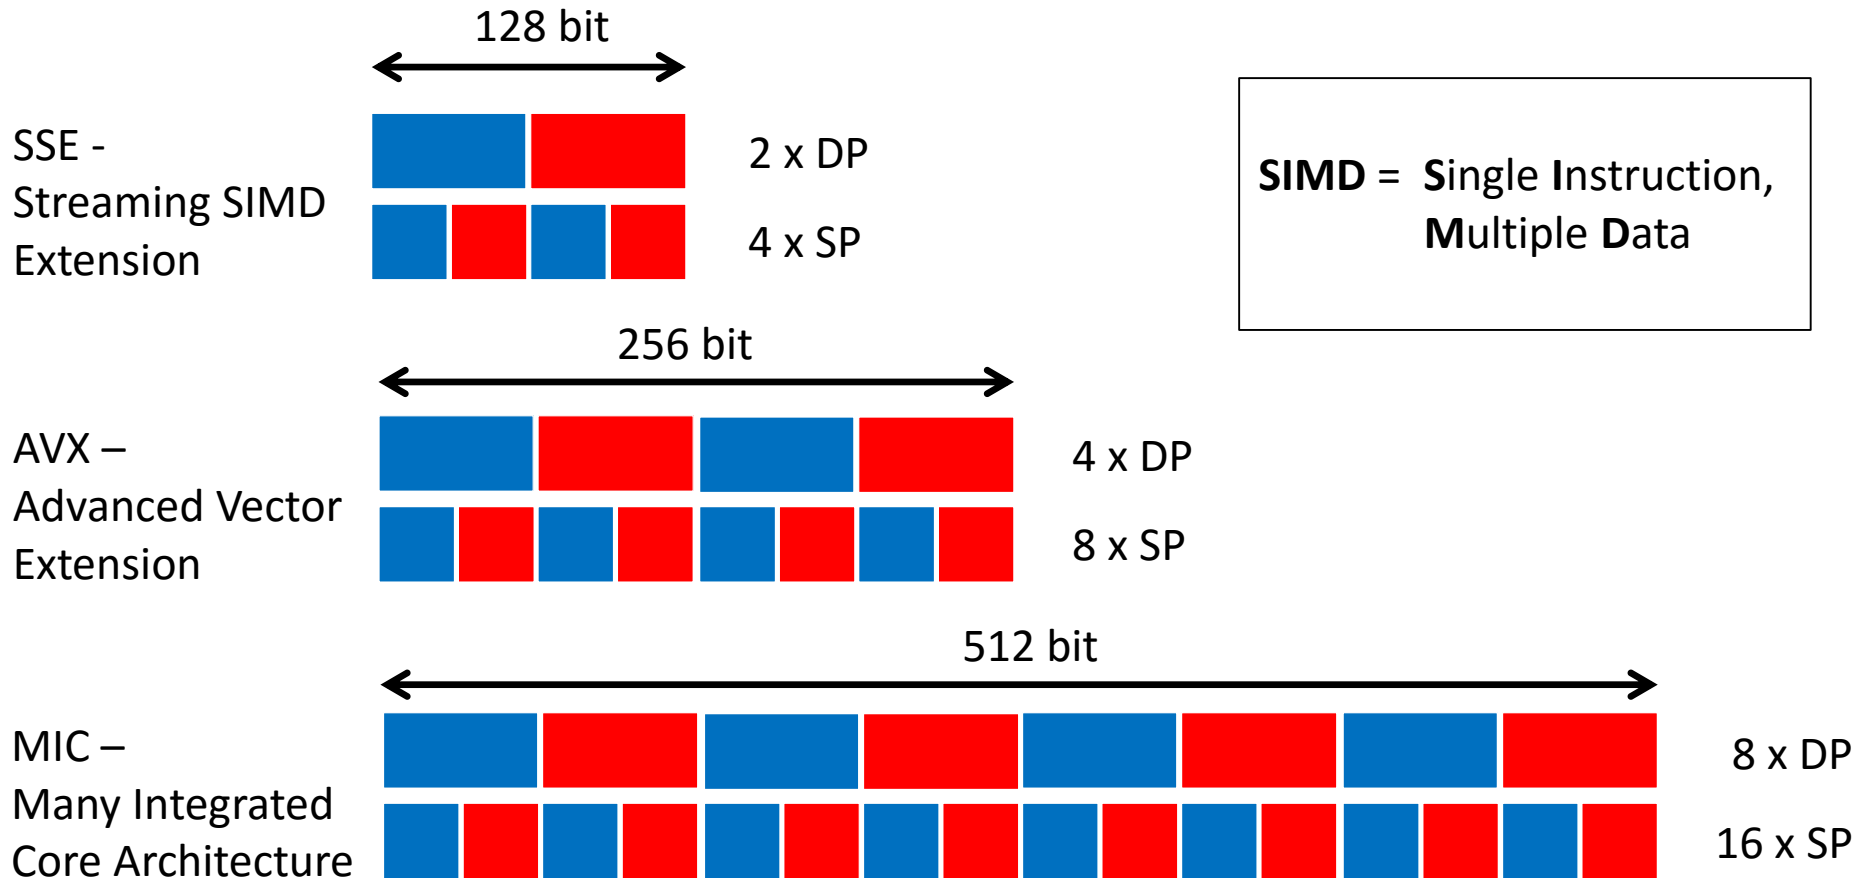
1. Why supercomputers?
2. Modern processors
3. Basic optimization techniques for serial code
4. Data access optimization
5. Parallel computers
6. Energy efficiency
7. Parallelization and optimization strategies
8. Parallel algorithms
9. Distributed-memory programming with MPI
10. Shared-memory programming with OpenMP
11. Hybrid programming (MPI + OpenMP)

12. Heterogeneous architectures (GPUs, Xeon Phis)

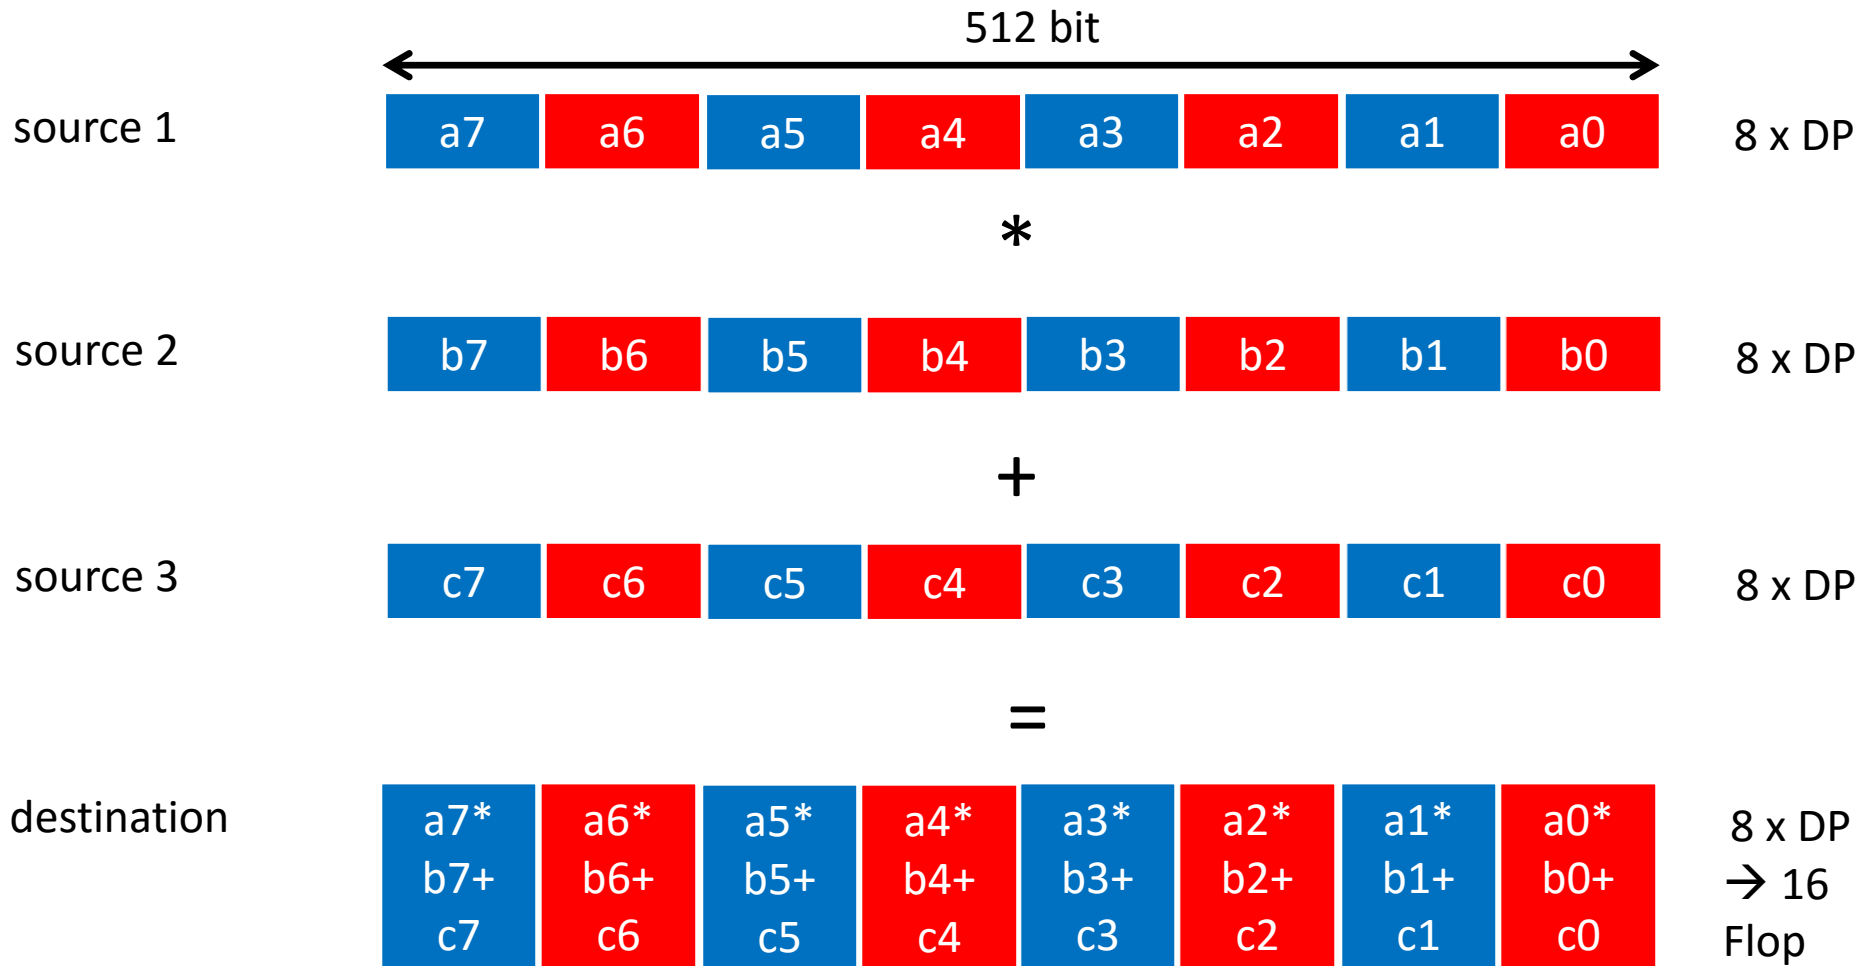
- Motivation
- RWTH Environment
- GPGPU – Basics
 - Overview
 - GPU Architecture

- Programming Model
- Execution Model
- Memory Model
- Misc on OpenACC
- Summary
- GPGPU – Advanced
 - Kernel optimizations
 - GPU-CPU interaction
- **Intel Xeon Phi**
 - Motivation
 - KNL Architecture
 - **Vectorization**
- Links

SIMD Vector Capabilities



SIMD Fused Multiply Add

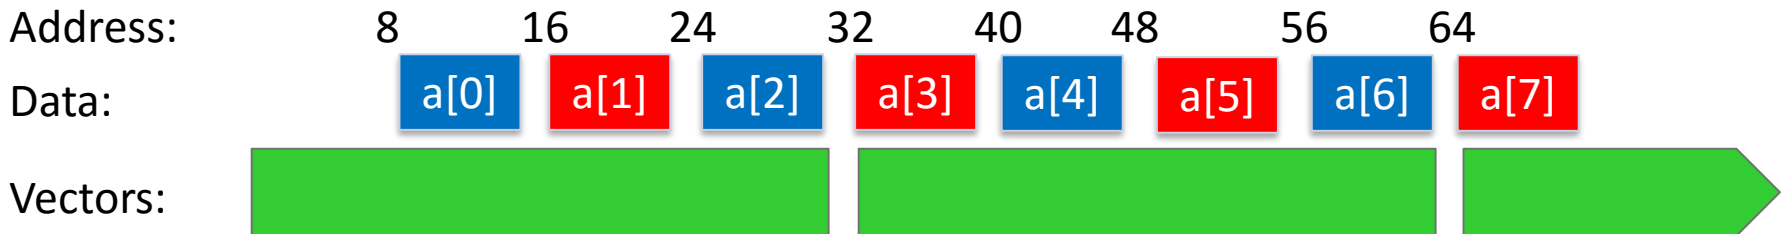


■ Vectorization works best on aligned data structures.

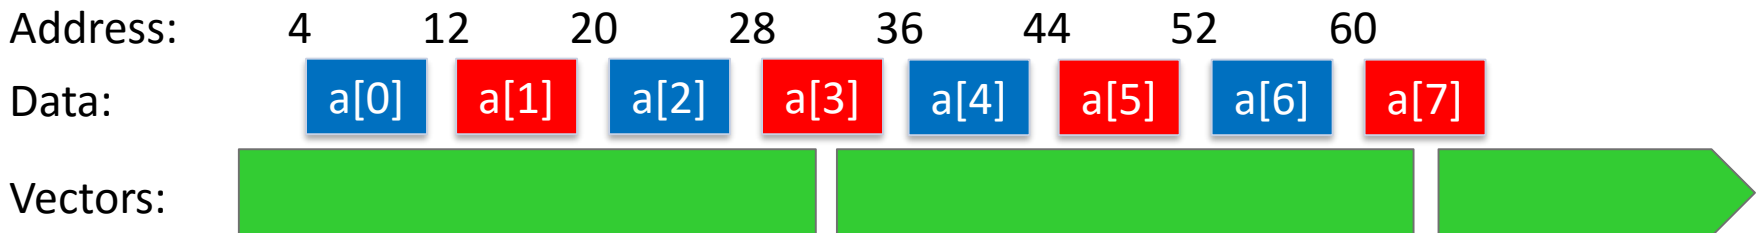
Good alignment



Bad alignment



Very bad alignment



■ Intel auto-vectorizer

- Combination of loop unrolling and SIMD instructions to get vectorized loops
- No guarantee given, compiler might need some hints

■ C/C++ aliasing: use `restricted` keyword

■ Compiler feedback

- Use `-vec-report [n]` to control the diagnostic information of the vectorizer
- concentrate on hotspots for optimization

Example: `-vec-report=3`

driver.c(96): (col. 3) remark: loop was not vectorized: not inner loop

driver.c(64): (col. 5) remark: loop was not vectorized: existence of vector dependence

driver.c(68): (col. 7) remark: vector dependence: assumed FLOW dependence between f line 68 and dx line 65

driver.c(65): (col. 7) remark: vector dependence: assumed ANTI dependence between dx line 65 and f line 68

driver.c(68): (col. 7) remark: vector dependence: assumed FLOW dependence between f line 68 and dy line 66

■ Vectorize a loop nest

→ Cut loop into chunks that fit a SIMD vector register

→ No parallelization of the loop body

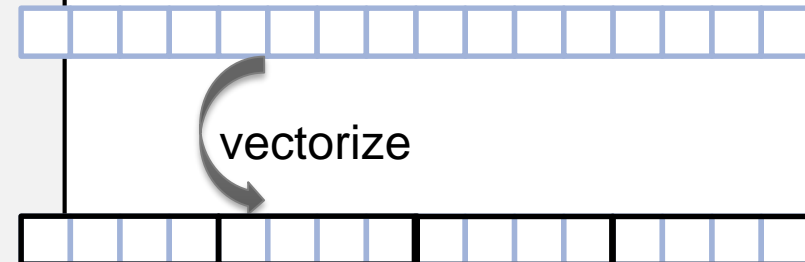
C/C++

```
#pragma omp simd [clause[[,] clause],...]  
for-loops
```

Fortran

```
!$omp simd [clause[[,] clause],...]  
do-loops
```

```
void sprod(float *a, float *b, int n)  
{  
    float sum = 0.0f;  
    #pragma omp simd reduction(+:sum)  
    for (int k=0; k<n; k++)  
        sum += a[k] * b[k];  
    return sum;  
}
```



Source: Michael Klemm, Intel

■ Parallelize & vectorize loop nest

C/C++

```
#pragma omp for simd [clause[,] clause],...  
for-loops
```

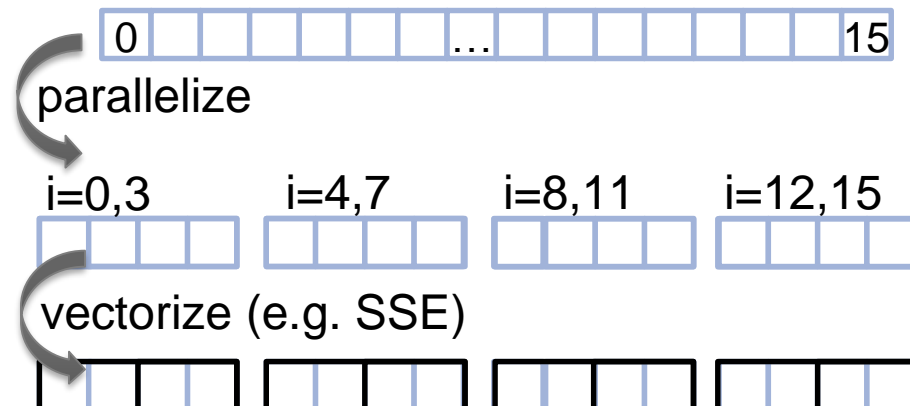
1. Parallelization

→ Distribution of iterations across implicit tasks of Parallel Region

2. Vectorization

→ Conversion of chunks of iterations to SIMD loops

```
void vecAdd(double *a, double *b,  
            double *c, int n)  
{  
    #pragma omp for simd  
    for (int i=0; i<n; i++)  
        c[i] = a[i] * b[i];  
}
```



■ Use Structure of Arrays (SoA) instead of Array of Structures (AoS)

→ Color structure

```
struct Color{ //AoS
    float r;
    float g;
    float b;
}
Color* c;
```



```
struct Colors{ //SoA
    float* r;
    float* g;
    float* b;
}
```



■ Detailed information: Intel Vectorization Guide

<http://software.intel.com/en-us/articles/a-guide-to-auto-vectorization-with-intel-c-compilers/>

■ **How does a Knights Landing look like?**

- How does the memory layout look like?
- In which modes can it be operated?
- How many threads/ vector-widths are available?

■ **Which optimization strategies should be applied?**

- Which impact does vectorization have?
- How can vectorization be achieved?

1. Why supercomputers?
2. Modern processors
3. Basic optimization techniques for serial code
4. Data access optimization
5. Parallel computers
6. Energy efficiency
7. Parallelization and optimization strategies
8. Parallel algorithms
9. Distributed-memory programming with MPI
10. Shared-memory programming with OpenMP
11. Hybrid programming (MPI + OpenMP)

12. Heterogeneous architectures (GPUs, Xeon Phis)

- Motivation
- RWTH Environment
- GPGPU – Basics
 - Overview
 - GPU Architecture

- Programming Model
- Execution Model
- Memory Model
- Misc on OpenACC
- Summary
- GPGPU – Advanced
 - Kernel optimizations
 - GPU-CPU interaction
- Intel Xeon Phi
 - Motivation
 - KNL Architecture
 - Vectorization
- **Links**

- **GPGPU Community:** <http://gpgpu.org/>
- **GPU Computing Community:** <http://gpucomputing.net/>
- **GPU Science:** <http://gpuscience.com/>
- **GPU Technology Conference (GTC) On-Demand**
<http://www.gputechconf.com/gtcnew/on-demand-GTC.php>

Collection of presentations given in the context of different GPU conferences & workshops

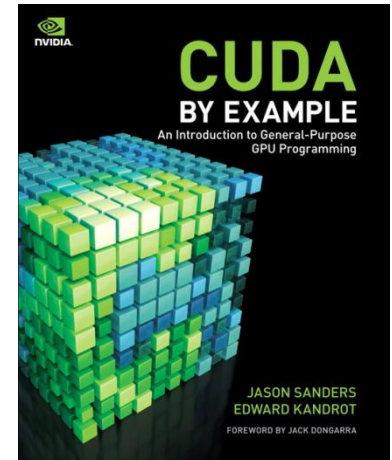
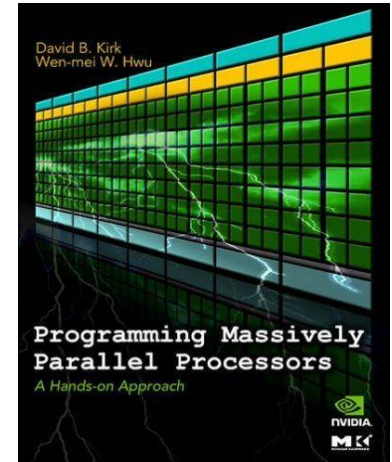
■ Webinars

Videos on different GPU topics

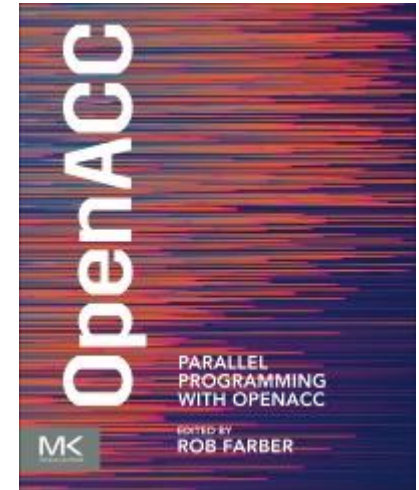
<https://developer.nvidia.com/gpu-computing-webinars>

- **Nvidia CUDA Zone (Toolkit, Profiler, SDK, documentation,...):**
http://www.nvidia.com/object/cuda_home_new.html
 - CUDA Programming Guide
 - CUDA Best Practice Guide
- **PGI's CUDA Fortran:**
<http://www.pgroup.com/resources/cudafortran.htm>
- **PGI's CUDA-x86:** <http://www.pgroup.com/resources/cuda-x86.htm>
- **Volkov2010:** <http://www.cs.berkeley.edu/~volkov/volkov10-GTC.pdf>

- David Kirk und Wen-Mei W. Hwu: *Programming Massively Parallel Processors – A Hands-on Approach* (2010)
- Jason Sanders und Edward Kandrot: *CUDA by Example: An Introduction to General-Purpose GPU Programming* (2010)



- **Rob Farber: Parallel Programming with OpenACC, 1st Edition, Morgan Kaufmann (2016)**
- **OpenACC Home, Specification, Quick Reference Card:** www.openacc.org/
- **OpenACC@PGI:** <http://www.pgroup.com/resources/accel.htm>
 - Tutorial presentations and articles, e.g. “OpenACC Features in the PGI C Compiler”
 - Tutorial videos and webinars
 - All kind of articles: <http://www.pgroup.com/resources/articles.htm>
- **OpenACC@NVIDIA:** <https://developer.nvidia.com/openacc>
 - Blog & Examples: <https://developer.nvidia.com/blog/tags/4621>



■ Dr. Dobb's

→ <http://www.drdobbs.com/parallel/easy-gpu-parallelism-with-openacc/240001776>

→ <http://www.drdobbs.com/parallel/the-openacc-execution-model/240006334>

■ Videos/ Webinars

→ <http://developer.nvidia.com/gpu-computing-webinars>

→ OpenACC and the new PGI Accelerator

→ Getting the most of OpenACC directives – Optimization with PGI Accelerator

→ <http://www.openacc.org/video>

■ Code examples

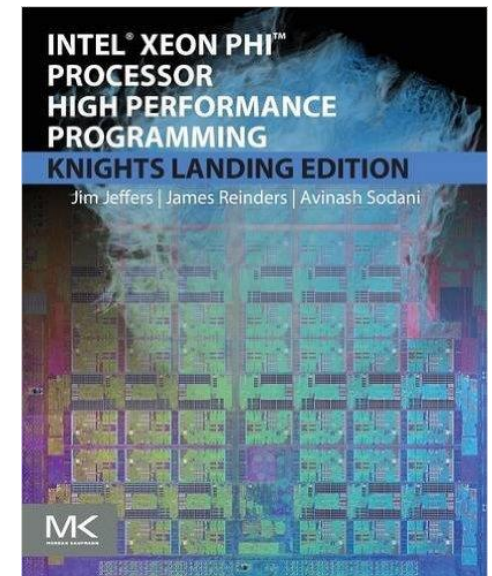
→ Case Studies: <http://www.openacc.org/Downloads>

■ Intel

→ Developer: <http://software.intel.com/en-us/mic-developer>

■ Book

→ J. Jeffers, J. Reinders, A. Sodani: Intel Xeon Phi Processor High Performance Programming (March 2016)



■ GPU cluster

- 28 nodes with each
2 NVIDIA Quadro 6000 GPUs (Fermi)

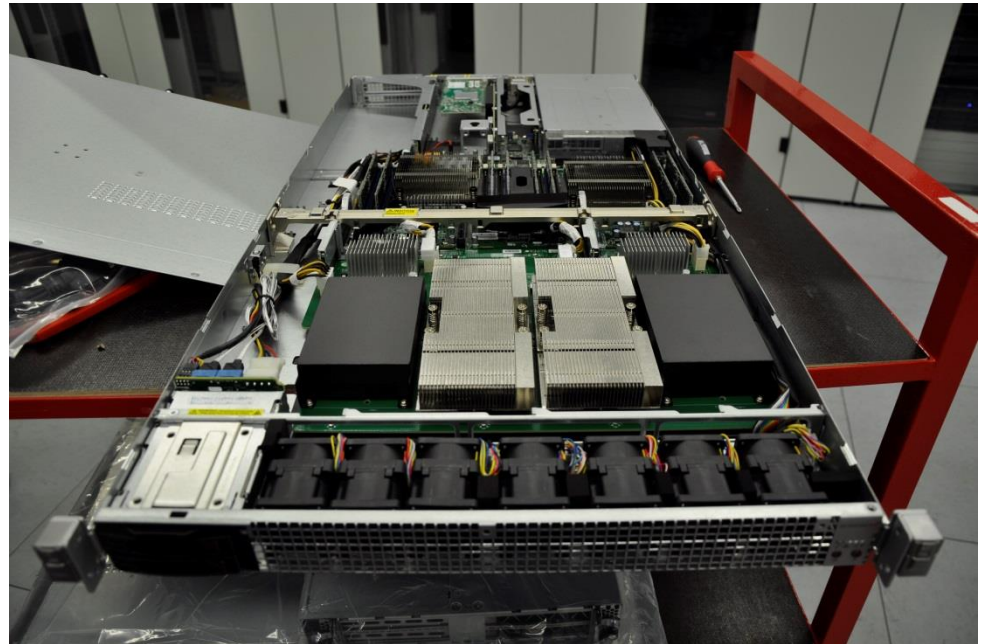
→ VR + HPC

- 2 nodes with each
2 NVIDIA K20X GPUs (Kepler)

- 10 nodes with each
2 NVIDIA P100 SXM2 (Pascal)
with 16GB



aixCAVE, VR, RWTH Aachen, since June 2012



		4 dialogue nodes	24 rendering nodes	1 head node
Name		linuxgpud[1-4]	linuxgpus[01-24]	linuxgpum1
Devices	#	2		1
	details/GPU	NVIDIA Quadro 6000 (Fermi) 448 cores 1.15 GHz 6 GB RAM ECC on max. GFlops: 1030.4 (SP), 515.2 (DP)		
Host	processor	2 x Intel Xeon X5650 EP (Westmere) (12-core CPU) @ 2.67GHz		
	RAM	24 GB		48 GB
Network		QDR InfiniBand		

■ Environment as on compute cluster (modules,...)

■ CUDA Toolkit: 7.5

→ CUDA

→ OpenCL

```
module load cuda  
→ directory: $CUDA_ROOT
```

■ PGI Compiler

→ CUDA Fortran

→ PGI OpenACC

```
module load pgi  
(module switch intel pgi)
```

■ CUDA Debugging

→ TotalView

→ DDT

→ Eclipse

```
module load totalview  
module load ddt
```

■ **Doc:** <https://doc.itc.rwth-aachen.de/display/CC/GPU+cluster>

■ **Interactive mode**

→ Short runs/tests only, debugging

→ 2 dialogue nodes (`linuxgpud1`, `linuxnvc01`): 24/7

→ 2 dialogue nodes (`linuxgpud[2,3]`): Mon – Fri, 8am – 8pm

■ **Batch mode**

→ No interaction, commands are queued + scheduled

→ For performance tests, long runs

→ 24 rendering nodes
4 dialogue nodes } Mon – Fri, 8pm – 8am; Sat + Son, whole day

→ 2 dialogue nodes (`linuxgpud4`, `linuxnvc02`): Mon – Fri, 8am – 8pm
→ for short test runs during daytime

■ **Note: reboot at switch from interactive to batch mode**

■ Jump from frontend cluster node to GPU dialogue node:

```
ssh -Y linuxgpud[1-3]
```

■ GPUs are set to “exclusive mode” (per process)

- Only one person can access GPU
- If occupied, e.g. message “all CUDA-capable devices are busy or unavailable”
- If not set a certain device in (CUDA) code, automatically scheduled to other GPU within node (if available)

■ Debugging

- Be aware: debugger run usually on GPU with ID 0 (fails if GPU is occupied)
- Use `CUDA_VISIBLE_DEVICES=1` to use other GPU

GPU Cluster - See what is running: nvidia-smi



```
linuxgpud1$> nvidia-smi  
Mon Oct 17 12:41:01 2011
```

nvidia-smi -q
Lists GPU details

ECC (SB: single bit,
DB: double bit)

NVIDIA-SMI 2.285.05 Driver Version: 285.05.09									
<hr/>									
Nb.	Name	GPU ID + type		Bus Id	display mode		Volatile	ECC SB / DB	
Fan	Temp	Power Usage /Cap		Memory Usage			GPU Util.	Compute M.	
<hr/>									
0.	Quadro 6000			0000:02:00.0	Off		0		0
30%	80 C	P0	Off / Off	4% 208MB / 5375MB			99%	E. Process	
<hr/>									
1.	Quadro 6000			0000:85:00.0	On		0		0
36%	84 C	P8	Off / Off	0% 22MB / 5375MB			0%	E. Process	
<hr/>									
Compute processes:									
GPU	PID	Process name		usage					
<hr/>									
0.	30234	nbody		process running on GPU					
<hr/>									

- **Create batch compute job for LSF**
- **Select appropriate queue to get scheduled on GPU-cluster**
-q gpu (over night), -a gpu (over day)
- **Exclusive nodes**
 - Nodes are allocated exclusively → at least 2 GPUs for one job
 - Please use resources reasonably!
- **Submit your job**
`bsub < myGPUScript.sh`
 - Starts running as soon as: batch mode starts and job is scheduled
- **Display pending reason: `bjobs -p`**
 - During daytime: `Dispatch windows closed`
- **More documentation**
 - RWTH Compute Cluster User's Guide: <https://doc.itc.rwth-aachen.de/display/CC/Using+the+batch+system>