

# Parallel Programming

## Processes and Threads

**Diego Fabregat-Traver** and Prof. Paolo Bientinesi

HPAC, RWTH Aachen  
`fabregat@ices.rwth-aachen.de`

WS16/17



## References

- Modern Operating Systems, 4th Edition. *Andrew S. Tanenbaum, Herbert Bos*. Chapters 1.5, 2.1, and 2.2.

Only if you want to know more. This slides are more than enough for this course!

# Outline


---

- 1 Concurrency
- 2 Processes
- 3 Threads
- 4 MPI
- 5 OpenMP

# Concurrency

---


- A property of computing systems in which several tasks are executing simultaneously
- Tasks are *in progress* at the same time
- Maybe running on one single processor, maybe on more than one
- Typical examples: web server, multiple programs running in your desktop, ...



# Concurrency

## Time-sharing or Multitasking systems


- CPU executes multiple processes by switching among them
- Switching occurs often enough for users to interact with each program while running
- In multi-core / multi-computer, processes may indeed be running in parallel.



# Process

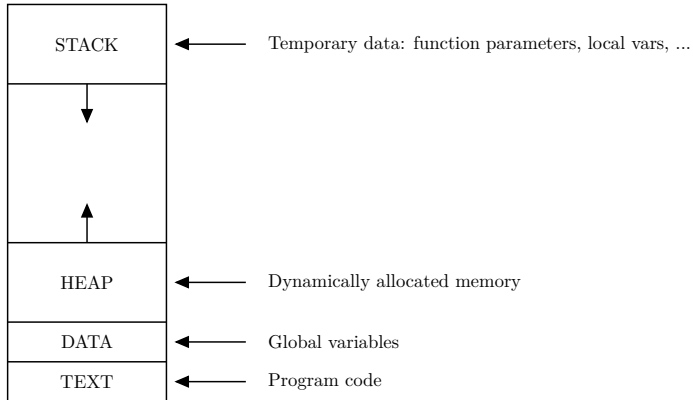
---

- A process is an instance of a program in execution
- States of a process
  - New: the process is being created
  - Ready: waiting to be assigned to a processor
  - Running: instructions are being executed
  - Waiting: waiting for some event to occur (e.g., I/O completion)
  - Terminated: has finished execution



# Process

- Associated address space
  - Program itself (text section)
  - Program's data (data section)
  - Stack, heap



# Process

---

- Process control block
  - Process ID
  - Process status
  - CPU registers (PC, ...)
  - Open files, memory management, ...
- Stores context to ensure a process can continue its execution properly after switching by restoring this context.



# Thread

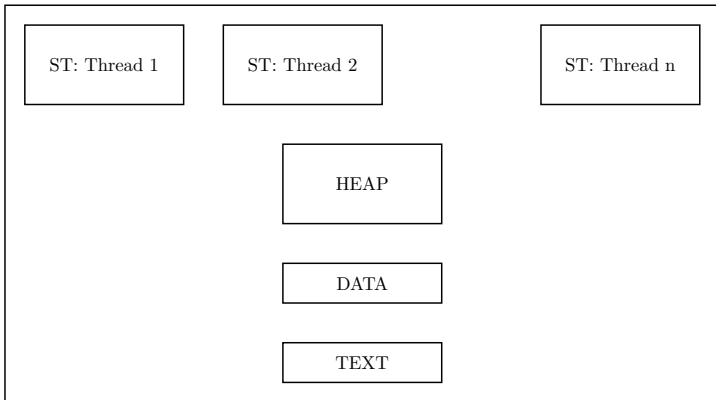
---

- Basic unit of CPU utilization
  - Flow of control within a process
- A thread includes
  - Thread ID
  - Program counter
  - Register set
  - Stack
- Shares resources with other threads within the same process
  - Text section
  - Data section
  - Other OS resources (open files, ...)

# Thread

---

## Single-threaded vs multi-threaded



# Processes vs Threads

---

- Processes:
  - Independent
  - Have separate address spaces
  - Heavier (carry more information)
  - Creation, context switching, ... is more expensive
  - Communicate via system-provided inter-process communication mechanisms
- Threads:
  - Exist within a process
  - Share address spaces
  - Lighter (faster creation, context switchin, ...)
  - Communicate via shared variables

# Programming models

# Single Program Multiple Data

---

- Most common programming model
- The same program is executed on multiple processors
- Different control flow based on the process/thread ID

```
if (process_id == 0)
    do_something()
else
    do_something_else()
```

# Message Passing

---

- Multiple processes (not necessarily running on different nodes)
- Each with its own private address spaces
- Access to (remote) data of other processes via sending/receiving messages (explicit communication)

```
if (process_id == SENDER)
    send_to(RECEIVER, data);

if (process_id == RECEIVER)
    recv_from(SENDER, data);
```

- Well-suited for distributed memory
- MPI (*Message Passing Interface*) is the *de-facto* standard

# Message Passing Interface (MPI)

---

- Specified and managed by the MPI Forum
  - Library offers a collection of communication primitives
  - Language bindings for C/C++ and Fortran
  - [www.mpi-forum.org](http://www.mpi-forum.org)
- Relatively low-level programming model
  - Data distribution and communication must be done manually
  - Primitives are easy to use, but designing parallel programs is hard
- Communication modes
  - Point-to-point (messages between two processors)
  - Collective (messages among groups of processors)
    - $1 \rightarrow n$  (e.g., broadcast)
    - $n \rightarrow 1$  (e.g., reduce)
    - $n \rightarrow n$  (e.g., allreduce )

- OpenMP
  - Higher level interface based on:
    - compiler directives
    - library routines
    - runtime
  - Emphasis on high-performance computing
- IEEE POSIX Threads (PThreads)
  - Standard UNIX threading API. Also used in Windows.
  - Over 60 functions: `pthread_create`, `pthread_join`, `pthread_exit`, etc.




- Specified and managed by the OpenMP ARB
- Assumes shared memory
- Allows definition of shared/private variables
- Language extensions based on:
  - Compiler directives
  - Library of routines
  - Runtime for the creation and management of threads

```
#pragma omp parallel for
for( i = 0; i < n; i++ )
    z[i] = a * x[i] + y[i]
```

- Currently available for C/C++ and Fortran
- [www.openmp.org](http://www.openmp.org)

# Hybrid programming

- Multiple processes, each spawning a number of threads
  - Inter-process communication via message passing (MPI)
  - Intra-process (thread) communication via shared memory
- Especially well-suited for hybrid architectures. For instance:
  - one process per shared-memory node, and
  - one thread per core



# Summary

---

- Process: instance of a program in execution
  - Container: code, data, process control block
  - Independent, heavier than threads
  - Communicate via inter-process mechanisms
- Threads: Unit of execution within a process
  - Share code and global address space
  - Private stack, lightweight
  - Communicate via shared variables
- Single Process Multiple Data (SPMD) paradigm
  - Message Passing via MPI
  - Shared-memory multithreading via OpenMP