

Parallel Programming

Prof. **Paolo Bientinesi**

`pauldj@ices.rwth-aachen.de`

WS16/17



Preliminaries

- Distributed-memory architecture
- Topologies
- Back in the days: Node \equiv CPU \equiv process
- Nowadays: NODE \rightarrow CPUs \rightarrow multi cores \rightarrow many processes
- Assumption: fully connected topology
- Assumption: each process can simultaneously send and receive
- Assumption: messages in opposite directions do not cause a conflict

What is “MPI”?

- A **library**, not a language, not a program.

"Minimal" MPI

<code>MPI_Init(...)</code>	MPI Initialization
<code>MPI_Comm_size(...)</code>	How many processes are there?
<code>MPI_Comm_rank(...)</code>	What rank am I?
<code>MPI_Send(...)</code>	Send data to another process
<code>MPI_Recv(...)</code>	Receive data from another process
<code>MPI_Finalize()</code>	MPI termination

What is “MPI”?

- A **library**, not a language, not a program.
- In fact, it's the specification of a library, not the actual implementation.
- MPI defines the interface, the functionality and the semantics of functions that deliver a message passing mechanism.
- Idea: clear separation between data communication and application.
- Both open-source and proprietary implementations.
- De-facto standard for distributed-memory parallelism.
- www.mpi-forum.org

`int MPI_Init(...)`

- `MPI_Init(&argc, &argv);`
- First MPI function
- Args not specified; an implementation might use them
- Query: `MPI_Initialized`

`int MPI_Finalize()`

- Last MPI function
- No arguments
- Query: `MPI_Finalized`

```
int MPI_Comm_size(MPI_Comm comm, int *size)
```

- Returns the number of processes in the communicator `comm`
- Communicator: for now `MPI_COMM_WORLD` \equiv “everybody”

```
int MPI_Comm_rank(MPI_Comm comm, int *rank)
```

- Returns the rank of the calling process within the communicator
- The rank is THE unique process identifier!
- NOTE: each process (rank) can be multi-threaded

Send ↔ Recv

- Objective: data movement
- MPI_Send and MPI_Recv must be matched
- Blocking communication

Necessary information:

Send ↔ Recv

- Objective: data movement
- MPI_Send and MPI_Recv must be matched
- Blocking communication

Necessary information:

Send	Recv
dest	source
*buffer	*target
size	size
datatype	datatype
tag	tag
comm	comm

```
int MPI_Send(*buffer, count, datatype, dest, tag, comm)
```

- `*buffer` is an address!
- `count` is indispensable; so is `datatype`
- `dest` is a rank (in `comm`)
- `tag` is an integer

```
int MPI_Recv(*target, count, datatype, source, tag, comm, *status)
```

- `*target`, `datatype` as for the `Send`
- `count` is the size of `target`. Actual size: `MPI_Get_count`
- `source` is either a rank (in `comm`) or `MPI_ANY_SOURCE`
- `tag` is either an integer or `MPI_ANY_TAG`
- `*status` on exit, contains info about the message

Before 1994

- Before MPI, no standards
- Different computers, different needs
⇒ **many** message passing environments
- N-cube, P4, PICL, PVM, ISIS, Express, Zipcode; Intel NX, IBM EUI, IBM CCL, ...
- **A lot of duplication!**
- **No portability whatsoever**

History

- **[1992]** First “MPI Forum” meeting (Supercomputing '92)
- **[1993–94]** Seven “MPI Forum” meetings. Working on the MPI standard

History

- **[1992]** First “MPI Forum” meeting (Supercomputing '92)
- **[1993–94]** Seven “MPI Forum” meetings. Working on the MPI standard
- **[1994]** Release of the first MPI standard: MPI-1
- **[1995]** First implementations of MPI: MPICH, LAM MPI, ...

History

- **[1992]** First “MPI Forum” meeting (Supercomputing '92)
- **[1993–94]** Seven “MPI Forum” meetings. Working on the MPI standard
- **[1994]** Release of the first MPI standard: MPI-1
- **[1995]** First implementations of MPI: MPICH, LAM MPI, ...
- **[1998]** Release of the second MPI standard: MPI-2
More than 100 new functions!
- **[2002]** Complete implementations of MPI-2
Dynamic process management, 1-sided communication, MPI-I/O

History

- **[1992]** First “MPI Forum” meeting (Supercomputing '92)
- **[1993–94]** Seven “MPI Forum” meetings. Working on the MPI standard
- **[1994]** Release of the first MPI standard: MPI-1
- **[1995]** First implementations of MPI: MPICH, LAM MPI, ...
- **[1998]** Release of the second MPI standard: MPI-2
More than 100 new functions!
- **[2002]** Complete implementations of MPI-2
Dynamic process management, 1-sided communication, MPI-I/O
- **[2012]** Release of MPI-3
Non-blocking collectives, sparse collectives, ...

Thanks to Jesper Larsson Träff (TU Wien).