

Parallel Programming

Prof. **Paolo Bientinesi**

`pauldj@aices.rwth-aachen.de`

WS 16/17



Dense Matrix-Vector Product: GEMV

1D matrix distribution

$$y := Ax, \quad x, y \in \mathbb{R}^n, \quad A \in \mathbb{R}^{n \times n}$$

Dense Matrix-Vector Product: GEMV

1D matrix distribution

$$y := Ax, \quad x, y \in \mathbb{R}^n, \quad A \in \mathbb{R}^{n \times n}$$

- A is partitioned along the rows, and distributed over p processes.
Process i owns block A_i .

$$A = \left[\begin{array}{c} \hline A_0 \\ \hline A_1 \\ \hline \vdots \\ \hline A_{p-1} \\ \hline \end{array} \right],$$

Dense Matrix-Vector Product: GEMV

1D matrix distribution

$$y := Ax, \quad x, y \in \mathbb{R}^n, \quad A \in \mathbb{R}^{n \times n}$$

- A is partitioned along the rows, and distributed over p processes. Process i owns block A_i .
- x is partitioned in p blocks; process i owns block x_i .

$$A = \begin{bmatrix} \overline{A_0} \\ \overline{A_1} \\ \vdots \\ \overline{A_{p-1}} \end{bmatrix}, \quad x = \begin{pmatrix} \overline{x_0} \\ \overline{x_1} \\ \vdots \\ \overline{x_{p-1}} \end{pmatrix},$$

Dense Matrix-Vector Product: GEMV

1D matrix distribution

$$y := Ax, \quad x, y \in \mathbb{R}^n, \quad A \in \mathbb{R}^{n \times n}$$

- A is partitioned along the rows, and distributed over p processes. Process i owns block A_i .
- x is partitioned in p blocks; process i owns block x_i .
- **Goal:** Compute y , and distribute it the same way as x .
Rationale: y overwrites x .

$$A = \begin{bmatrix} \frac{A_0}{\hline} \\ \frac{A_1}{\hline} \\ \vdots \\ \frac{A_{p-1}}{\hline} \end{bmatrix}, \quad x = \begin{pmatrix} \frac{x_0}{\hline} \\ \frac{x_1}{\hline} \\ \vdots \\ \frac{x_{p-1}}{\hline} \end{pmatrix}, \quad y = \begin{pmatrix} \frac{y_0}{\hline} \\ \frac{y_1}{\hline} \\ \vdots \\ \frac{y_{p-1}}{\hline} \end{pmatrix}$$

Dense Matrix-Vector Product: GEMV

1D matrix distribution

$$y := Ax, \quad x, y \in \mathbb{R}^n, \quad A \in \mathbb{R}^{n \times n}$$

- A is partitioned along the rows, and distributed over p processes. Process i owns block A_i .
- x is partitioned in p blocks; process i owns block x_i .
- **Goal:** Compute y , and distribute it the same way as x .
Rationale: y overwrites x .

$$A = \begin{bmatrix} \overline{A_0} \\ \overline{A_1} \\ \vdots \\ \overline{A_{p-1}} \end{bmatrix}, \quad x = \begin{pmatrix} \overline{x_0} \\ \overline{x_1} \\ \vdots \\ \overline{x_{p-1}} \end{pmatrix}, \quad y = \begin{pmatrix} \overline{y_0} \\ \overline{y_1} \\ \vdots \\ \overline{y_{p-1}} \end{pmatrix}$$

Note: A_i, x_i and y_i indicate a *block* of rows, not a single row.

GEMV: Cost

1D matrix distribution

Algorithm

GEMV: Cost

1D matrix distribution

Algorithm

- 1 $x = \text{Allgather}(x_i)$ x becomes available to every process
- 2 $y_i = A_i x$ local computation

Parallel cost (lower bound for $T_p(n)$)

GEMV: Cost

1D matrix distribution

Algorithm

- 1 $x = \text{Allgather}(x_i)$ x becomes available to every process
- 2 $y_i = A_i x$ local computation

Parallel cost (lower bound for $T_p(n)$)

- 1 $\lceil \log_2(p) \rceil \alpha + \frac{p-1}{p} n \beta \approx \log_2(p) \alpha + n \beta$
- 2 $2 \frac{n^2}{p} \gamma$

Sequential cost

GEMV: Cost

1D matrix distribution

Algorithm

- 1 $x = \text{Allgather}(x_i)$ x becomes available to every process
- 2 $y_i = A_i x$ local computation

Parallel cost (lower bound for $T_p(n)$)

- 1 $\lceil \log_2(p) \rceil \alpha + \frac{p-1}{p} n \beta \approx \log_2(p) \alpha + n \beta$
- 2 $2 \frac{n^2}{p} \gamma$

Sequential cost

- $T_1(n) = 2n^2 \gamma$

GEMV: Scalability

1D matrix distribution

Speedup

GEMV: Scalability

1D matrix distribution

Speedup

$$S_p(n) = \frac{T_1(n)}{T_p(n)} = \frac{2n^2\gamma}{\log_2(p)\alpha + n\beta + 2\frac{n^2}{p}\gamma}$$

Efficiency

GEMV: Scalability

1D matrix distribution

Speedup

$$S_p(n) = \frac{T_1(n)}{T_p(n)} = \frac{2n^2\gamma}{\log_2(p)\alpha + n\beta + 2\frac{n^2}{p}\gamma}$$

Efficiency

$$E_p(n) = \frac{S_p(n)}{p} = \frac{1}{1 + \frac{p \log_2(p)}{2n^2} \frac{\alpha}{\gamma} + \frac{p}{2n} \frac{\beta}{\gamma}}$$

Strong scalability

GEMV: Scalability

1D matrix distribution

Speedup

$$S_p(n) = \frac{T_1(n)}{T_p(n)} = \frac{2n^2\gamma}{\log_2(p)\alpha + n\beta + 2\frac{n^2}{p}\gamma}$$

Efficiency

$$E_p(n) = \frac{S_p(n)}{p} = \frac{1}{1 + \frac{p \log_2(p)}{2n^2} \frac{\alpha}{\gamma} + \frac{p}{2n} \frac{\beta}{\gamma}}$$

Strong scalability

$$\lim_{p \rightarrow \infty} E_p(n) = 0 \quad \times$$

Weak scalability

local memory = $M \Rightarrow$ combined memory = Mp
largest problem (n_M) solvable with p processes?
 $n_M^2 = Mp$

GEMV: Scalability

1D matrix distribution

Speedup

$$S_p(n) = \frac{T_1(n)}{T_p(n)} = \frac{2n^2\gamma}{\log_2(p)\alpha + n\beta + 2\frac{n^2}{p}\gamma}$$

Efficiency

$$E_p(n) = \frac{S_p(n)}{p} = \frac{1}{1 + \frac{p \log_2(p)}{2n^2} \frac{\alpha}{\gamma} + \frac{p}{2n} \frac{\beta}{\gamma}}$$

Strong scalability

$$\lim_{p \rightarrow \infty} E_p(n) = 0 \quad \times$$

Weak scalability

local memory = $M \Rightarrow$ combined memory = Mp
largest problem (n_M) solvable with p processes?
 $n_M^2 = Mp$

$$\lim_{p \rightarrow \infty} E_p(n_M) = \frac{1}{1 + \frac{\log_2(p)}{2M} \frac{\alpha}{\gamma} + \frac{\sqrt{p}}{2\sqrt{M}} \frac{\beta}{\gamma}} = 0 \quad \times$$

GEMV, distribution by columns

- A is partitioned along the columns, and distributed over p processes. Process i owns block A_i . Note that A_i now denotes a block of columns.
- x is partitioned as before; process i owns x_i .
- **Goal:** Compute y , and distribute it the same way as x .

$$A = \left[\begin{array}{c|c|c|c} A_0 & A_1 & \dots & A_{p-1} \end{array} \right], \quad x = \begin{pmatrix} \frac{x_0}{x_1} \\ \vdots \\ \frac{x_{p-1}}{x_{p-1}} \end{pmatrix}, \quad y = \begin{pmatrix} \frac{y_0}{y_1} \\ \vdots \\ \frac{y_{p-1}}{y_{p-1}} \end{pmatrix}$$

Cost & Scalability

Algorithm

Cost & Scalability

Algorithm

① $y^{(i)} = A_i x_i$

local computation

② $y = \text{Reduce-Scatter}(y^{(i)})$

reduction: sum of $y^{(i)}$'s
+ scatter

Parallel cost

Cost & Scalability

Algorithm

① $y^{(i)} = A_i x_i$

local computation

② $y = \text{Reduce-Scatter}(y^{(i)})$

reduction: sum of $y^{(i)}$'s
+ scatter

Parallel cost

① $2 \frac{n^2}{p} \gamma$

② $\lceil \log_2(p) \rceil \alpha + \frac{p-1}{p} n \beta + \frac{p-1}{p} n \gamma \approx \log_2(p) \alpha + n(\beta + \gamma)$

Analysis

Cost & Scalability

Algorithm

① $y^{(i)} = A_i x_i$

local computation

② $y = \text{Reduce-Scatter}(y^{(i)})$

reduction: sum of $y^{(i)}$'s
+ scatter

Parallel cost

① $2 \frac{n^2}{p} \gamma$

② $\lceil \log_2(p) \rceil \alpha + \frac{p-1}{p} n \beta + \frac{p-1}{p} n \gamma \approx \log_2(p) \alpha + n(\beta + \gamma)$

Analysis

Compared to the previous case (A partitioned by rows), $T_p(n)$ now has one extra term ($n\gamma$); this algorithm is therefore also not scalable ❌

GEMV: 2D matrix distribution

- A is partitioned according to a 2D mesh of processes, with $p = r \times c$.
 P_{ij} , the (i, j) process in the mesh, owns block A_{ij} .

$$\begin{pmatrix} P_{00} & P_{01} & \cdots & P_{0,(c-1)} \\ P_{10} & P_{11} & \cdots & P_{1,(c-1)} \\ \vdots & \vdots & \ddots & \vdots \\ P_{(r-1),0} & P_{(r-1),1} & \cdots & P_{(r-1),(c-1)} \end{pmatrix} \leftarrow \begin{pmatrix} A_{00} & A_{01} & \cdots & A_{0,(c-1)} \\ A_{10} & A_{11} & \cdots & A_{1,(c-1)} \\ \vdots & \vdots & \ddots & \vdots \\ A_{(r-1),0} & A_{(r-1),1} & \cdots & A_{(r-1),(c-1)} \end{pmatrix}$$

- As before, x and y are partitioned in p blocks;
 x is mapped to the mesh by columns, y by rows.

Example: $p = 2 \times 3$

$$\left(\begin{array}{c|c|c} A_{00} & A_{01} & A_{02} \\ \hline A_{10} & A_{11} & A_{12} \end{array} \right), \quad \left(\begin{array}{c|c|c} x_0 & x_2 & x_4 \\ \hline x_1 & x_3 & x_5 \end{array} \right), \quad \left(\begin{array}{c|c|c} y_0 & y_1 & y_2 \\ \hline y_3 & y_4 & y_5 \end{array} \right)$$

Cost

2D matrix distribution

Algorithm

Cost

2D matrix distribution

Algorithm

- 1 $x_I = \text{Allgather}(x_i)$ within columns
- 2 $y_J = A_{ij}x_I$
- 3 $y_j = \text{Reduce-scatter}(y_J)$ within rows

x_I is a “block”

local computation

Parallel cost (lower bound for $T_p(n)$)

Cost

2D matrix distribution

Algorithm

- 1 $x_I = \text{Allgather}(x_i)$ within columns x_I is a “block”
- 2 $y_J = A_{ij}x_I$ local computation
- 3 $y_j = \text{Reduce-scatter}(y_J)$ within rows

Parallel cost (lower bound for $T_p(n)$)

- 1 $\lceil \log_2(r) \rceil \alpha + \frac{r-1}{p} n \beta \approx \log_2(r) \alpha + \frac{n}{c} \beta$
- 2 $2 \frac{n^2}{p} \gamma$
- 3 $\lceil \log_2(c) \rceil \alpha + \frac{c-1}{p} n \beta + \frac{c-1}{p} n \gamma \approx \log_2(c) \alpha + \frac{n}{r} \beta + \frac{n}{r} \gamma$

Sequential cost

- $T_1(n) = 2n^2 \gamma$

Scalability

2D matrix distribution

Parallel cost

$$2\frac{n^2}{p}\gamma + (\log_2(r) + \log_2(c))\alpha + \left(\frac{n}{c} + \frac{n}{r}\right)\beta + \frac{n}{r}\gamma$$

Scalability

2D matrix distribution

Parallel cost $2 \frac{n^2}{p} \gamma + (\log_2(r) + \log_2(c)) \alpha + \left(\frac{n}{c} + \frac{n}{r} \right) \beta + \frac{n}{r} \gamma$

assuming $r = c = \sqrt{p}$ $2 \frac{n^2}{p} \gamma + \log_2(p) \alpha + \frac{n}{\sqrt{p}} (2\beta + \gamma)$

Speedup $S_p(n) = \frac{p}{1 + \frac{p \log_2(p)}{2n^2} \frac{\alpha}{\gamma} + \frac{\sqrt{p}}{2n} \frac{2\beta + \gamma}{\gamma}}$

Scalability

2D matrix distribution

Parallel cost $2 \frac{n^2}{p} \gamma + (\log_2(r) + \log_2(c)) \alpha + \left(\frac{n}{c} + \frac{n}{r} \right) \beta + \frac{n}{r} \gamma$

assuming $r = c = \sqrt{p}$ $2 \frac{n^2}{p} \gamma + \log_2(p) \alpha + \frac{n}{\sqrt{p}} (2\beta + \gamma)$

Speedup $S_p(n) = \frac{p}{1 + \frac{p \log_2(p)}{2n^2} \frac{\alpha}{\gamma} + \frac{\sqrt{p}}{2n} \frac{2\beta + \gamma}{\gamma}}$

Efficiency $E_p(n) = S_p(n)/p = \dots$

Scalability

2D matrix distribution

Parallel cost $2 \frac{n^2}{p} \gamma + (\log_2(r) + \log_2(c)) \alpha + \left(\frac{n}{c} + \frac{n}{r} \right) \beta + \frac{n}{r} \gamma$

assuming $r = c = \sqrt{p}$ $2 \frac{n^2}{p} \gamma + \log_2(p) \alpha + \frac{n}{\sqrt{p}} (2\beta + \gamma)$

Speedup $S_p(n) = \frac{p}{1 + \frac{p \log_2(p)}{2n^2} \frac{\alpha}{\gamma} + \frac{\sqrt{p}}{2n} \frac{2\beta + \gamma}{\gamma}}$

Efficiency $E_p(n) = S_p(n)/p = \dots$

Strong scalability $\lim_{p \rightarrow \infty} E_p(n) = 0 \quad \times$

Scalability

2D matrix distribution

Parallel cost $2 \frac{n^2}{p} \gamma + (\log_2(r) + \log_2(c)) \alpha + \left(\frac{n}{c} + \frac{n}{r} \right) \beta + \frac{n}{r} \gamma$

assuming $r = c = \sqrt{p}$ $2 \frac{n^2}{p} \gamma + \log_2(p) \alpha + \frac{n}{\sqrt{p}} (2\beta + \gamma)$

Speedup $S_p(n) = \frac{p}{1 + \frac{p \log_2(p)}{2n^2} \frac{\alpha}{\gamma} + \frac{\sqrt{p}}{2n} \frac{2\beta + \gamma}{\gamma}}$

Efficiency $E_p(n) = S_p(n)/p = \dots$

Strong scalability $\lim_{p \rightarrow \infty} E_p(n) = 0 \quad \times$

Weak scalability
local memory = $M \Rightarrow$ combined memory = Mp
largest problem (n_M) solvable with p processes?
 $n_M^2 = Mp$

$$\lim_{p \rightarrow \infty} E_p(n_M) = \frac{1}{1 + \frac{\log_2(p)}{2M} \frac{\alpha}{\gamma} + \frac{1}{2\sqrt{M}} \frac{2\beta + \gamma}{\gamma}} = 0 \quad \times \quad \dots \quad \checkmark$$

Exercise

2D matrix distribution

- A is partitioned according to a 2D mesh of processes, with $p = r \times c$. Process P_{ij} owns block A_{ij} .
- x and y are partitioned in p blocks, and mapped to the mesh by columns.
- Sketch out the algorithm to compute y .
- Determine the cost of the algorithm.
- Study the scalability of the algorithm.

Parallel Matrix Distribution

How to store a (large) matrix using $p = r \times c$ processes?

- 1D, block of rows (or columns)
- 1D (block) cyclic, either by rows or columns
- 2D, $r \times c$ quadrants
- 2D (block) cyclic

Bad idea

Bad idea

Not so good

Good idea!

References

- Applet:
<http://acts.nersc.gov/scalapack/hands-on/datadist.html>
- “Introduction to High-Performance Scientific Computing”
by Victor Eijkhout (free download).
- “A Comprehensive Approach to Parallel Linear Algebra Libraries”
(Technical Report, University of Texas). Chapter 3, pages 19–40.