

Parallel Programming

Prof. **Paolo Bientinesi**

`pauldj@aices.rwth-aachen.de`

WS 16/17



Scenario

Process P_i owns matrix A_i , with $i = 0, \dots, p - 1$.

Objective

$\begin{cases} \text{Even}(i) : & \text{compute } T_i := A_i + A_{(i+1) \bmod p} \\ \text{Odd}(i) : & \text{compute } T_i := A_i - A_{(i+1) \bmod p} \end{cases}$

Scenario

Process P_i owns matrix A_i , with $i = 0, \dots, p - 1$.

Objective

$\begin{cases} \text{Even}(i) : & \text{compute } T_i := A_i + A_{(i+1) \bmod p} \\ \text{Odd}(i) : & \text{compute } T_i := A_i - A_{(i+1) \bmod p} \end{cases}$

Scenario

1D domain, logically split among p processes.

Objective

Run a finite difference scheme, e.g.,

$$u(x_i) := \frac{u(x_{i-1}) - 2u(x_i) + u(x_{i+1}))}{h^2}.$$

Scenario

Process P_i owns matrix A_i , with $i = 0, \dots, p - 1$.

Objective

$\begin{cases} \text{Even}(i) : & \text{compute } T_i := A_i + A_{(i+1) \bmod p} \\ \text{Odd}(i) : & \text{compute } T_i := A_i - A_{(i+1) \bmod p} \end{cases}$

Scenario

1D domain, logically split among p processes.

Objective

Run a finite difference scheme, e.g.,

$$u(x_i) := \frac{u(x_{i-1}) - 2u(x_i) + u(x_{i+1}))}{h^2}.$$

\Rightarrow **point-to-point communication**

Anatomy of MPI_Send and MPI_Recv

```
int MPI_Send(  
    *buffer, count, datatype,           ← “data”  
    destination, tag, communicator      ← “envelope”  
);
```

```
int MPI_Recv(  
    *buffer, count, datatype,           ← “data”  
    source, tag, commmunicator,        ← “envelope”  
    *status  
);
```

Anatomy of MPI_Send and MPI_Recv

```
int MPI_Send(  
    *buffer, count, datatype,           ← "data"  
    destination, tag, communicator      ← "envelope"  
);
```

```
int MPI_Recv(  
    *buffer, count, datatype,           ← "data"  
    source, tag, commmunicator,        ← "envelope"  
    *status  
);
```

message = data + envelope (+ info)
matching envelopes \Rightarrow data transfer

Note: Meaning of `count`: send \neq recv

`count` in send = size of message vs. `count` in receive = size of buffer.

Point-to-point communication

Send

- `MPI_Ssend`
- `MPI_Send`
- `MPI_Isend`
- `⋮`
- `MPI_Bsend`

Receive

- `MPI_Recv`
- `MPI_Irecv`

Send+Receive

- `MPI_Sendrecv`
- `MPI_Sendrecv_replace`

Send Modes

The stress is on the buffer being sent: “When I can I safely overwrite it?”

Send Modes

The stress is on the buffer being sent: “When I can I safely overwrite it?”

- `MPI_Ssend`: The program execution is blocked until a matching receive is posted. The buffer is usable as soon as the call completes.

Send Modes

The stress is on the buffer being sent: “When I can I safely overwrite it?”

- `MPI_Ssend`: The program execution is blocked until a matching receive is posted. The buffer is usable as soon as the call completes.
- `MPI_Send`: MPI attempts to copy the outgoing message onto a local (hidden) buffer. If possible, the execution continues and the send buffer is immediately usable, otherwise same as `Ssend`.

Send Modes

The stress is on the buffer being sent: “When I can I safely overwrite it?”

- `MPI_Ssend`: The program execution is blocked until a matching receive is posted. The buffer is usable as soon as the call completes.
- `MPI_Send`: MPI attempts to copy the outgoing message onto a local (hidden) buffer. If possible, the execution continues and the send buffer is immediately usable, otherwise same as `Ssend`.
- `MPI_Isend`: The execution continues immediately. The send buffer should not be accessed until the `MPI_request` allows it. To be used in conjunction with `MPI_Wait` or `MPI_Test*`.

Note: Careful with multithreading!!

*: See also `MPI_Waitany`, `MPI_Waitall`, `MPI_Waitsome`, `MPI_Testany`, `MPI_Testall`, `MPI_Testsome`.

Recv Modes

The stress is on the incoming buffer: “When I can I safely access it?”

Recv Modes

The stress is on the incoming buffer: “When I can I safely access it?”

- `MPI_Recv`: The program execution is blocked until a matching send is posted. The incoming buffer is usable as soon as the call completes.

Recv Modes

The stress is on the incoming buffer: “When I can I safely access it?”

- `MPI_Recv`: The program execution is blocked until a matching send is posted. The incoming buffer is usable as soon as the call completes.
- `MPI_Irecv`: The execution continues Immediately. The incoming buffer should not be modified until the `MPI_request` allows it. To be used in conjunction with `MPI_Wait` or `MPI_Test*`.

*: See also `MPI_Waitany`, `MPI_Waitall`, `MPI_Waitsome`, `MPI_Testany`, `MPI_Testall`, `MPI_Testsome`.